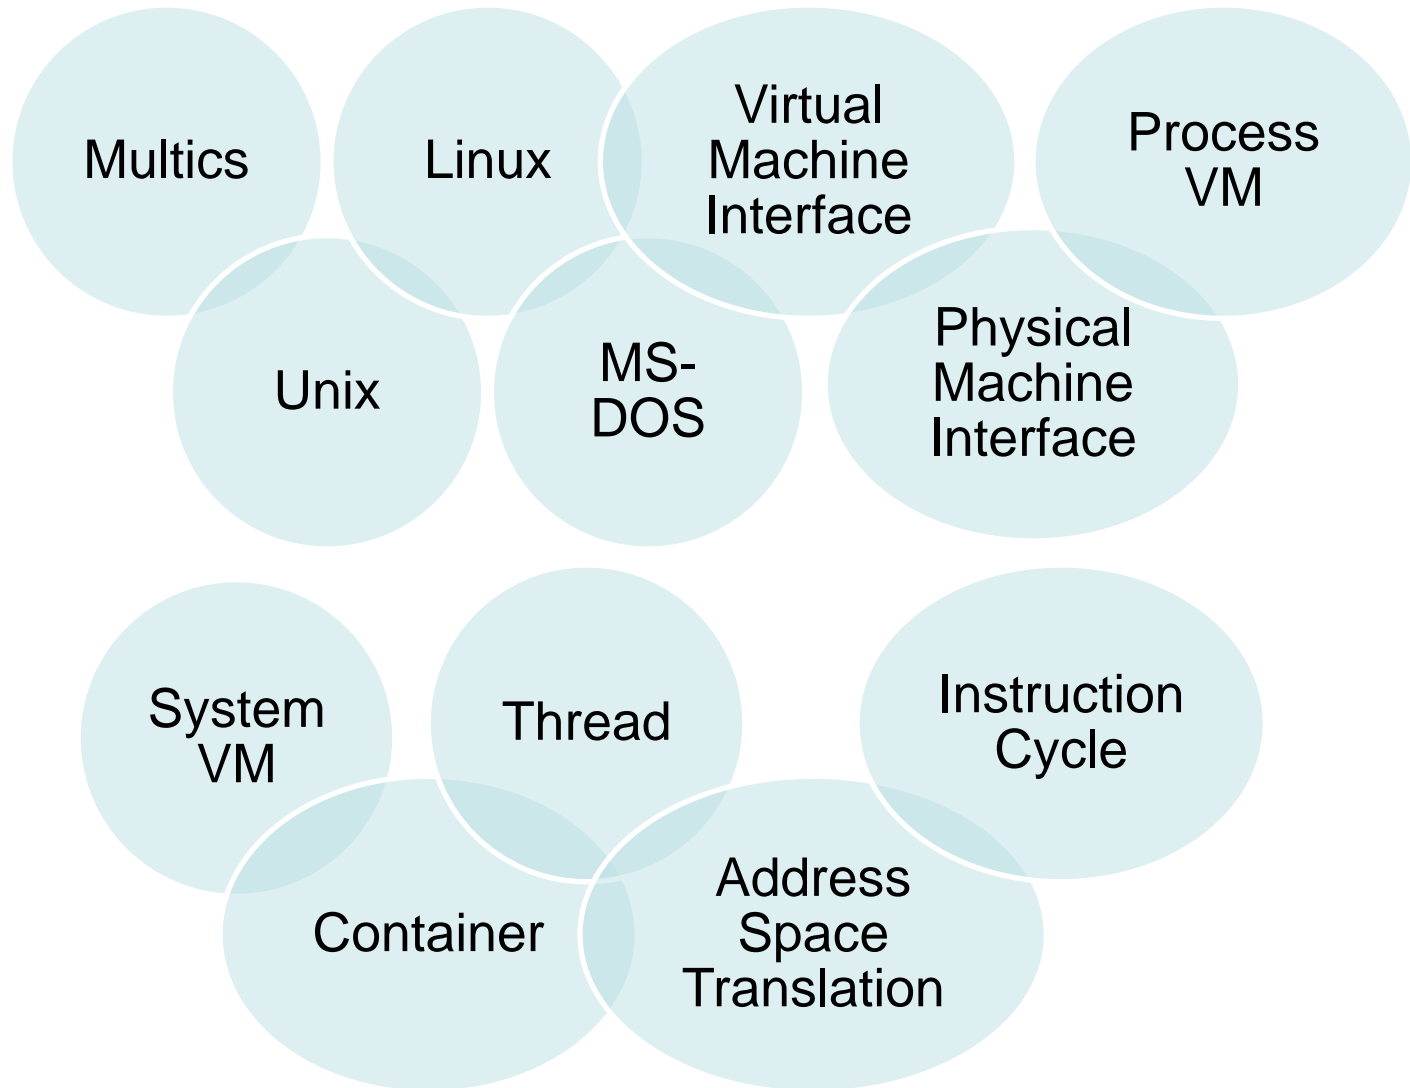


4 Concepts, Base & Bound

21 November 2024
Lecture 3

Slides adapted from John Kubiatowicz (UC Berkeley)

Concept Review



Topics for Today

- Four Concepts
 - Threads (last week)
 - Address space (continued)
 - Processes
 - Dual Mode Operation
- Protection Mechanisms
 - Address ranges: Base and Bound
- Interrupt and Process Context

Multiprogramming – Multiple Threads of Control

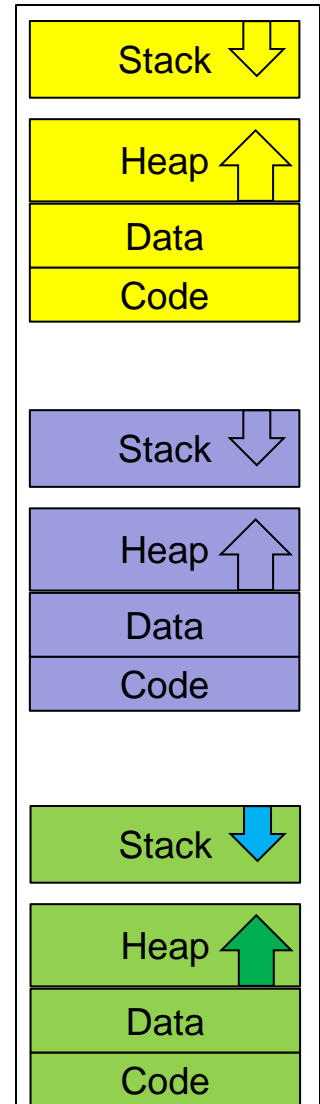
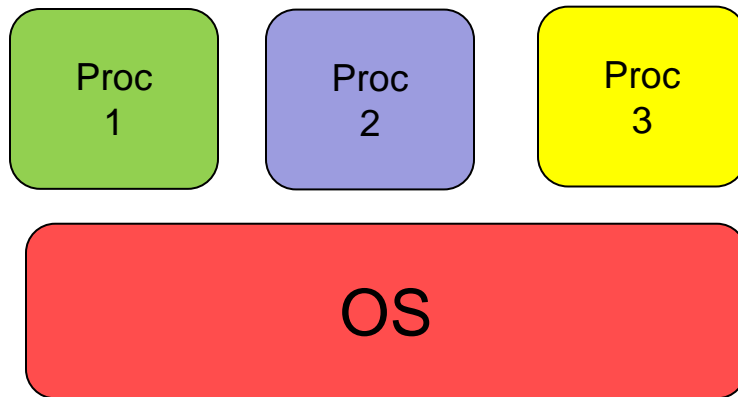




Image source: <http://moziru.com/images/game-clipart-musical-chair-5.gif>



Basic Problems of Concurrency

Basic problem of concurrency is resources:

- **Hardware**: single CPU, single DRAM, single I/O devices
- **Multiprogramming API**: processes think they have exclusive access to shared resources

OS must coordinate all activity

- Multiple processes, I/O interrupts, ...
- How can it keep all these things straight?

Basic Idea: Use Virtual Machine abstraction

- Simple machine abstraction for processes
- Multiplex these abstract machines

Dijkstra did this for the **“THE system”**

- Few thousand lines vs 1 million lines in OS 360 (1K bugs)

Simple Multiprogramming (little protection)



All virtual CPUs share **same non-CPU resources**

- I/O devices the same
- Memory the same

Consequence of sharing

- Each thread can access the data of **every other thread** (**good** for sharing, **bad** for protection)
- Threads can **share instructions** (**good** for sharing, **bad** for protection)
- Can threads overwrite OS functions?

The (**unprotected**) model found in

- Embedded applications
- Windows 3.1/Early Macintosh (switch only with *yield*)
- Windows 95-ME (switch with both *yield* and *timer*)

Clear divisions between entities

<https://www.flickr.com/photos/tabor-roeder/5404004415/in/photostream/>





Protection (stronger)

- OS protects itself from user programs
 - **Reliability**: Compromising the OS generally causes it to crash
 - **Security**: Limit the scope of what processes can do
 - **Privacy**: limit each process to the data it is permitted to access
 - **Fairness**: each should be limited to its appropriate share of system resources (CPU time, memory, I/O, etc)
- OS protects User programs from one another





Protection

- Primary Mechanism: limit the translation from program address space to physical memory space
 - Can only touch what is mapped into process *address space*



- Additional Mechanisms:
 - Privileged instructions, in/out instructions, special registers
 - syscall processing, subsystem implementation
 - (e.g., file access rights, etc)

Image source: <https://media.istockphoto.com/>

So Far

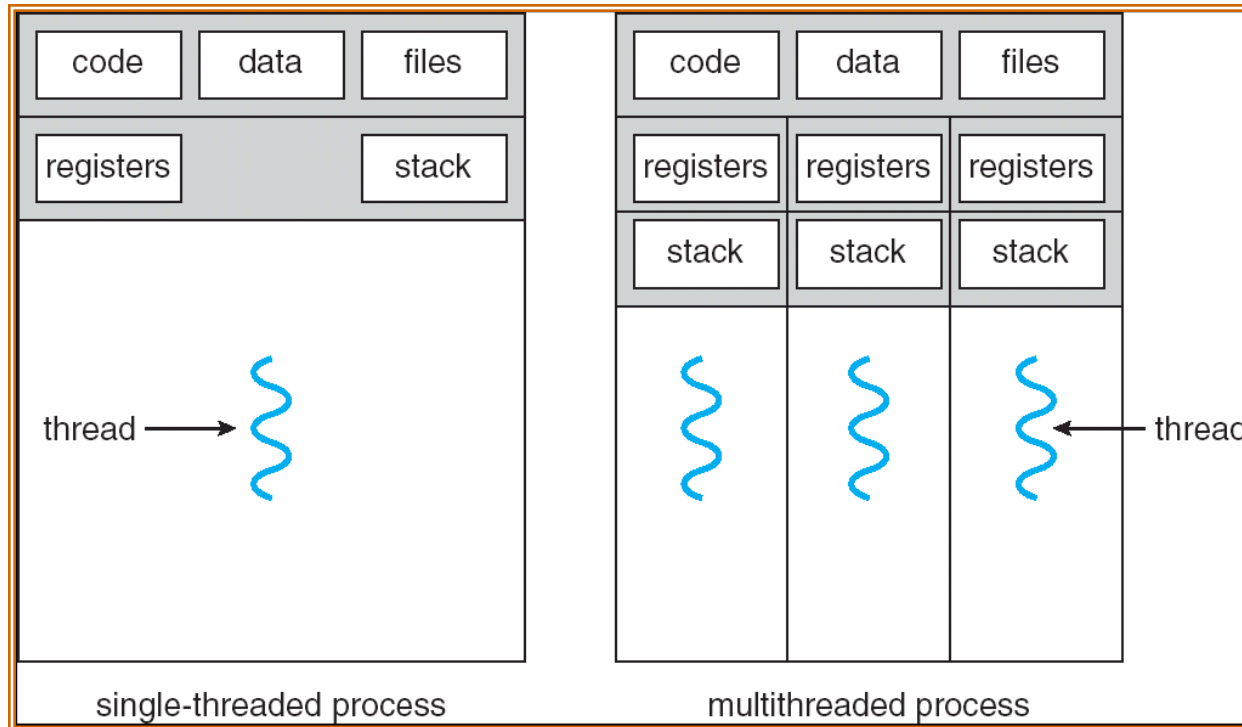
- Four Concepts
 - Threads (last week)
 - Address space (continued)
 - Processes
 - Dual Mode Operation
- Protection Mechanisms
 - Address ranges: Base and Bound
- Interrupt and Process Context

Third OS Concept: Process

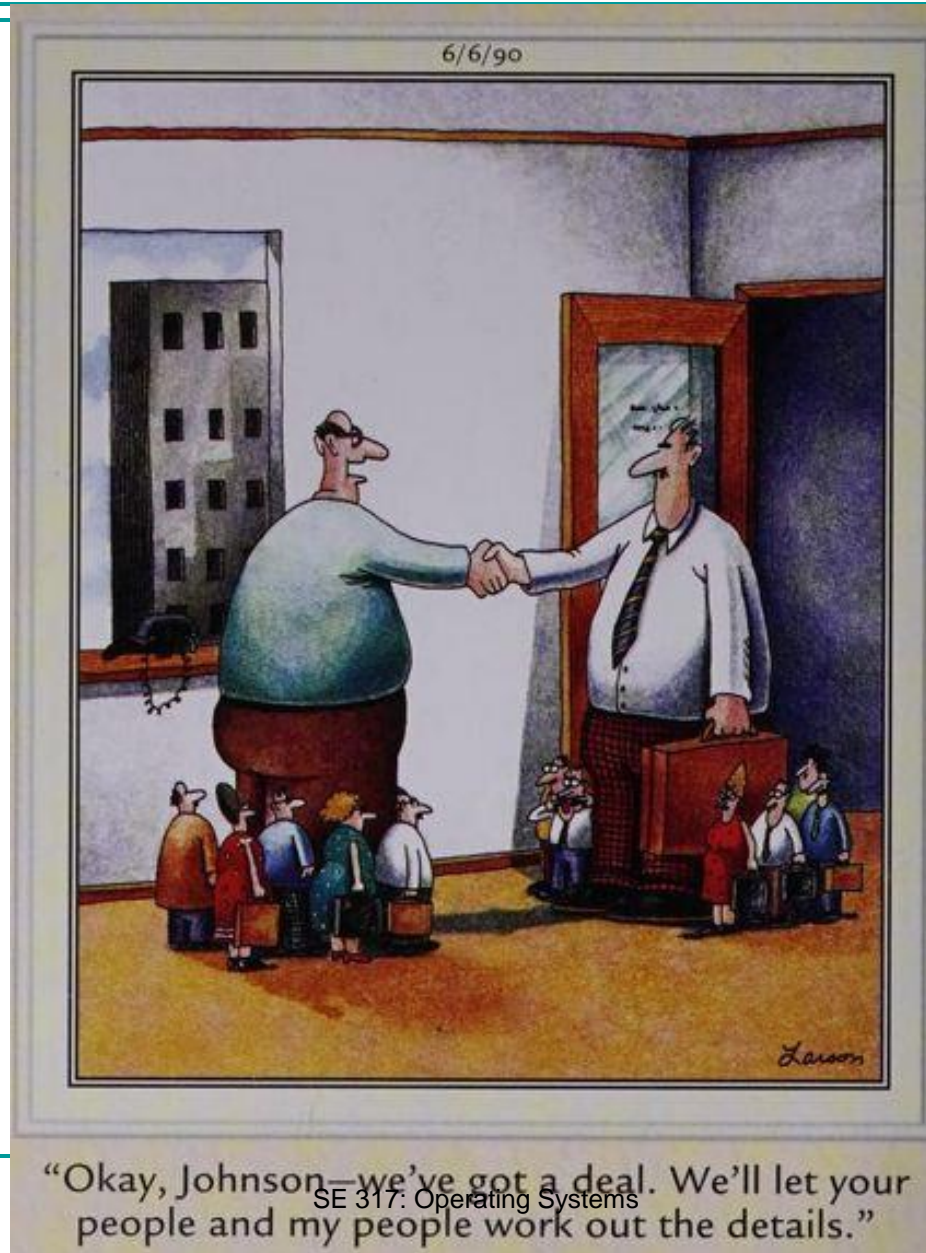


- **Process:** execution environment with Restricted Rights
 - **Address Space with One or More Threads**
 - Owns memory (address space), file descriptors, file system context, ...
 - Encapsulate one or more threads sharing process resources
- Why **processes**?
 - Protected from each other! OS Protected from them!
 - Processes provides memory protection, group threads
- Tradeoff between protection (process) and efficiency (thread)
 - Communication easier *within* a process
 - Communication harder *between* processes
- Application instance consists of one or more processes

Single and Multithreaded Processes



- Threads encapsulate concurrency: “Active” component
- Address spaces encapsulate protection: “Passive” part
 - Keeps buggy program from trashing the system
- Why have multiple threads per address space?



So Far

- Four Concepts
 - Threads (last week)
 - Address space (continued)
 - Processes
 - Dual Mode Operation
- Protection Mechanisms
 - Address ranges: Base and Bound
- Interrupt and Process Context



Fourth OS Concept: Dual Mode Operation

- **Hardware** provides at least two modes:
 - “**Kernel**” mode (or “supervisor” or “protected”)
 - “**User**” mode: Normal programs executed
- What hardware needed to support “**dual mode**” operation?

Bit of state (user/system mode bit)

Certain operations / actions **only permitted** in system/kernel mode



- In user mode they **fail or trap**

User → Kernel transition **sets system mode and saves the user PC**



- Operating system code carefully puts aside user state then performs the necessary operations

Kernel → User transition **clears system mode and restores appropriate user PC**

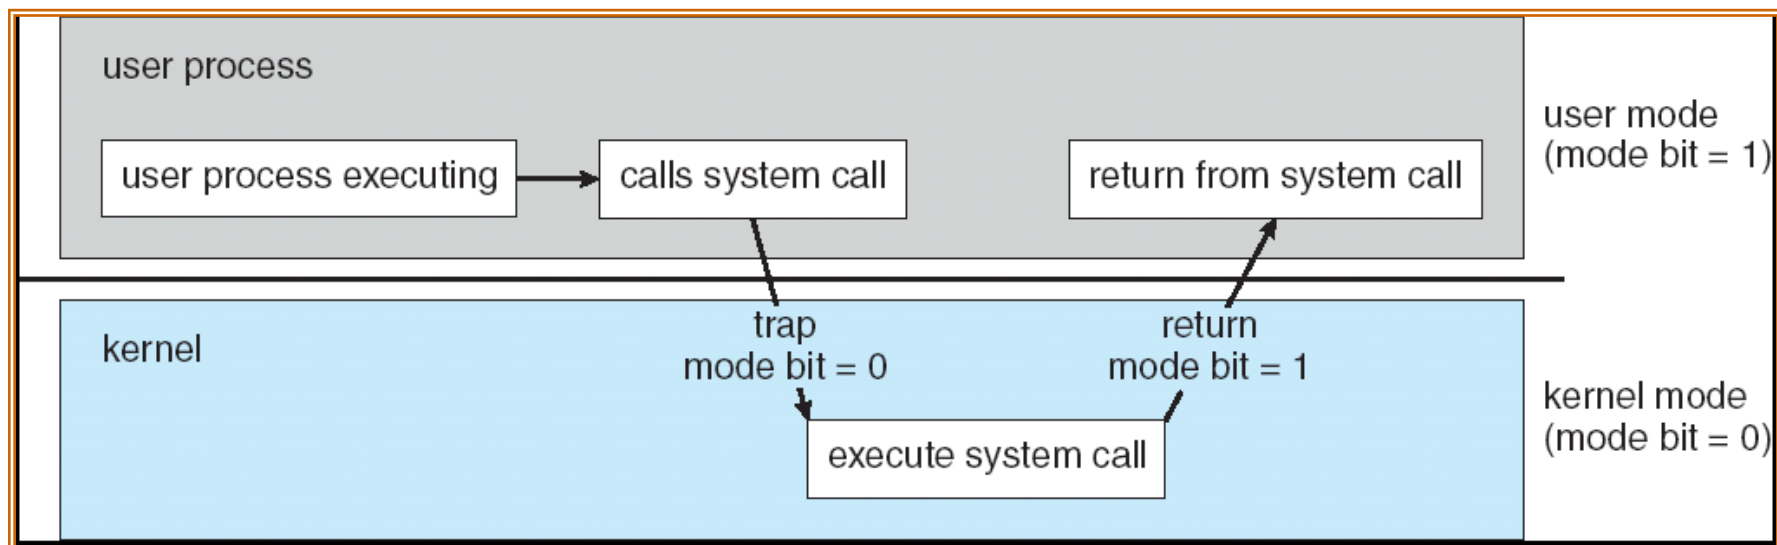


- return-from-interrupt



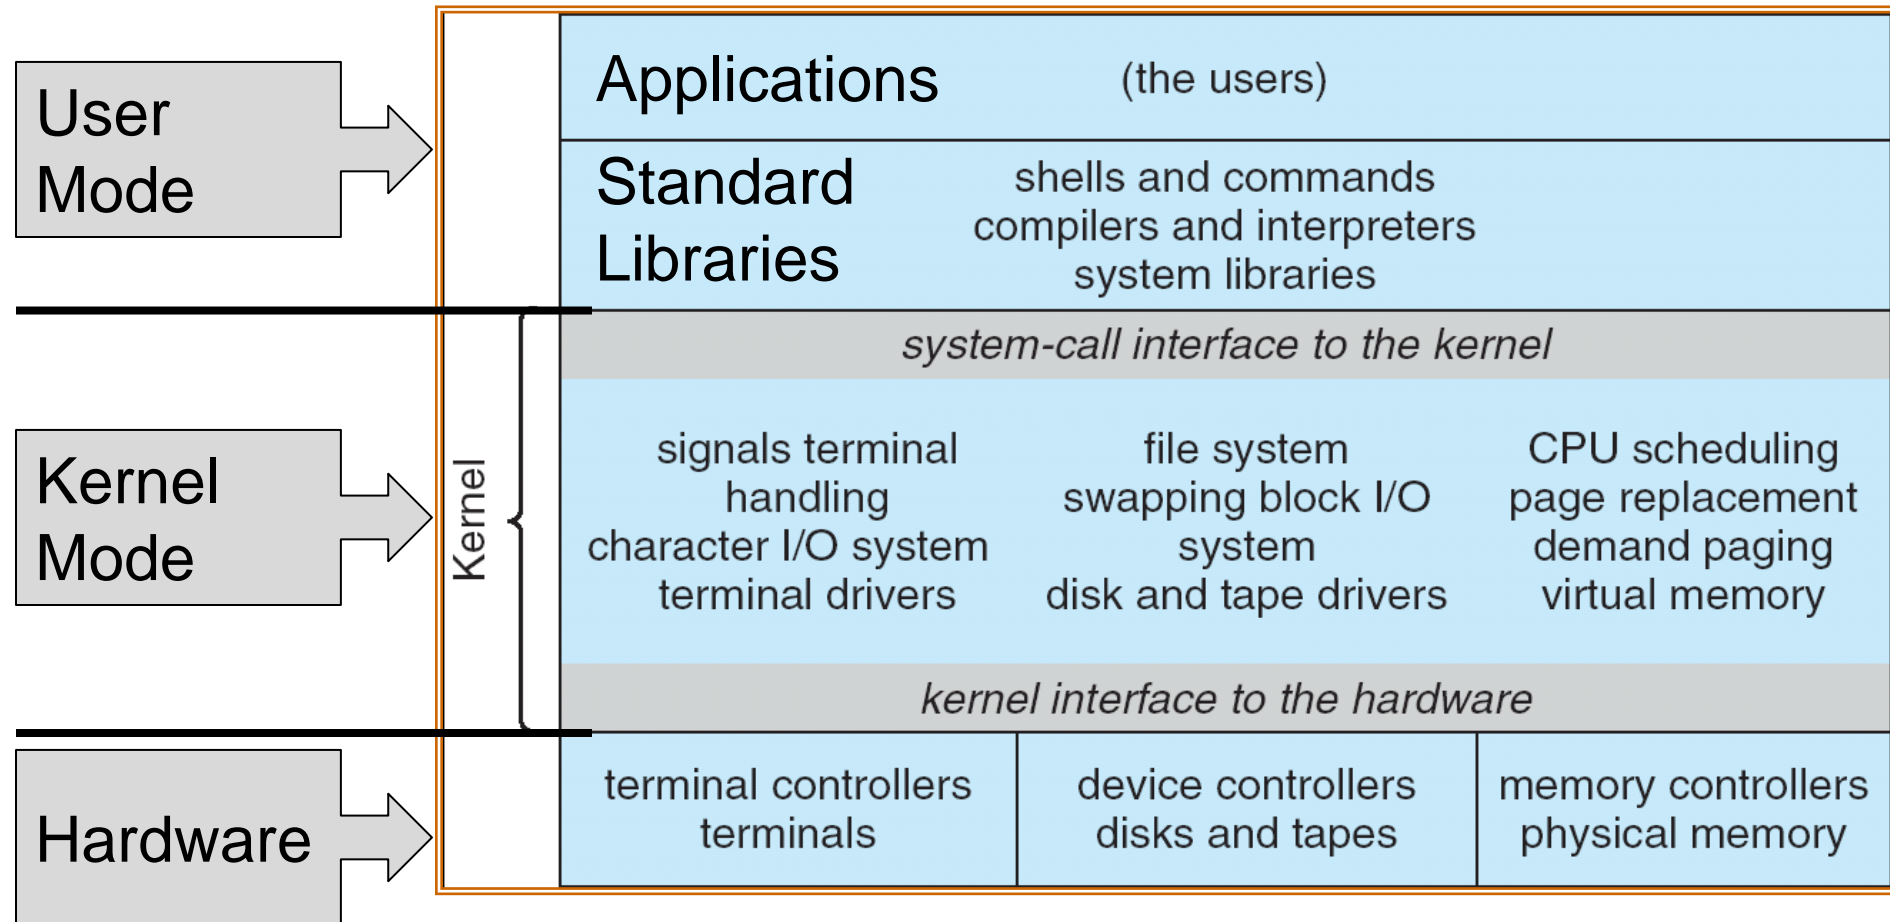
Dual Mode Operation

- Some instructions/ops **prohibited** in **user** mode:
 - Example: can't modify page tables in **user** mode
 - Attempt to modify \Rightarrow Exception generated
- Transitions from **user** mode to **kernel** mode:
 - System Calls, Interrupts, Other exceptions

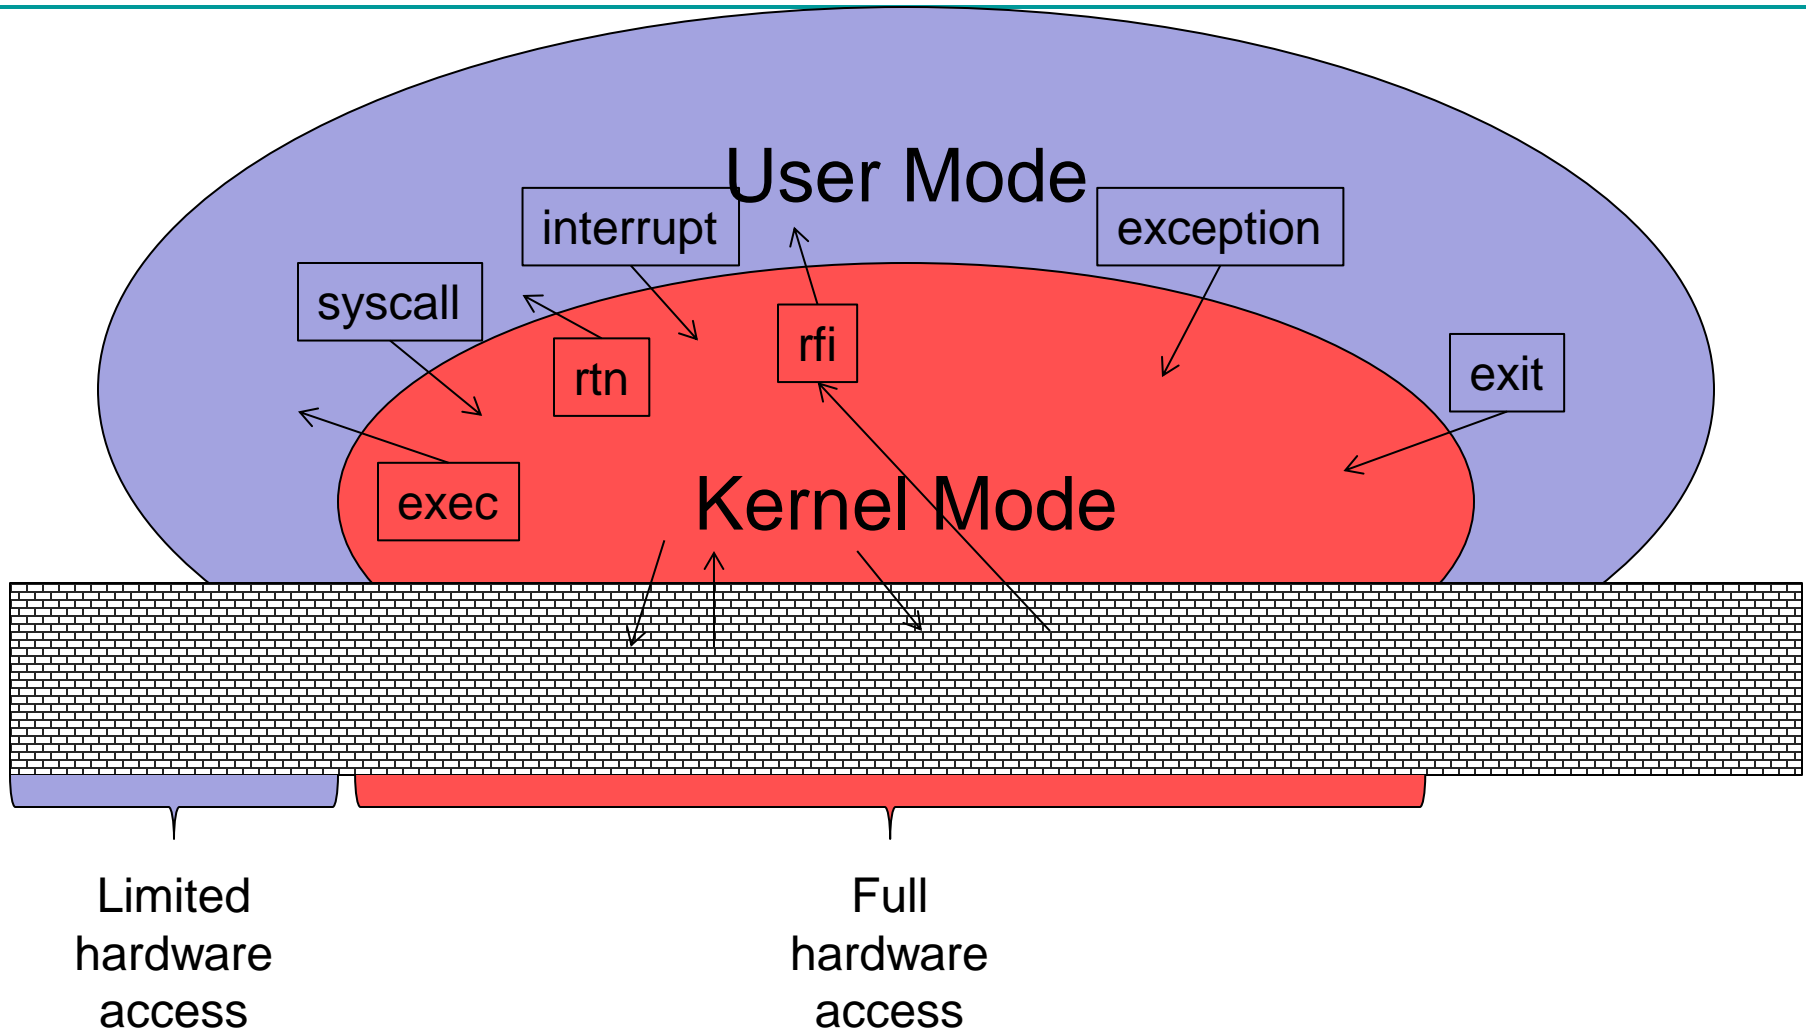


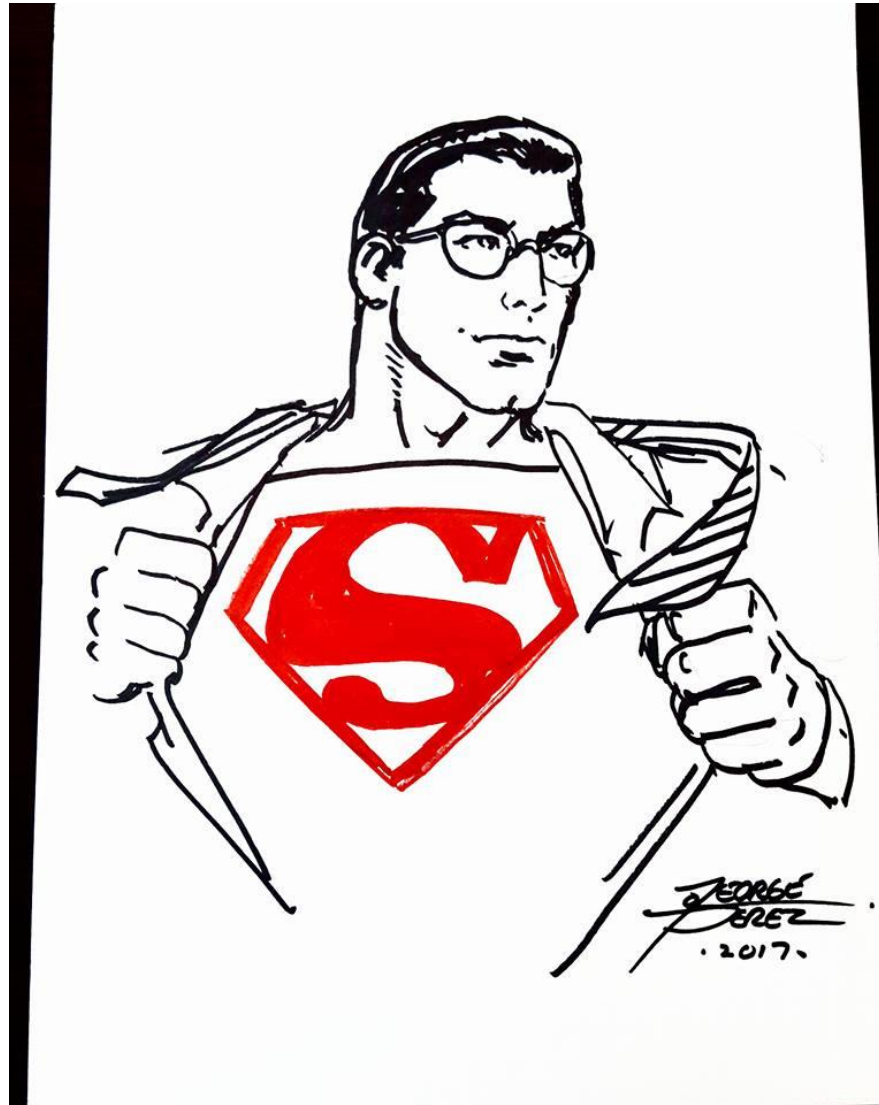


UNIX System Structure



User/Kernel (Privileged) Mode





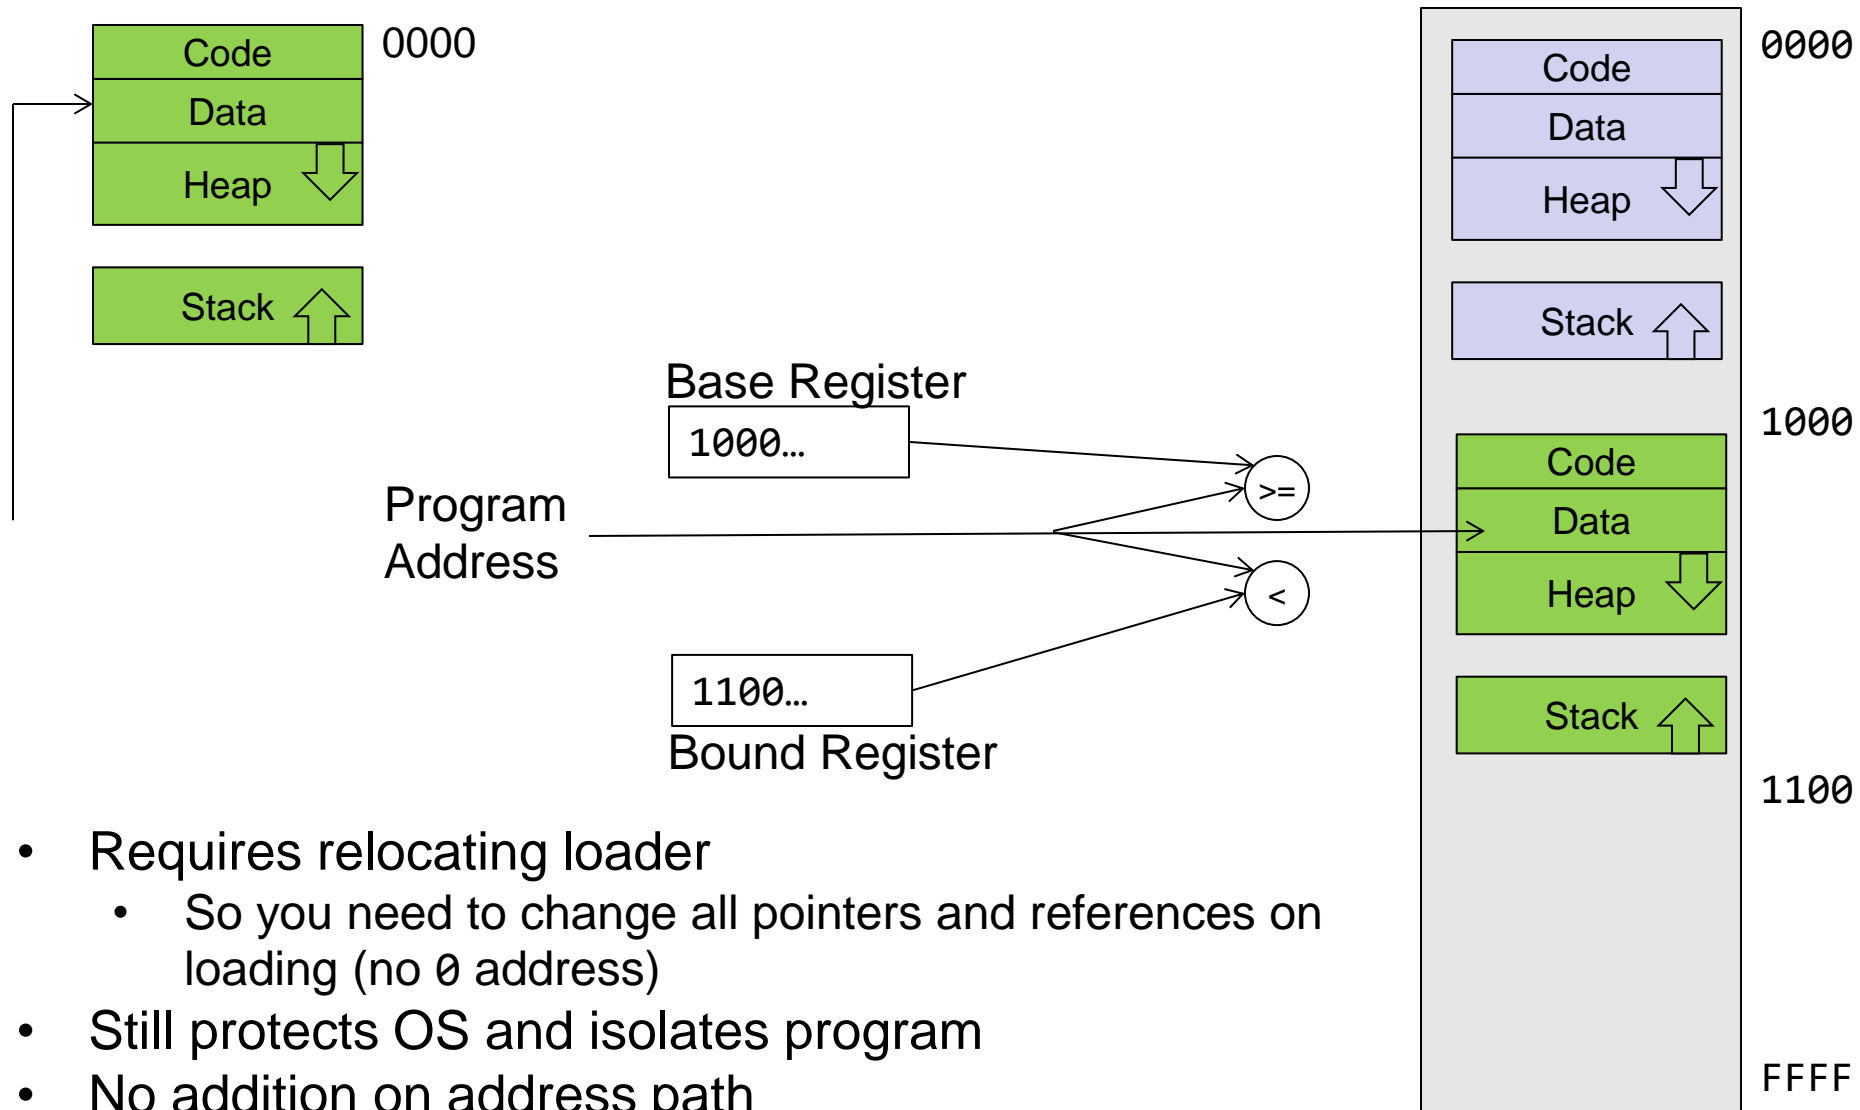
So Far

- Four Concepts
 - Threads (last week)
 - Address space (continued)
 - Processes
 - Dual Mode Operation
- Protection Mechanisms
 - Address ranges: Base and Bound
- Interrupt and Process Context

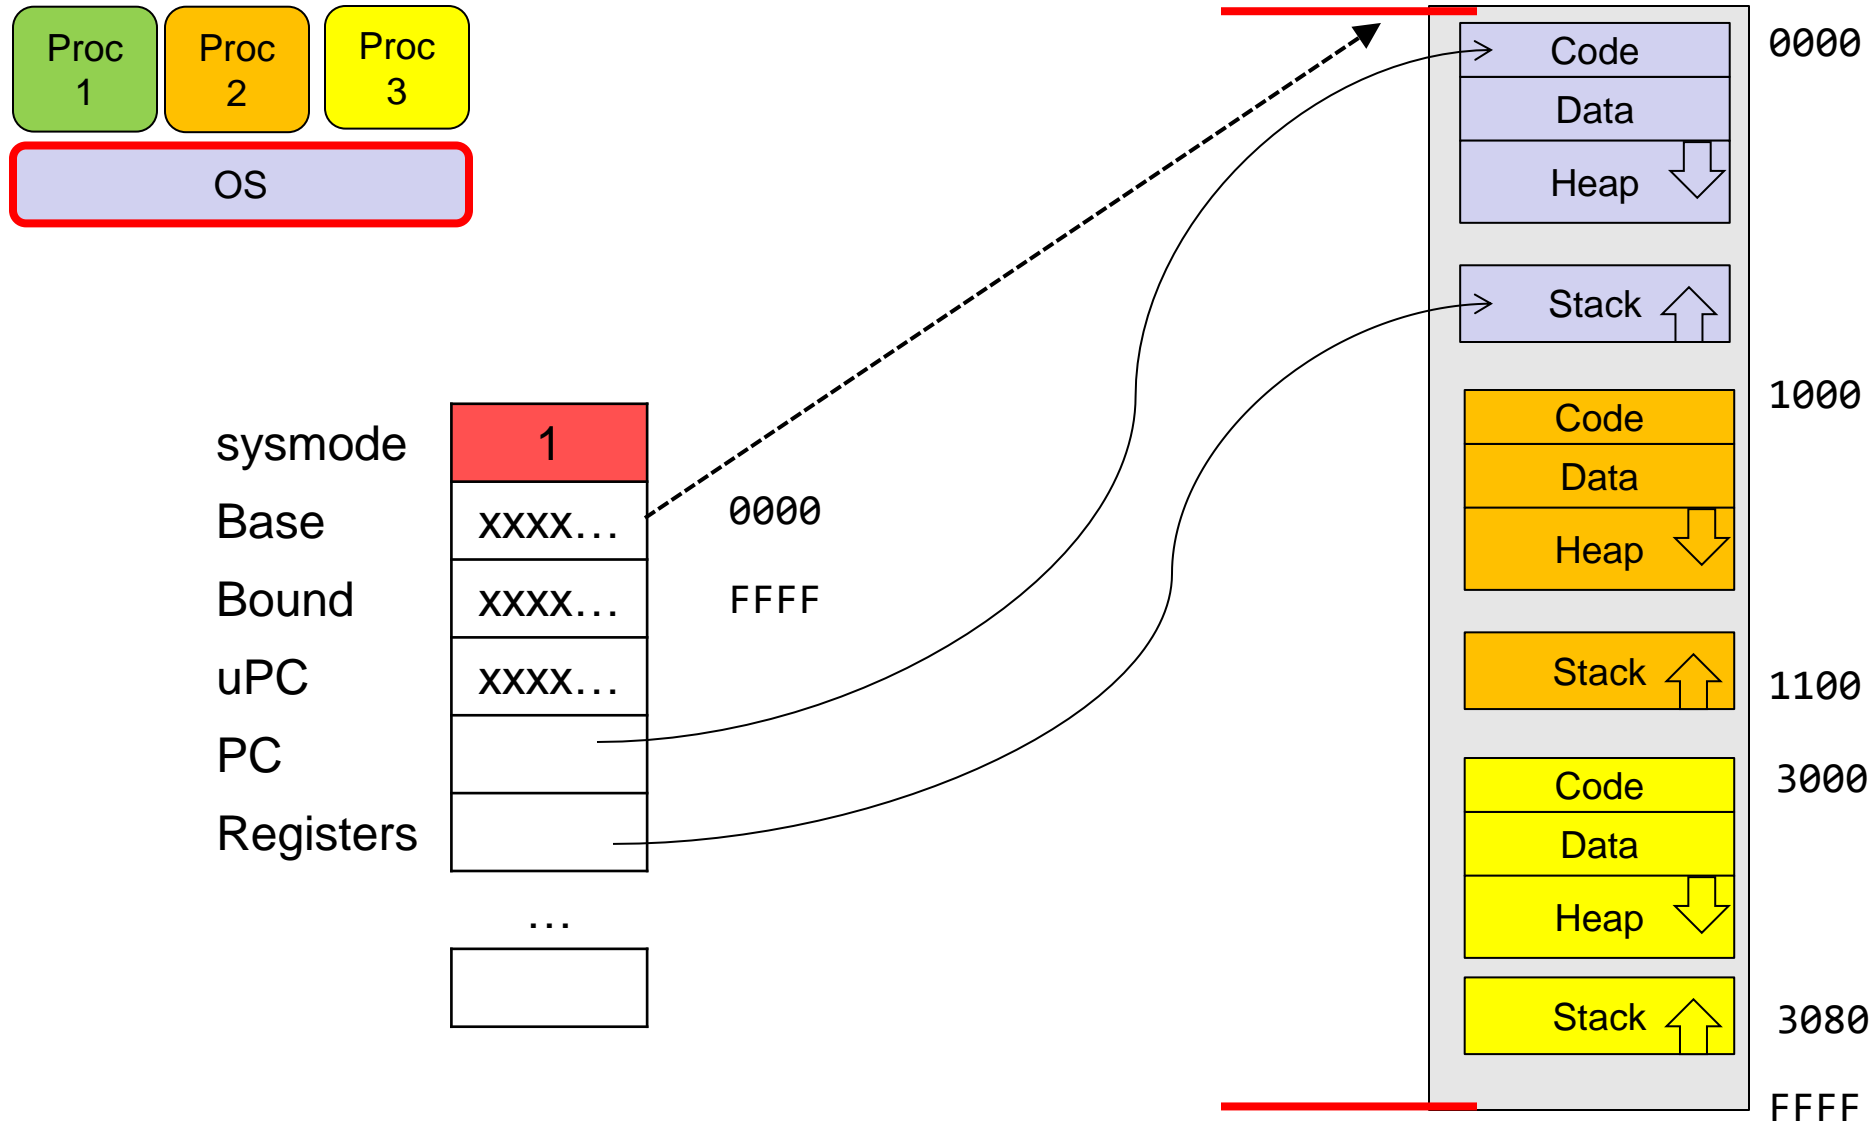
**EVERYTHING
THE LIGHT TOUCHES
IS OUR
KINGDOM.**



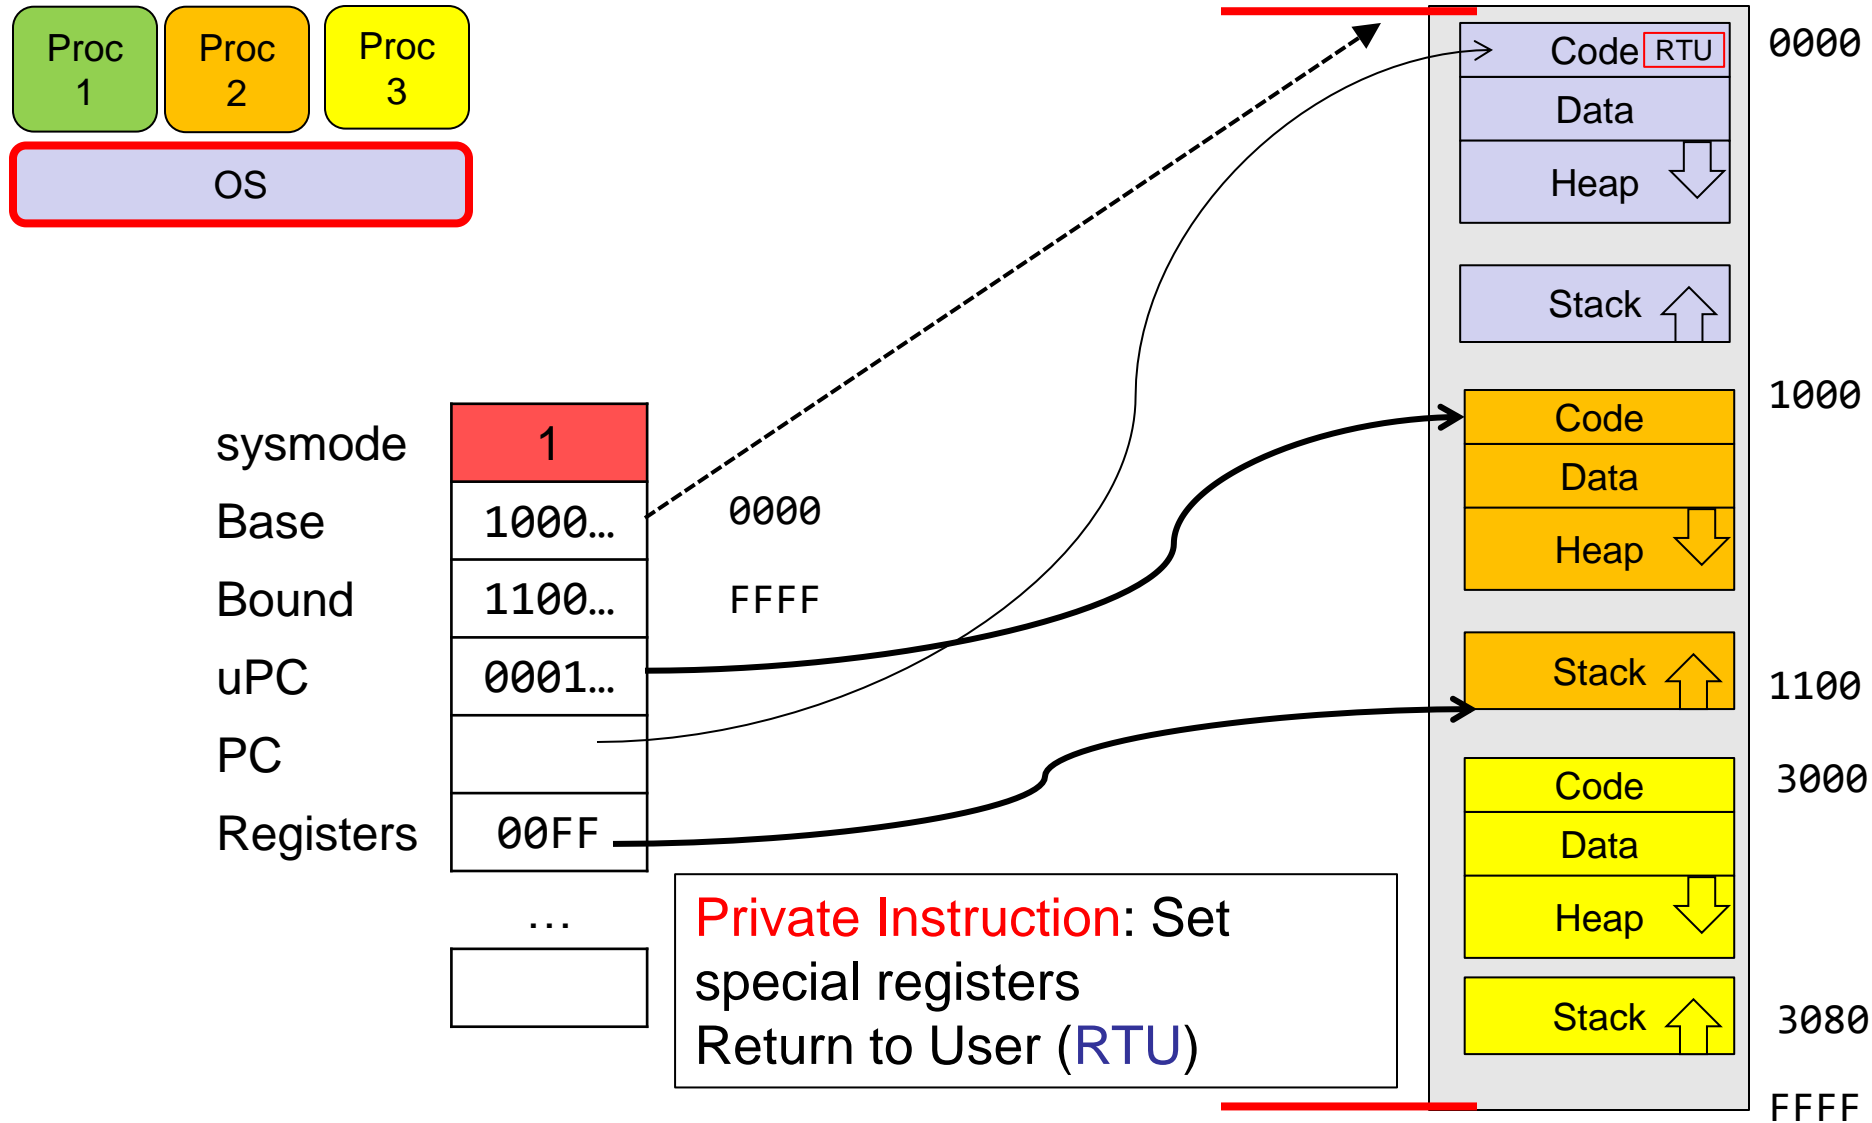
Simple Protection: Base and Bound (B&B)



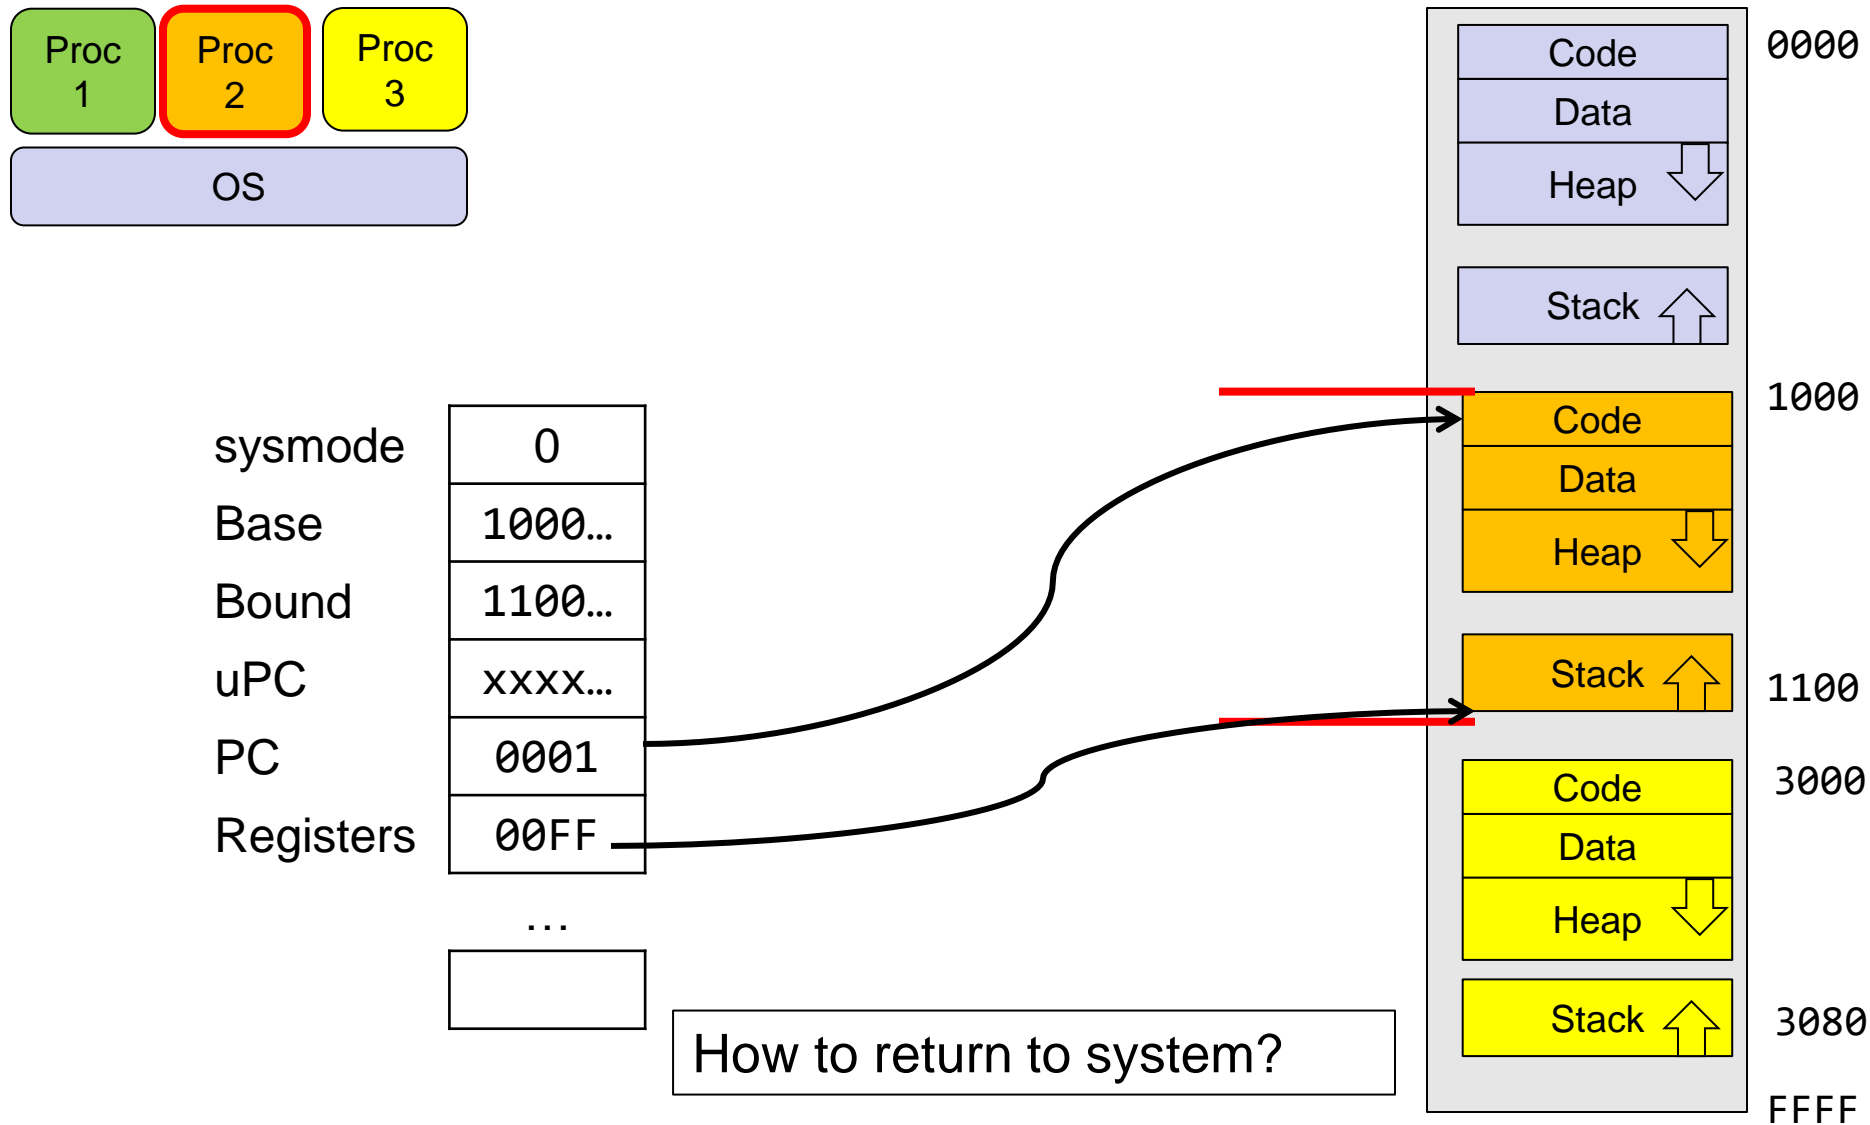
Tying it together: Simple B&B: OS Loads Process



Simple B&B: OS gets ready to switch



Simple B&B: “Return” to User



3 Types of Mode Transfer

- All 3 are a **Unprogrammed Control Transfers**
 - Where do they go?

Syscall

- Process requests a **system service** (e.g., `exit()`)
- Like a function call, but “**outside**” the process
- Does not have the address of the system function to call
- Like a **Remote Procedure Call (RPC)** (next year)
- Marshall the **syscall id** and **arguments** in registers and **exec syscall**

Interrupt

- External asynchronous event triggers context switch
- Examples: Timer, I/O device
- Independent of user process

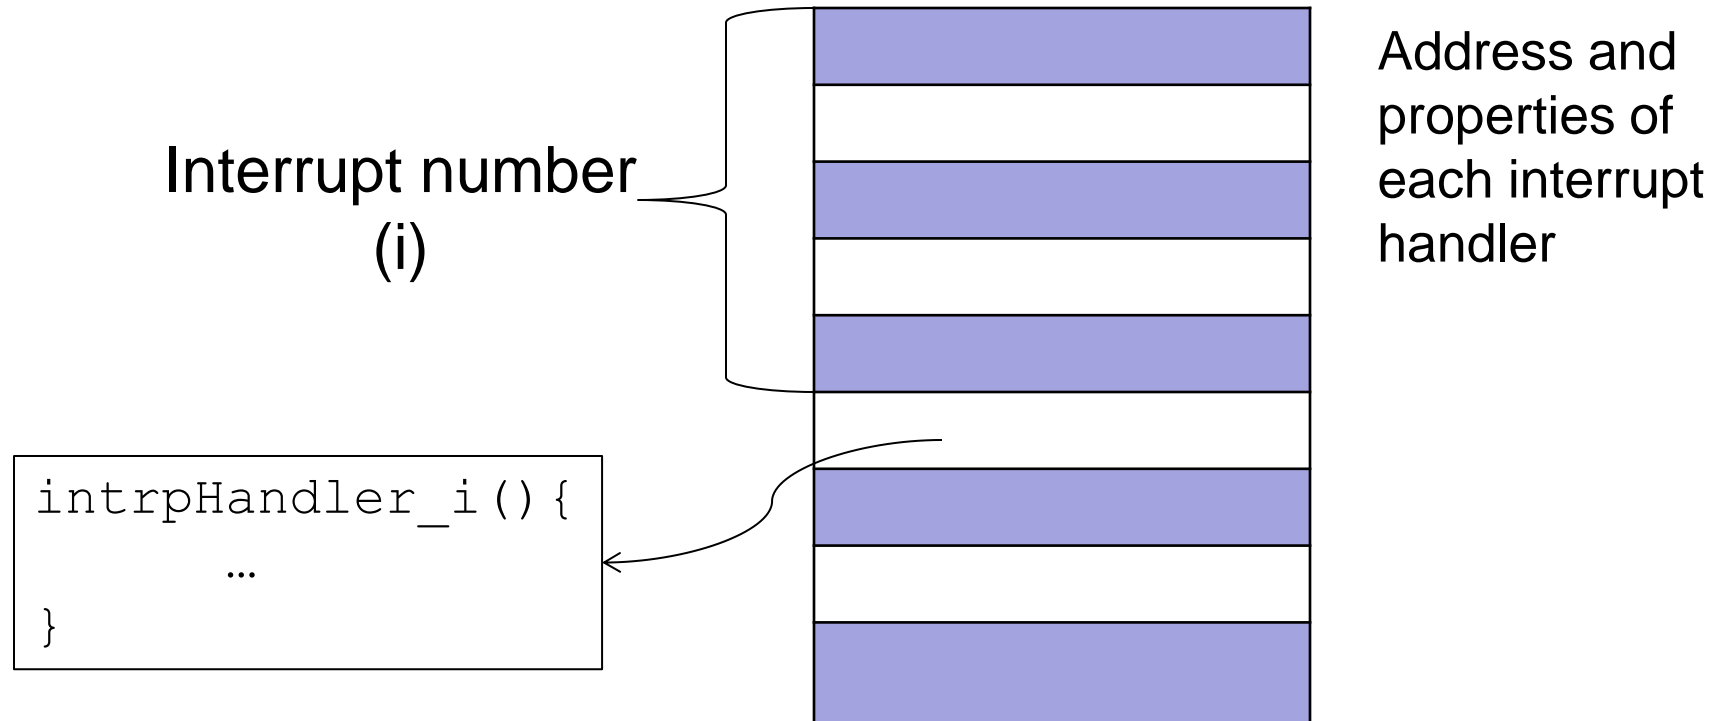
Trap or Exception

- Internal synchronous event in process triggers context switch
- e.g., Protection violation (segmentation fault), Divide by zero, etc.

How do we get the system
target address of the
“unprogrammed control
transfer?”

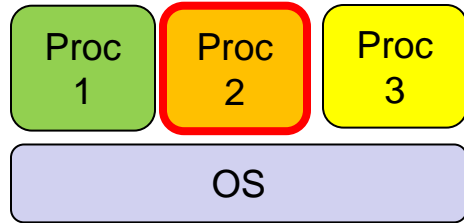


Interrupt Vector

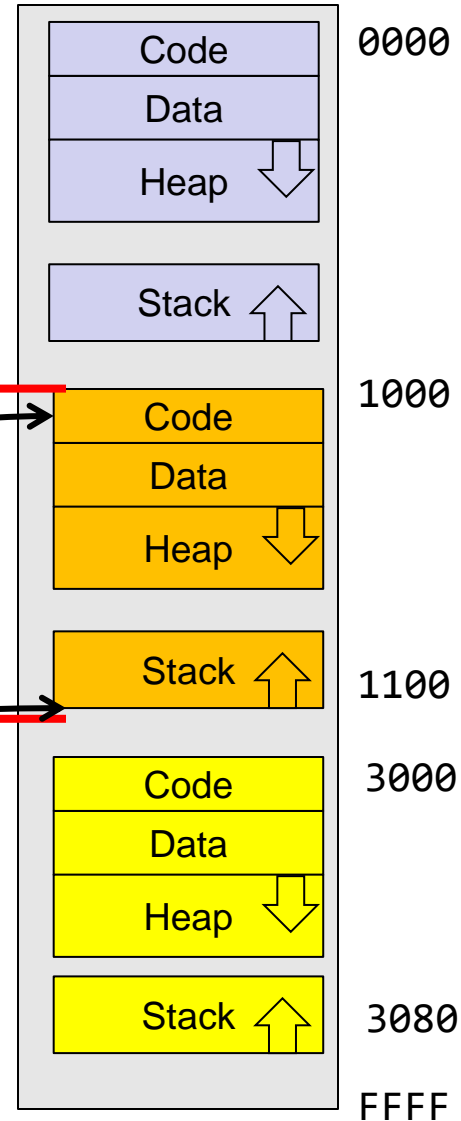


Where else do you see this dispatch pattern?

Simple B&B: User → Kernel

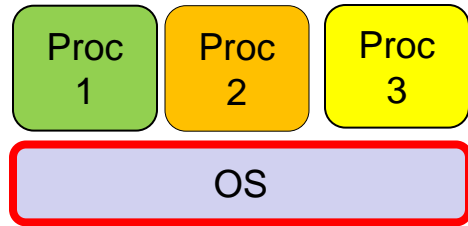


sysmode	0
Base	1000...
Bound	1100...
uPC	xxxx...
PC	000 1234
Registers	00FF
	...



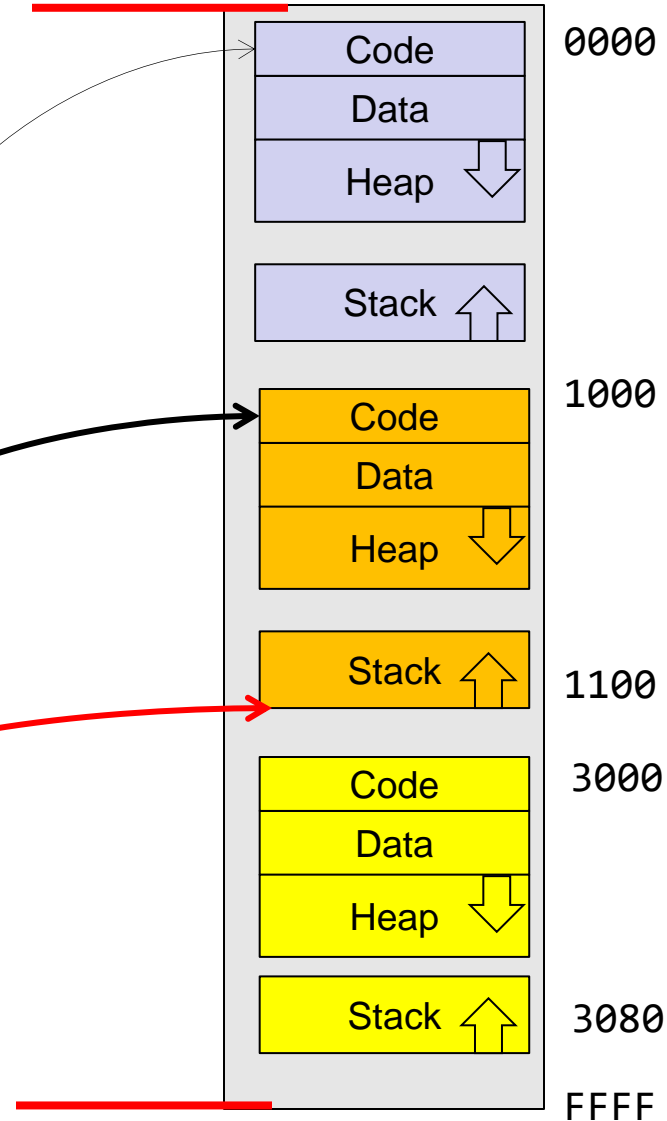
How to return to system?

Simple B&B: Interrupt

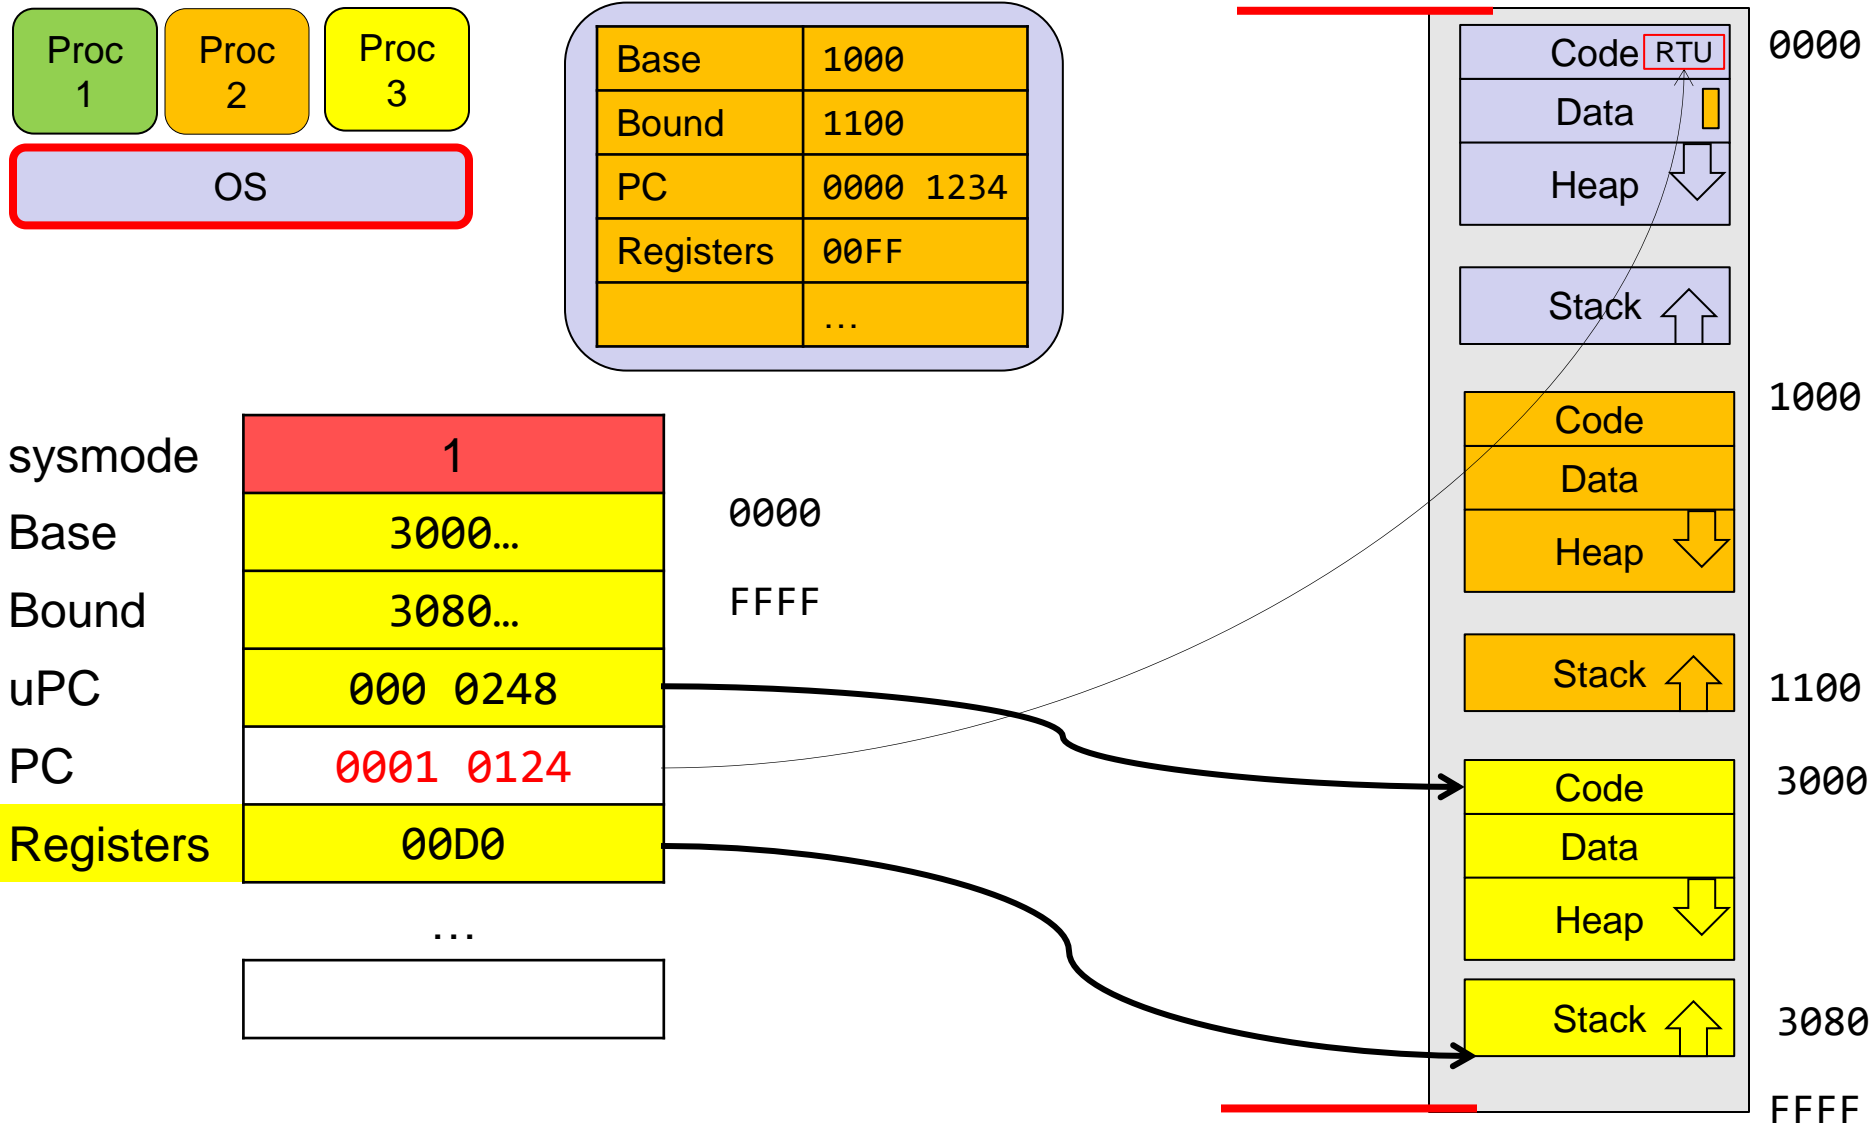


sysmode	1
Base	1000...
Bound	1100...
uPC	000 1234
PC	<i>intrpHandler_i</i>
Registers	00FF
	...

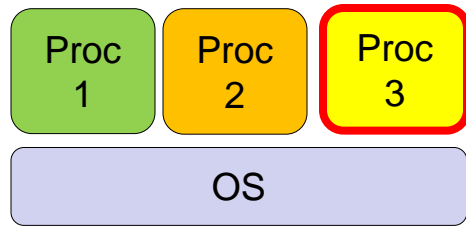
0000
FFFF



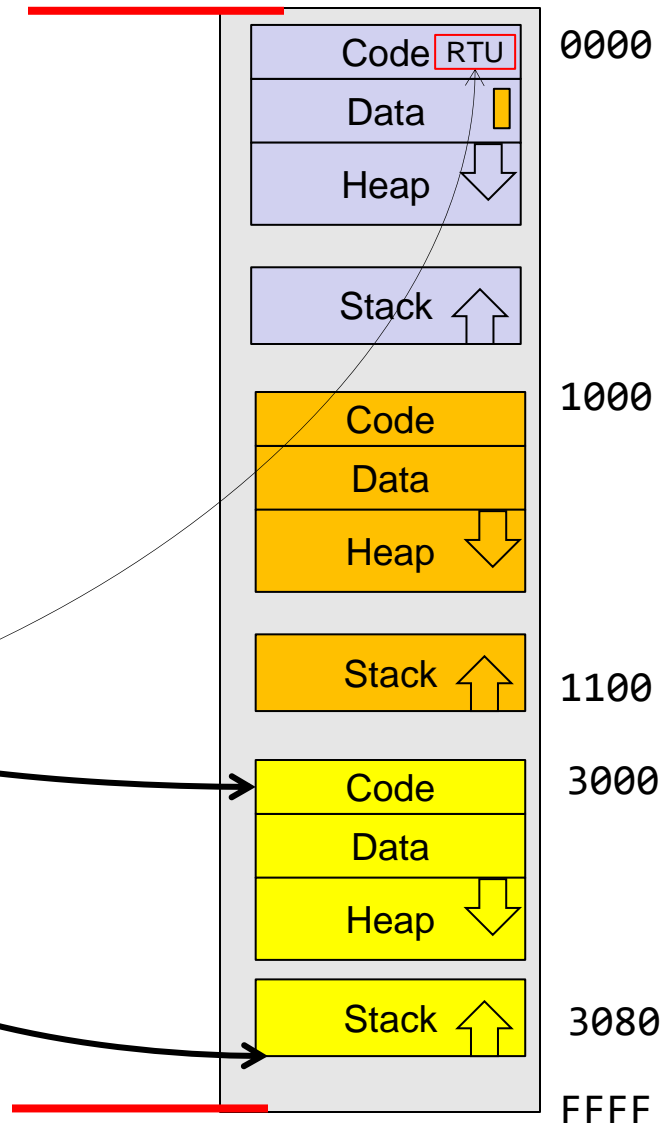
Simple B&B: Switch User Process



Simple B&B: “Resume”



Base	1000
Bound	1100
PC	0000 1234
Registers	00FF
	...

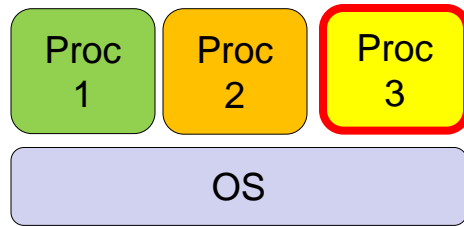


sysmode	1
Base	3000...
Bound	3080...
uPC	000 0248
PC	0001 0124
Registers	00D0
	...

0000

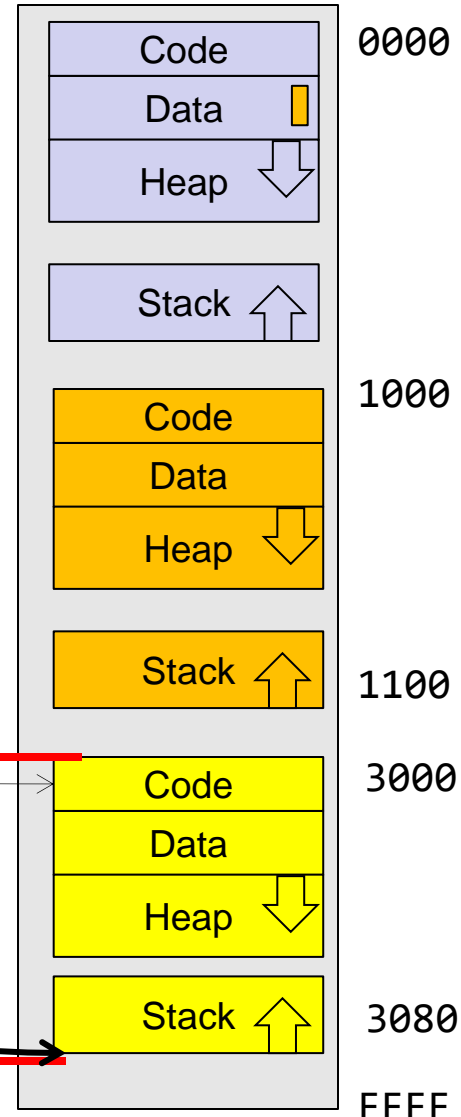
FFFF

Simple B&B: Proc 3 is running



Base	1000
Bound	1100
PC	0000 1234
Registers	00FF
	...

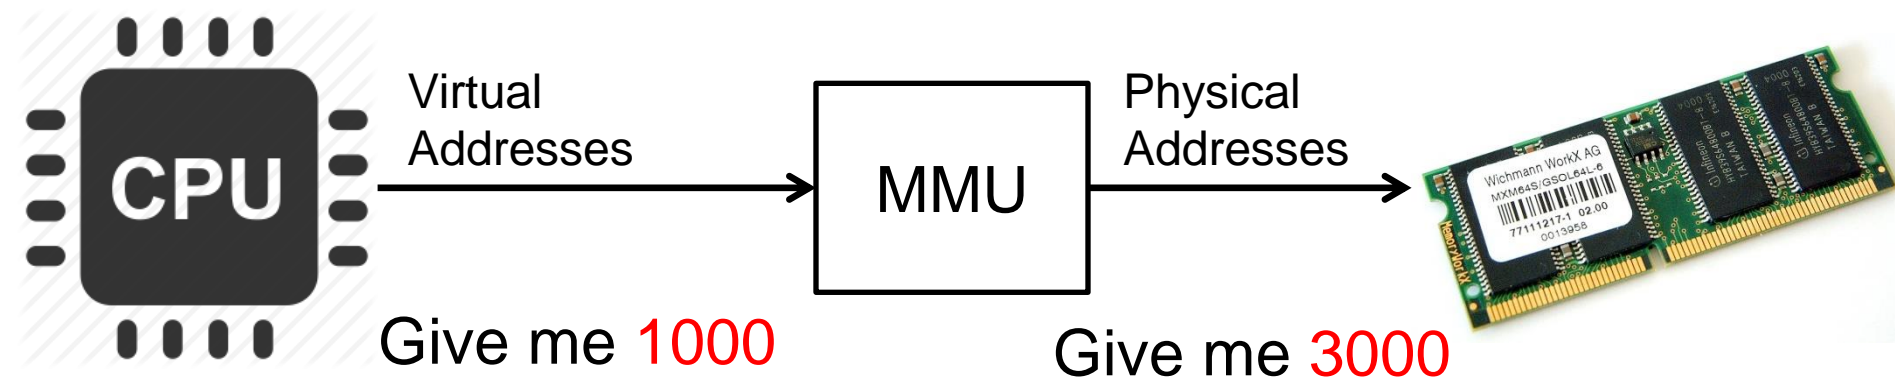
sysmode	0	
Base	3000...	0000
Bound	3080...	FFFF
uPC	xxxx	
PC	000 0248	
Registers	00D0	
	...	



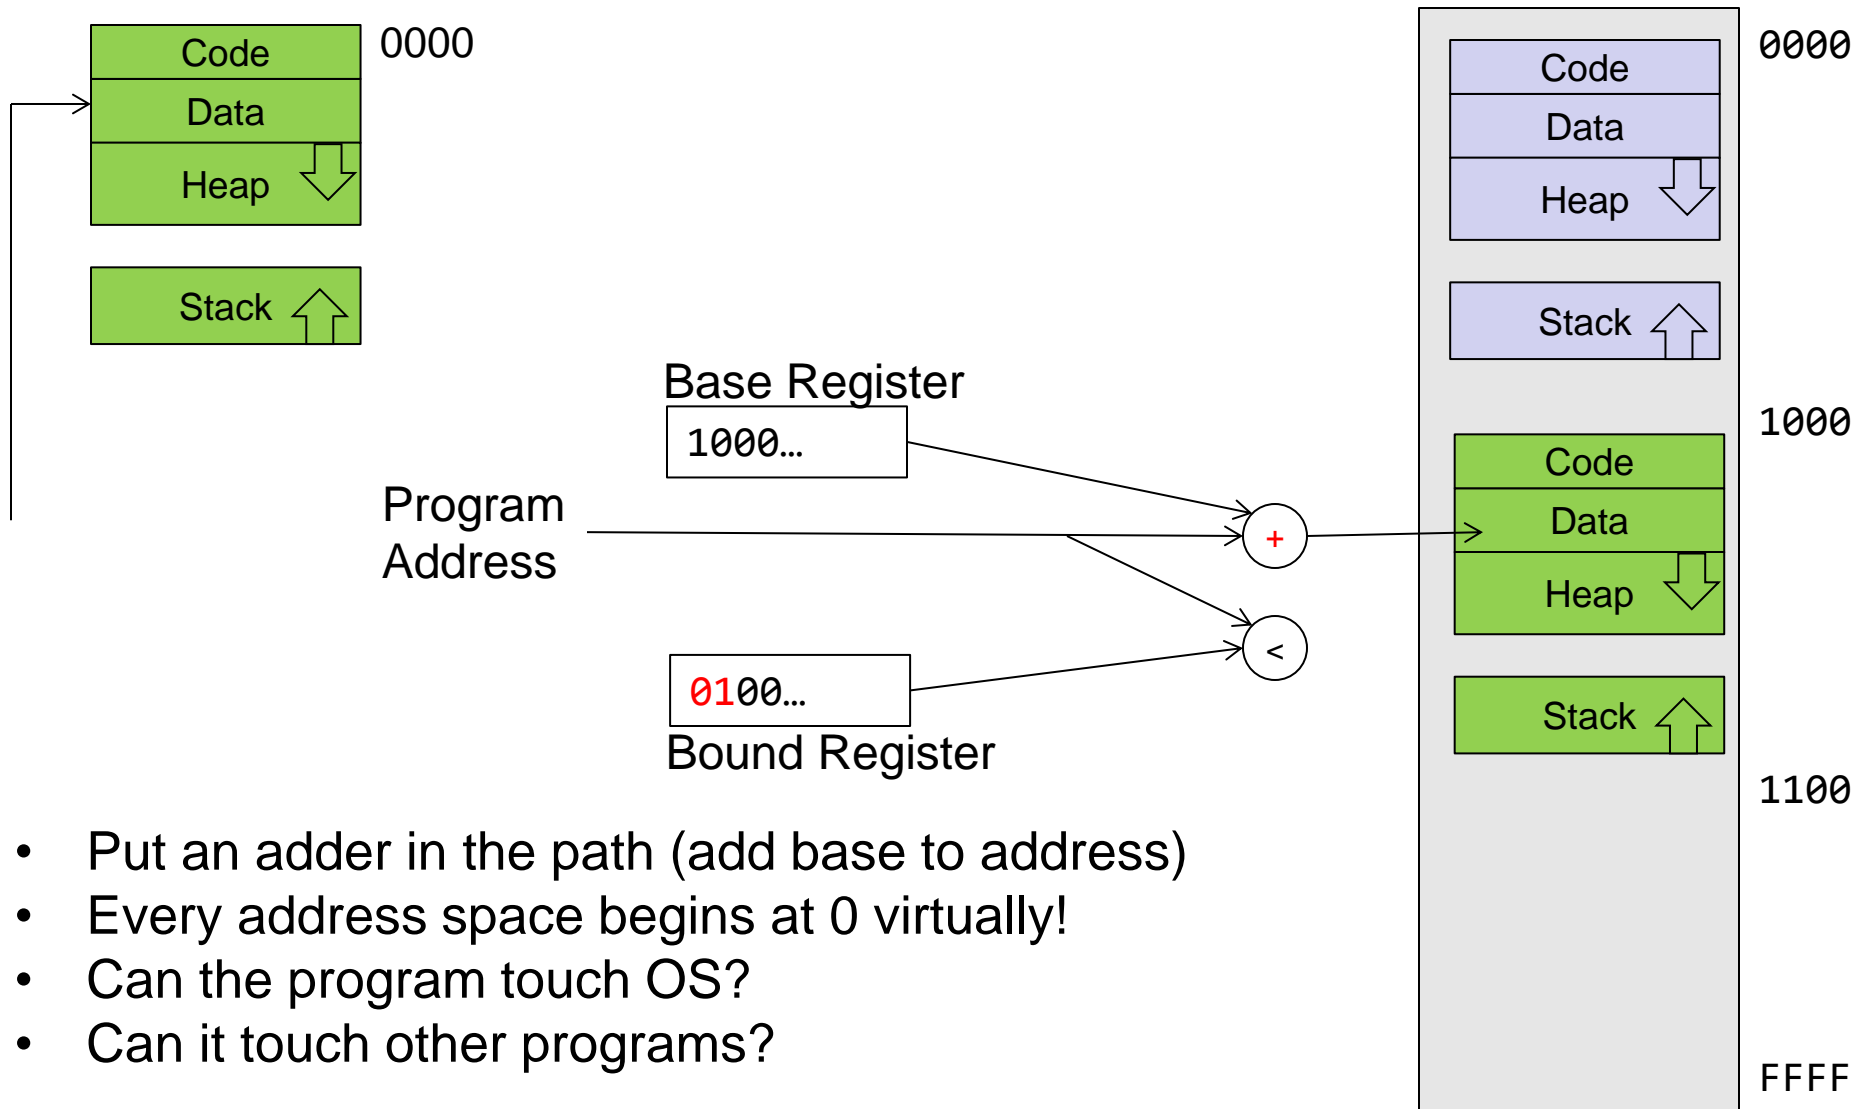
How do you save registers and set up the system stack?

Another Idea: Address Space Translator

- Program operates in an address space that is distinct from the physical memory space of the machine



Add an Adder to Base and Bound (B&B)

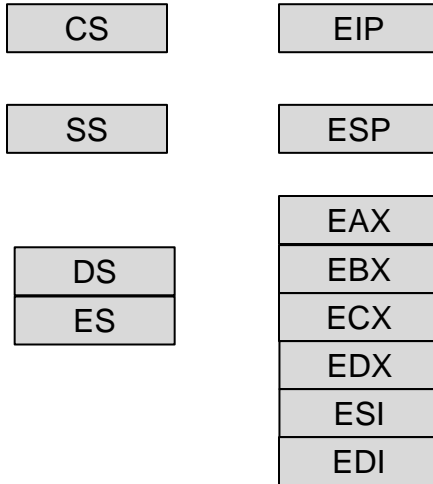


- Put an adder in the path (add base to address)
- Every address space begins at 0 virtually!
- Can the program touch OS?
- Can it touch other programs?

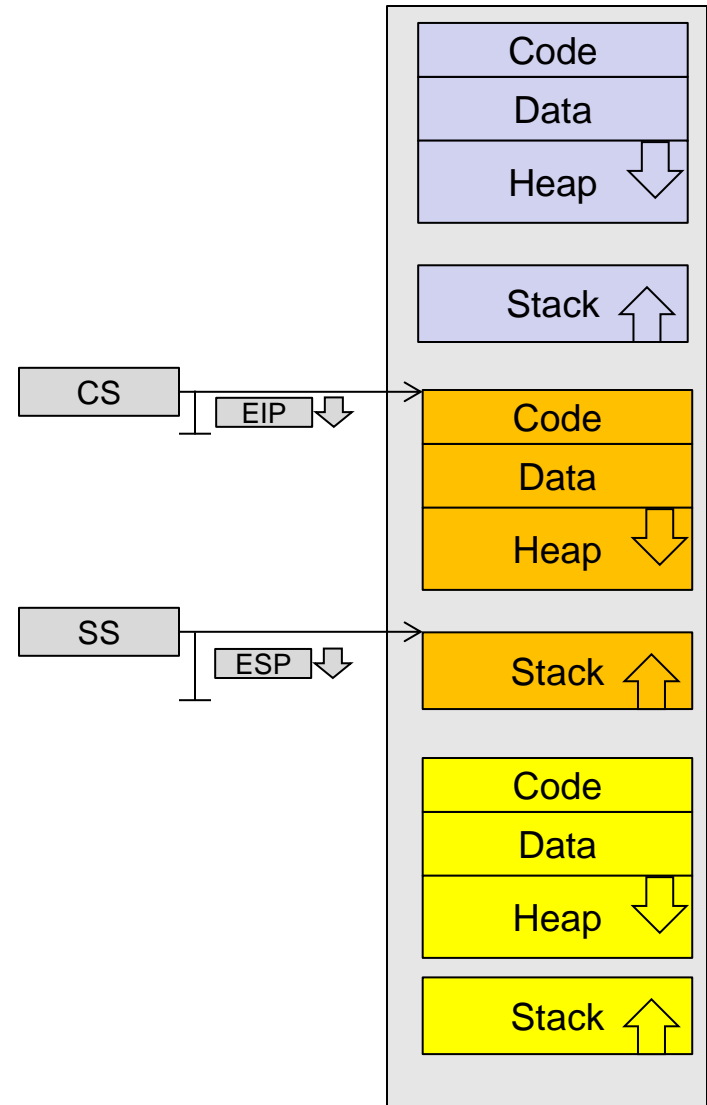
What's wrong with this simplistic address translation mechanism?

x86 – Segments and Stacks

Processor Registers



Start address, length, and access rights are associated with each segment

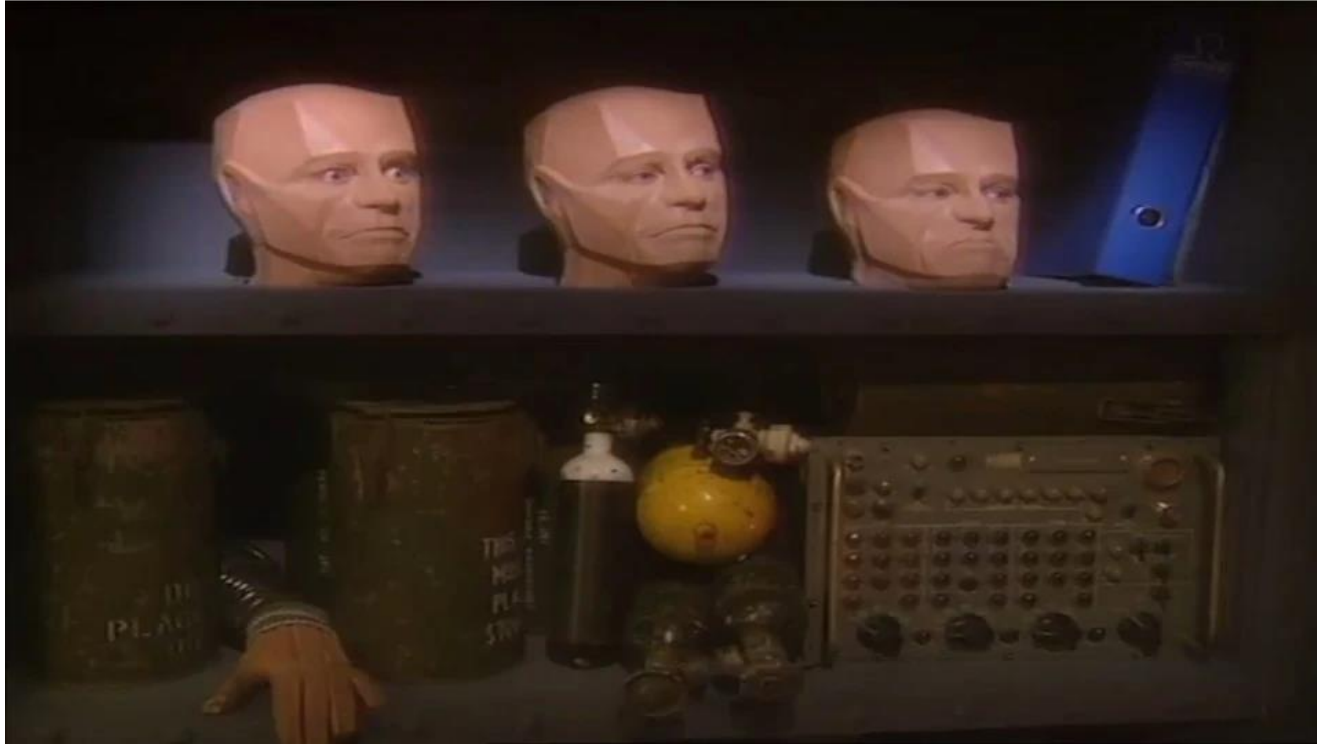


So Far

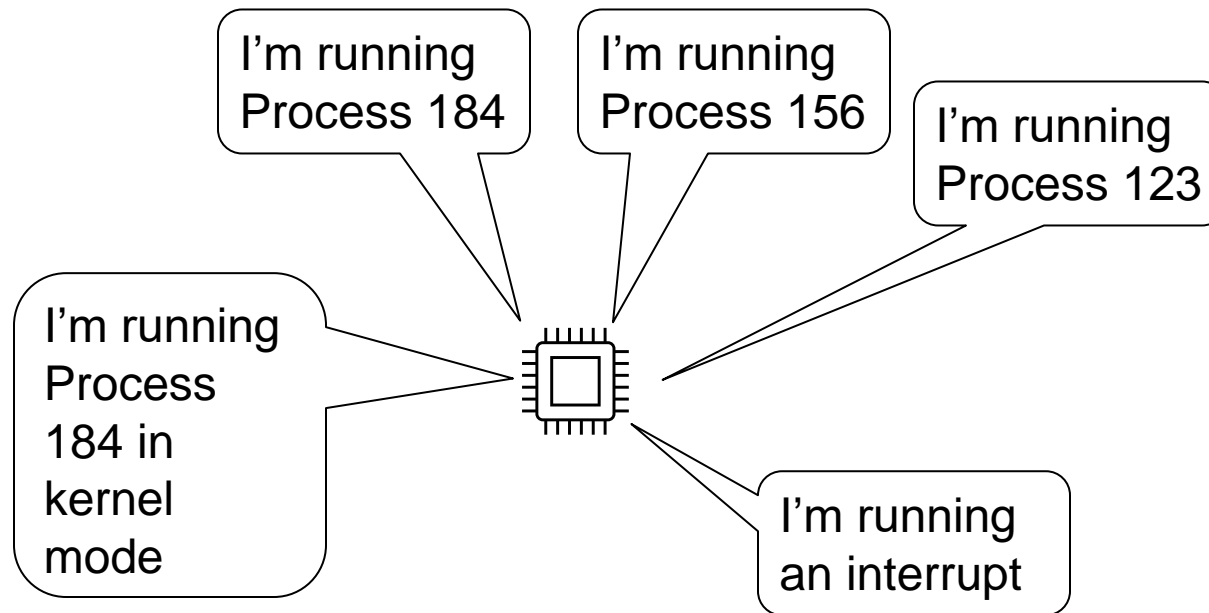
- Four Concepts
 - Threads (last week)
 - Address space (continued)
 - Processes
 - Dual Mode Operation
- Protection Mechanisms
 - Address ranges: Base and Bound
- Interrupt and Process Context

Which head to use now?

<https://vignette.wikia.nocookie.net/reddwarf/images/c/c9/Spareheads.jpg/revision/latest?cb=20140602034628>



Discern: Context



Process context versus Interrupt context

Conclusion

- Four Concepts
 - Threads (last week)
 - Address space (continued)
 - Processes
 - Dual Mode Operation
- Protection Mechanisms
 - Address ranges: Base and Bound
- Interrupt and Process Context