

Signals, Syscalls, I/O Basics, High and low, Drivers, Sockets

5 December 2024
Lecture 5

Slides adapted from John Kubiatowicz (UC Berkeley)

Concept Review

Segments

Process
context

Interrupt
context

Interrupt
handler

Kernel Stack

Process
Control Block

Scheduler

Simultaneous Multi
Threading
(SMT)

fork


pid_t



Topics for Today

- Signals
- Syscalls
- Basic Support for I/O (drivers, etc.)
 - Files and Streams
 - Low Level
- I/O and Drivers
- Sockets and networks

Signals: Simple Messaging



- **Signal**: Like a software interrupt
 - Simple (integer) message you can send a process
 - Interrupts normal execution
- Sent by:
 - One process to another (syscall or via `kill` program)
 - OS to a process (due to event)

- Sample signals:

SIGINT	Interrupt from keyboard (CTRL+C)
SIGQUIT	Quit from keyboard (CTRL+\)
SIGCONT	Continue a program
SIGKILL	Kill signal
SIGTTIN	Background process tries to read from STDIN

Signals – infloop.c

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
#include <signal.h>

void signal_callback_handler(int signum)
{
    printf("Caught signal %d - phew!\n",signum);
    exit(1);
}

int main() {
    signal(SIGINT, signal_callback_handler);
    while (1) {}
}
```

How do you stop this?

Look at top

Some more signals – signals-sample.c

```
void signal_handler(int sig) {
    switch (sig) {
        case SIGINT:

            printf("Received SIGINT (Interrupt from keyboard), signal number: %d\n", sig);
            break;

        case SIGTERM:

            printf("Received SIGTERM (Termination signal), signal number: %d\n", sig);
            break;

        case SIGQUIT:

            printf("Received SIGQUIT (Quit from keyboard), signal number: %d\n", sig);
            break;

        default:

            printf("Received signal number: %d\n", sig);

    }
}
```



signals-sample.c

```
int main() {
    if (signal(SIGINT, signal_handler) == SIG_ERR) {
        perror("Error registering SIGINT handler");
        exit(1);
    }

    if (signal(SIGTERM, signal_handler) == SIG_ERR) {
        perror("Error registering SIGTERM handler");
        exit(1);
    }

    if (signal(SIGQUIT, signal_handler) == SIG_ERR) {
        perror("Error registering SIGQUIT handler");
        exit(1);
    }

    printf("Running... Press Ctrl+C to send SIGINT, or send SIGTERM or SIGQUIT to this process.\n");
    while (1) {
        sleep(1);
    }

    return 0;
}
```

An aside about System Calls (syscalls)

What is a syscall?

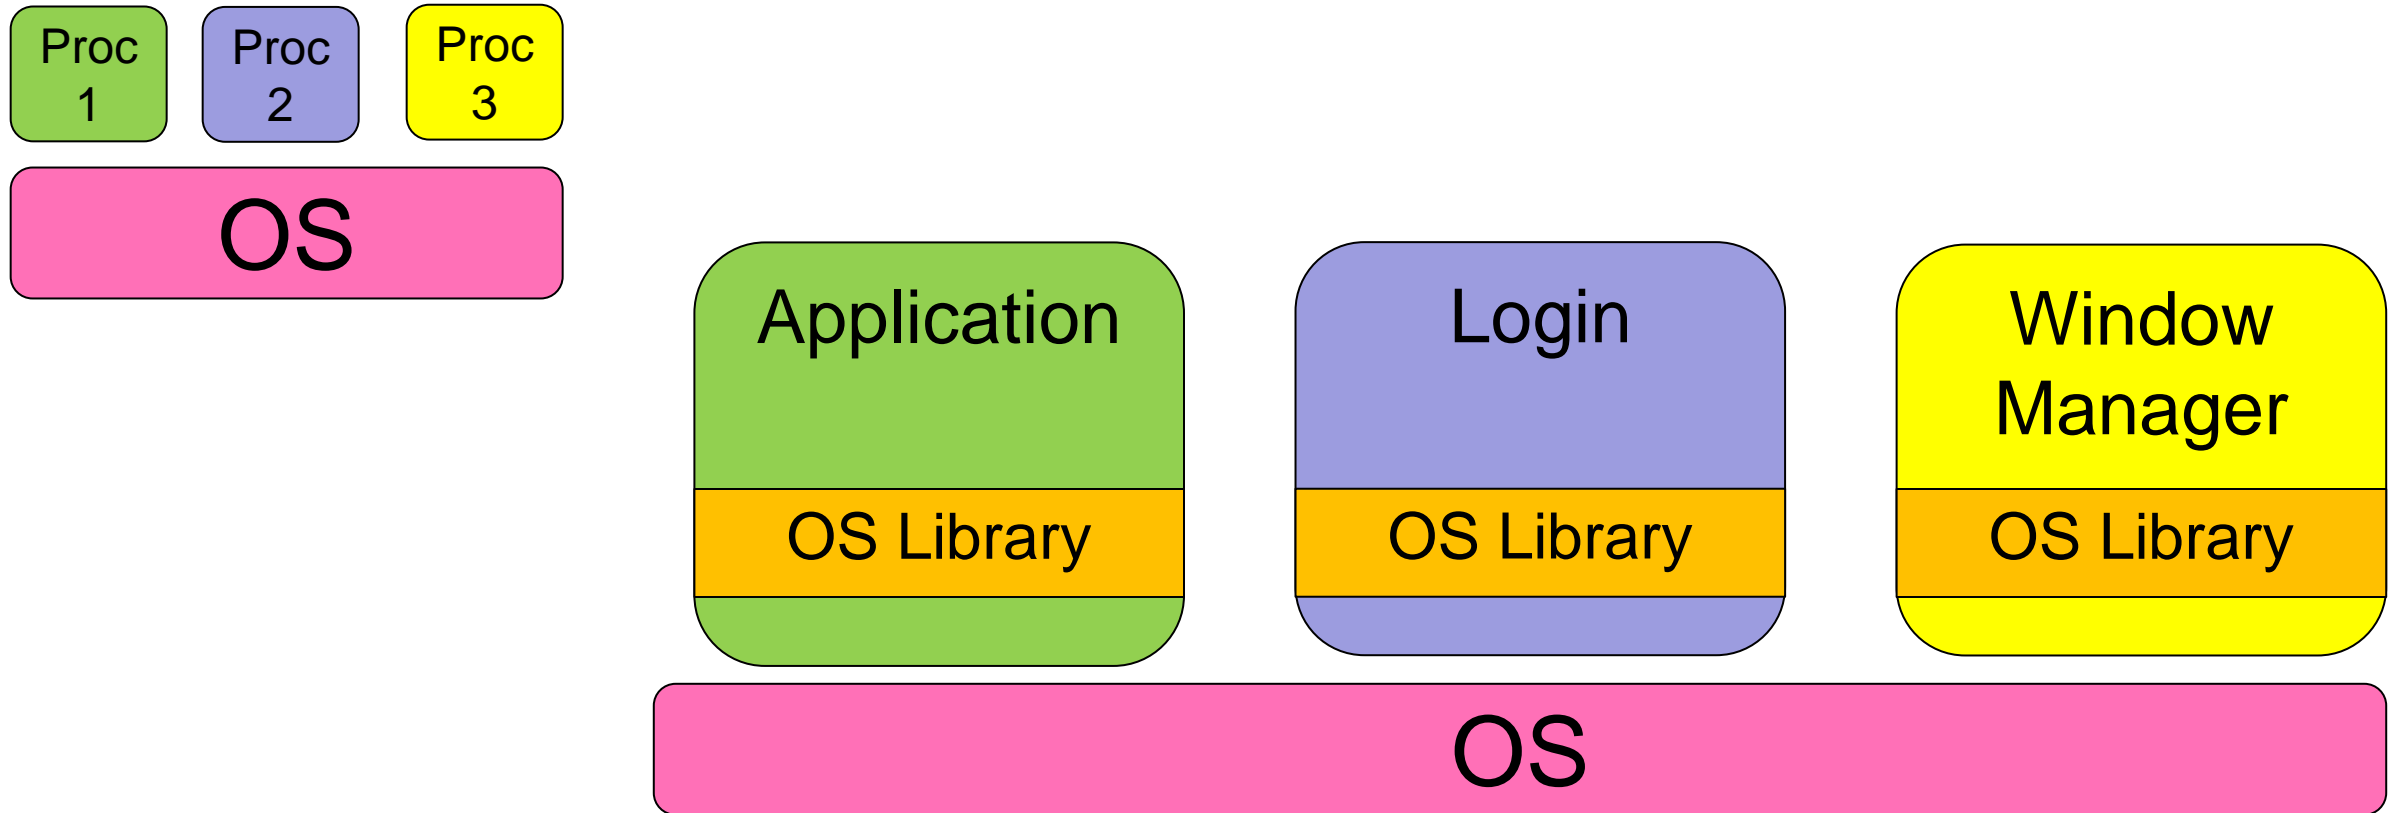
- Applications request services from the operating system via a **syscall**, but ...
- I've been writing applications and never saw a “syscall” in code?

Why?

Syscalls are buried
in the programming
language runtime
library (ex. `libc.a`)

Layering at work

OS run-time library

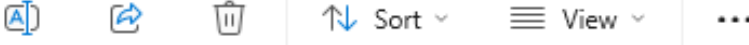

















In WSL (Ubuntu)

```
mjmay@RAMOT:/lib/x86_64-linux-gnu$ ls
Mcrt1.o      libcurl-gnutls.so.4.7.0  libjack.so.0.1.0      libpython3.10.so.1.0
Scrt1.o      libcurl.so.4             libjacknet.so.0        libquadmath.so.0
audit        libcurl.so.4.7.0         libjacknet.so.0.1.0    libquadmath.so.0.0.0
avahi        libcurses.a              libjackserver.so.0     libraw1394.so.11
awk          libcurses.so             libjackserver.so.0.1.0 libraw1394.so.11.1.0
bfd-plugins libdatatrie.so.1         libjansson.so.4        libreadline.so.8
caca         libdatatrie.so.1.4.0     libjansson.so.4.13.0   libreadline.so.8.1
crt1.o       libdb-5.3.so             libjavascriptcoregtk-4.0.so.18 libresolv.a
crti.o       libdbus-1.so.3           libjavascriptcoregtk-4.0.so.18.25.12 libresolv.so
crtn.o       libdbus-1.so.3.19.13    libjbig.so.0           libresolv.so.2
dri          libdconf.so.1            libjpeg.so.8           librom1394.so.0
e2fsprogs   libdconf.so.1.0.0       libjpeg.so.8.2.2       librom1394.so.0.3.0
enchant-2   libdebconfclient.so.0    libjson-c.so.5         librsvg-2.so.2
engines-3   libdebconfclient.so.0.0.0 libjson-c.so.5.1.0     librsvg-2.so.2.48.0
gawk        libdeflate.so.0         libjson-glib-1.0.so.0  librt.a
gconv       libdevmapper.so.1.02.1  libjson-glib-1.0.so.0.600.6 librt.so.1
gcrt1.o     libdl.a                 libk5crypto.so.3       librtmp.so.1
gdk-pixbuf-2.0 libdl.so.2              libk5crypto.so.3.1     libsamplerate.so.0
gedit       libdns-9.18.28-0ubuntu0.22.04.1-Ubuntu.so libkeyutils.so.1       libsamplerate.so.0.2.2
gio         libdns-export.so.1110   libkeyutils.so.1.9     libsasl2.so.2
girepository-1.0 libdns-export.so.1110.0.2 libkmod.so.2           libsasl2.so.2.0.25
glib-2.0    libdrm.so.2             libkmod.so.2.3.7       libseccomp.so.2
graphviz    libdrm.so.2.4.0         libkrb5.so.3           libseccomp.so.2.5.3
gstreamer1.0 libdrm_and_gnu.so.1     libkrb5.so.3.2         libsecret-1.so.0
```

In Windows 11

> This PC > Windows (C:) > Windows > System32 >



Name	Date modified	Type	Size
 wpnsruprov.dll	22/11/2024 9:27	Application extension	140 KB
 WpnUserService.dll	22/11/2024 9:27	Application extension	148 KB
 WpPortingLibrary.dll	07/05/2022 8:19	Application extension	36 KB
 WppRecorderUM.dll	22/11/2024 9:27	Application extension	64 KB
 WPTaskScheduler.dll	22/11/2024 9:27	Application extension	224 KB
 wpx.dll	13/11/2024 14:27	Application extension	1,462 KB
 ws2_32.dll	07/05/2022 8:19	Application extension	466 KB
 ws2help.dll	07/05/2022 8:19	Application extension	12 KB
 wsauth.dll	18/07/2024 1:49	Application extension	1,677 KB
 wscapi.dll	22/11/2024 9:27	Application extension	280 KB
 wscinterop.dll	07/05/2022 8:20	Application extension	256 KB
 wscisvif.dll	05/11/2023 23:14	Application extension	52 KB
 WSCClient.dll	07/05/2022 8:19	Application extension	40 KB
 wscproxystub.dll	05/11/2023 23:14	Application extension	36 KB
 wscntutil.dll	22/11/2024 9:27	Application extension	302 KB

A Narrow Waist

Compilers Word Processing Web Browsers

Application/
Service

Email Databases Web Servers

Portable OS Library

System Call Interface

Portable OS Kernel

Platform Support

Device Drivers

x86 PowerPC ARM

Ethernet (10/100/1000) 802.11 a/b/g/n SCSI IDE Graphics

User

System

Software

Hardware

OS



© Reuters

<https://www.independent.co.uk/news/world/asia/hundreds-of-cars-caught-up-in-beijing-traffic-during-golden-week-a6685961.html>

So Far

- Signals
- Syscalls
- Basic Support for I/O (drivers, etc.)
 - Files and Streams
 - Low Level
- I/O and Drivers
- Sockets and networks

Key Unix I/O Design Concepts

Uniformity



- file operations, device I/O, and interprocess communication through open, read/write, close
- Allows simple composition of programs
 - `find | grep | wc ...`

Open before use



- Provides opportunity for access control and arbitration
- Sets up the underlying machinery, i.e., data structures

Byte-oriented

- Even if blocks are transferred, **addressing** is in bytes



Key Unix I/O Design Concepts

Kernel buffered reads

- Streaming and block devices looks the same
- read blocks process, yielding processor to other task



Kernel buffered writes

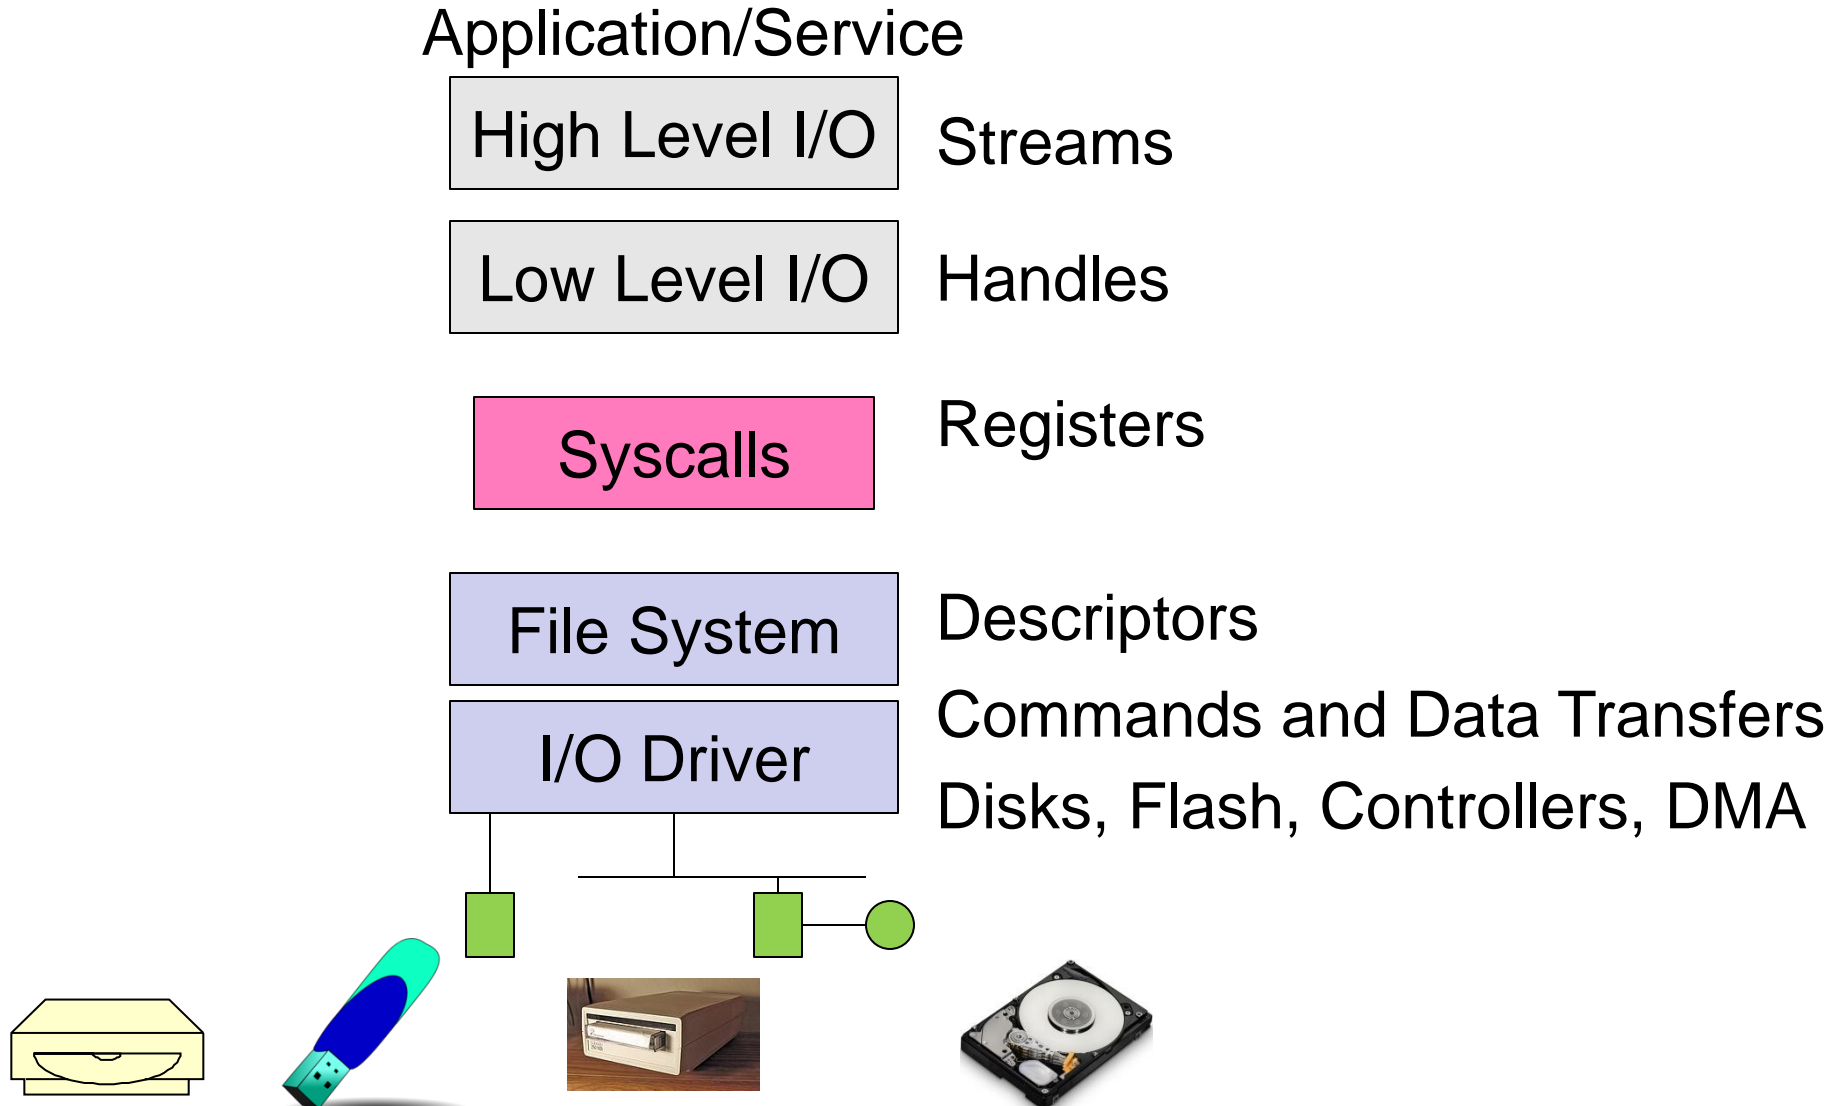
- Completion of out-going transfer decoupled from the application, allowing it to continue



Explicit close



I/O & Storage Layers



The file system abstraction

- File



Named collection of data in a file system

File data

- Text, binary, linearized objects

File Metadata:

- Size, Modification Time, Owner, Security info
- Basis for access control

- Directory



“Folder”
containing
files &
Directories


Hierarchical (graphical) naming

- Path through the directory graph
- Uniquely identifies a file or directory
 - `/home/mjmay/se317/public_html/fa24/index.html`

Links and
Volumes
(later)

C high level File API – streams

- Operate on “streams” - sequence of bytes, whether text or data, with a position



```
#include <stdio.h>
```

```
FILE *fopen( const char *filename, const char *mode );
```

```
int fclose( FILE *fp );
```

Mode Text	Binary	Description
r	rb	Open existing file for reading
w	wb	Open for writing; create if doesn't exist
a	ab	Open for appending; create if doesn't exist
r+	rb+	Open existing file for reading and writing
w+	wb+	Open for reading and writing; truncated to zero if exists, create otherwise
a+	ab+	Open for reading and writing; create if doesn't exist. Read from beginning, write as append

Connecting Processes, Filesystem, and Users

Every process has a “current working directory”

Absolute Paths

- `/home/mjmay/se317`

Relative paths

- `index.html`, `./index.html` - **current WD**
- `../index.html` - **parent of current WD**

Path aliases

- `~`, `~se317` - **home directory**

You are here!



C API Standard Streams

- Three predefined streams are opened implicitly when the program is executed.
 1. `FILE *stdin` – normal source of input, can be redirected
 2. `FILE *stdout` – normal source of output, can be redirected too
 3. `FILE *stderr` – diagnostics and errors
- `STDIN/STDOUT` enable composition in Unix
 - Recall: Use of pipe symbols connects `STDOUT` and `STDIN`
 - `find | grep | wc ...`

Picturing it






Image sources: <https://openclipart.org/detail/300835/baby-2>
<https://creazilla.com/media/clipart/34014/baby-is-crying>

C high level File API – stream ops

```
#include <stdio.h>
// character oriented
int fputc( int c, FILE *fp );    // rtn c or EOF on err
int fputs( const char *s, FILE *fp );    // rtn >0 or EOF

int fgetc( FILE * fp );
char *fgets( char *buf, int n, FILE *fp );
```

```
// block oriented
size_t fread(void *ptr, size_t size_of_elements,
             size_t number_of_elements, FILE *a_file);

size_t fwrite(const void *ptr, size_t size_of_elements,
             size_t number_of_elements, FILE *a_file);
```

```
// formatted
int fprintf(FILE *restrict stream, const char *restrict format, ...);
int fscanf(FILE *restrict stream, const char *restrict format, ... );
```


Example Stream Code

```
#include <stdio.h>
#include <string.h>

#define BUFLLEN 256
FILE *outfile;
char mybuf[BUFLLEN];

int storetofile(){
    char *instring;

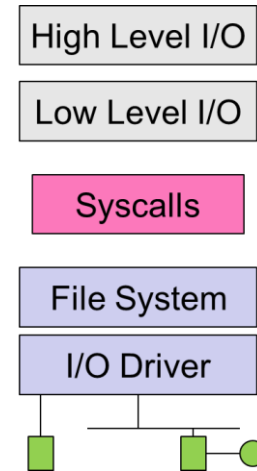
    outfile = fopen("/home/mjmay/Lecture5-Images/tokens", "w+");
    if (!outfile) {
        return (-1); // Error!
    }
    while (1) {
        instring = fgets(mybuf, BUFLLEN, stdin); // catches overrun!

        // check for error or end of file (^D)
        if ( !instring || strlen(instring) == 0) break;

        // write string to output file, exit on error
        if (fputs(instring, outfile) < 0) break;
    }
    fclose(outfile); // Flushes from userspace
}
```

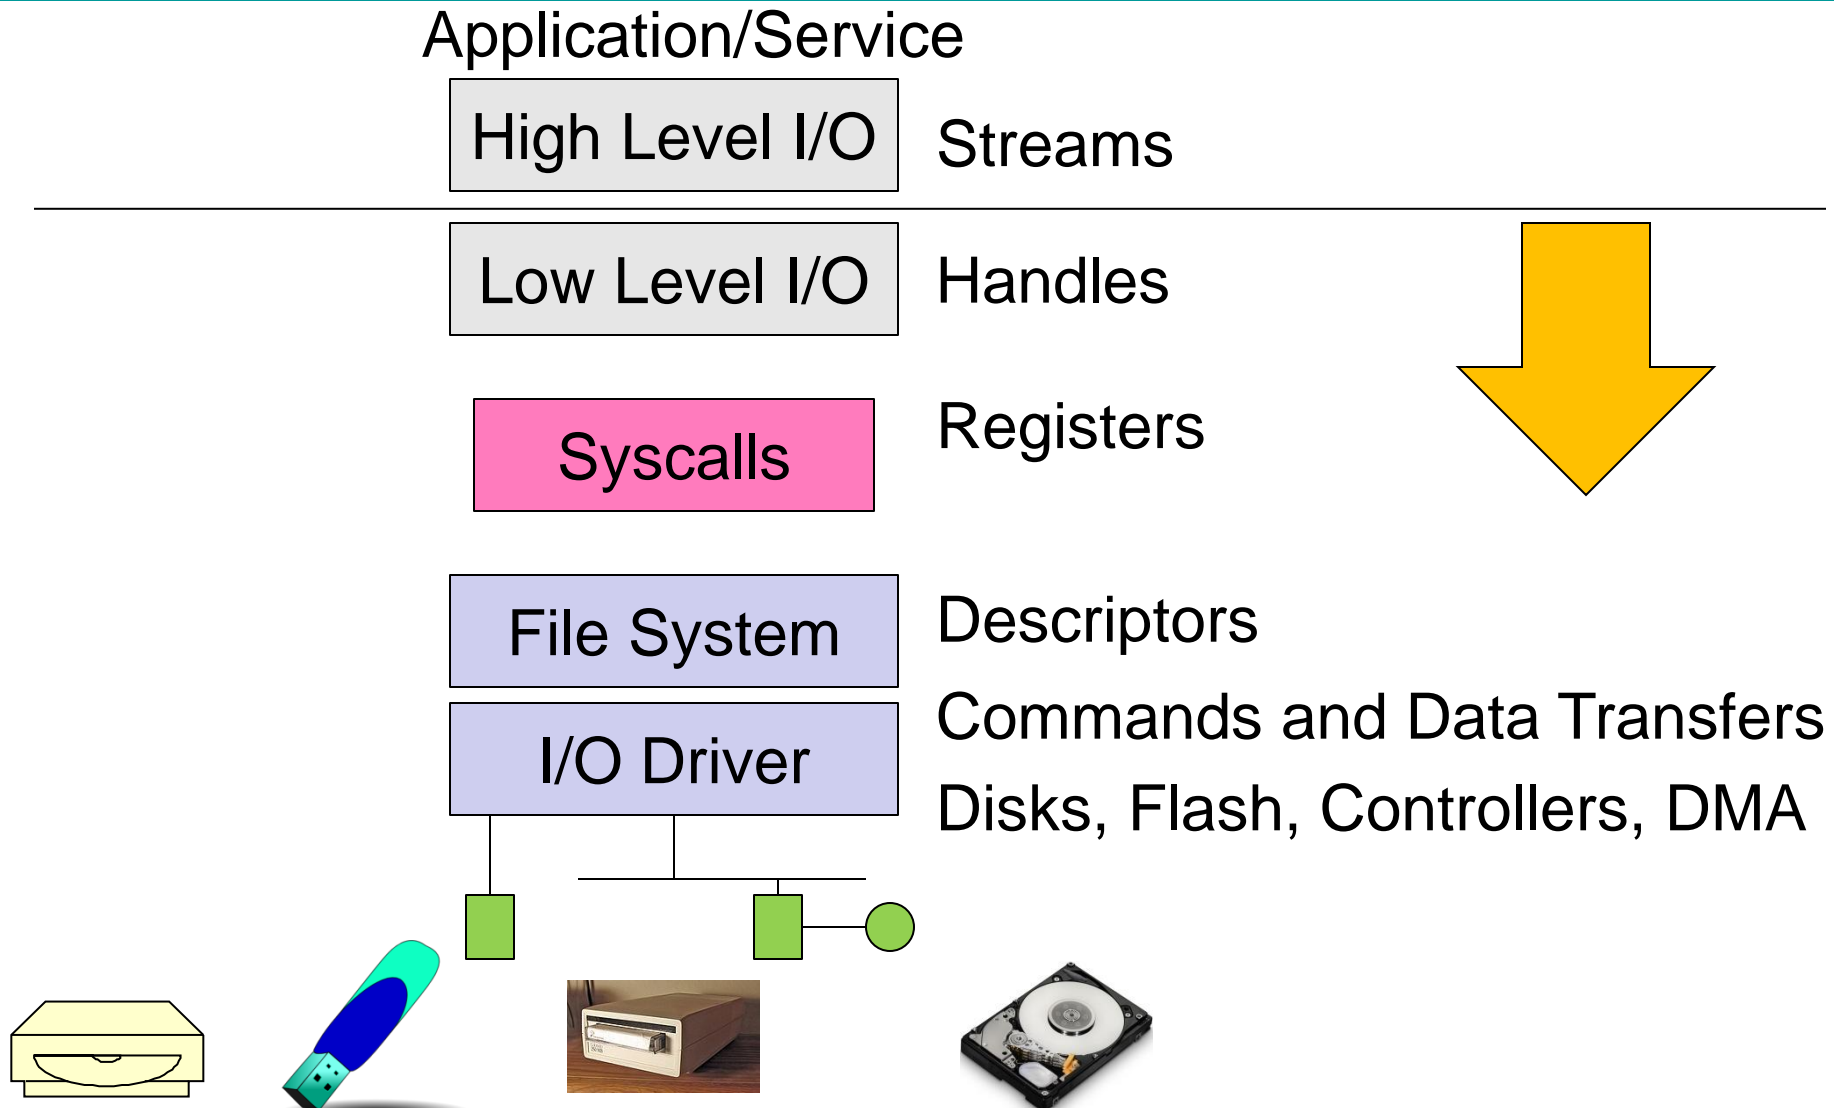
C Stream API positioning

```
int fseek(FILE *stream, long int offset, int whence);  
long int ftell (FILE *stream);  
void rewind (FILE *stream);
```



- Preserves high level abstraction of a uniform stream of objects
- Adds buffering for performance (**don't forget to flush fflush**)

What's below the surface?



C Low level I/O

- Operations on File Descriptors – as OS object representing the state of a file
 - User has a “handle” on the descriptor

http://www.gnu.org/software/libc/manual/html_node/Opening-and-Closing-Files.html

```
#include <fcntl.h>
#include <unistd.h>
#include <sys/types.h>

int open (const char *filename, int flags [, mode_t mode])
int creat (const char *filename, mode_t mode)
int close (int filedes)
```

Flags: Bit Vector of:

- Access Mode (Rd, Wr, ...)
- Open flags (Create,...)
- Operating modes (Appends,...)

mode: Bit Vector of Permission Bits

- User | Group | Other × R | W | X

C Low Level: standard descriptors

```
#include <unistd.h>
```

```
STDIN_FILENO - macro has value 0
```


```
STDOUT_FILENO - macro has value 1
```

```
STDERR_FILENO - macro has value 2
```

```
int fileno (FILE *stream)
```

```
FILE * fdopen (int filedes, const char *opentype)
```

- Crossing levels: File descriptors vs. streams
- Don't mix them!



C Low Level Operations

```
ssize_t read (int fildes, void *buffer, size_t maxsize)
/*returns bytes read, 0 => EOF, -1 => error*/
ssize_t write (int fildes, const void *buffer, size_t size)
/*returns bytes written*/

off_t lseek (int fildes, off_t offset, int whence)

int fsync (int fildes) /*- wait for i/o to finish*/
void sync (void) /*- wait for ALL to finish*/
```

- When write returns, data is on its way to disk and can be read, but it **may not actually be permanent!**

And lots more!

- TTYs versus files
- Memory mapped files
- File Locking
- Asynchronous I/O
- Generic I/O Control Operations
- Duplicating descriptors

```
int dup2 (int old, int new)
int dup (int old)
```

```
freopen(const char
*filename, const char
*mode, FILE *stream)
```

Example – lowio

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>

#define BUFSIZE 1024

int main (int argc, char *argv[]) {
    char buf[BUFSIZE];
    ssize_t writelen = write(STDOUT_FILENO, "I am a process.\n", 16);

    ssize_t readlen = read (STDIN_FILENO, buf, BUFSIZE);

    ssize_t strlen = snprintf(buf, BUFSIZE, "Got %zd chars\n", readlen);

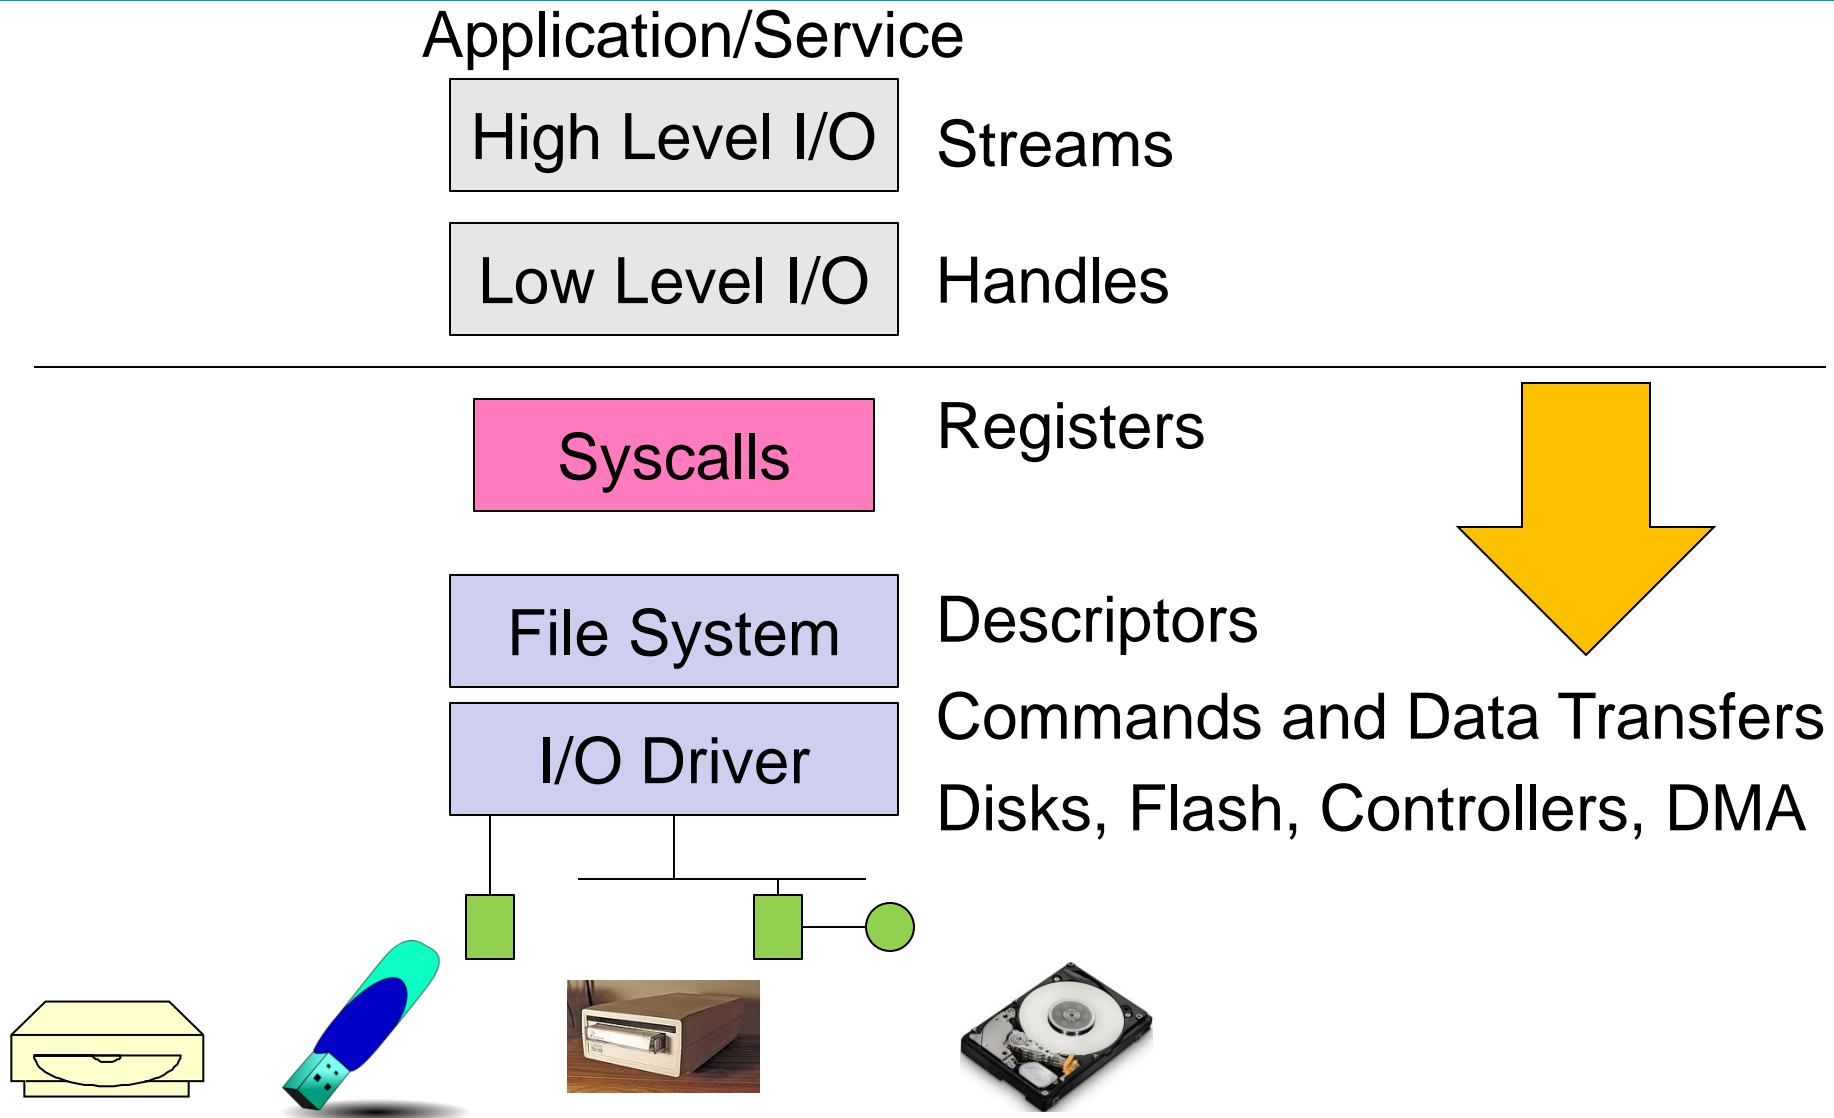
    writelen = strlen < BUFSIZE ? strlen : BUFSIZE;
    write (STDOUT_FILENO, buf, writelen);

    exit(0);
}
```


So Far

- Signals
- Syscalls
- Basic Support for I/O (drivers, etc.)
 - Files and Streams
 - Low Level
- I/O and Drivers
- Sockets and networks

What's below the surface?



Linux Syscalls

syscalls.kernelgrok.com

Search

Linux Syscall Reference

Show 10 entries

Search:

#	Name	Registers						Definition
		eax	ebx	ecx	edx	esi	edi	
0	sys_restart_syscall	0x00	-	-	-	-	-	kernel/signal.c:2058
1	sys_exit	0x01	int error_code	-	-	-	-	kernel/exit.c:1046
2	sys_fork	0x02	struct pt_regs *	-	-	-	-	arch/alpha/kernel/entry.S:716
3	sys_read	0x03	unsigned int fd	char __user *buf	size_t count	-	-	fs/read_write.c:391
4	sys_write	0x04	unsigned int fd	const char __user *buf	size_t count	-	-	fs/read_write.c:408
5	sys_open	0x05	const char __user *filename	int flags	int mode	-	-	fs/open.c:900
6	sys_close	0x06	unsigned int fd	-	-	-	-	fs/open.c:969
7	sys_waitpid	0x07	pid_t pid	int __user *stat_addr	int options	-	-	kernel/exit.c:1771
8	sys_creat	0x08	const char __user *pathname	int mode	-	-	-	fs/open.c:933
9	sys_link	0x09	const char __user *oldname	const char __user *newname	-	-	-	fs/namei.c:2520

Showing 1 to 10 of 338 entries

First Previous 1 2 3 4 5 Next Last

Generated from Linux kernel 2.6.35.4 using **Exuberant Ctags**, **Python**, and **DataTables**.
Project on **GitHub**. Hosted on **GitHub Pages**.

<https://filippo.io/linux-syscall-table/>

Linux Syscalls

Low level lib
parameters are set
up in registers and
syscall instruction is
issued

A type of
synchronous
exception that enters
well-defined entry
points into the kernel

Internal OS File Descriptor

- Internal Data Structure describing everything about the file
 - Where it resides
 - Its status
 - How to access it

lxr.free-electrons.com/source/include/linux/fs.h

```
875
876 struct file {
877     union {
878         struct llist_node    fu_llist;
879         struct rcu_head      fu_rcuhead;
880     } f_u;
881     struct path              f_path;
882     struct inode             *f_inode;    /* cached value */
883     const struct file_operations *f_op;
884
885     /*
886      * Protects f_ep_links, f_flags.
887      * Must not be taken from IRQ context.
888      */
889     spinlock_t              f_lock;
890     atomic_long_t           f_count;
891     unsigned int            f_flags;
892     fmode_t                 f_mode;
893     struct mutex            f_pos_lock;
894     loff_t                  f_pos;
895     struct fown_struct      f_owner;
896     const struct cred        *f_cred;
897     struct file_ra_state    f_ra;
898
899     u64                     f_version;
900 #ifdef CONFIG_SECURITY
901     void                    *f_security;
902 #endif
903     /* needed for tty driver, and maybe others */
904     void                    *private_data;
905
906 #ifdef CONFIG_EPOLL
907     /* Used by fs/eventpoll.c to link all the hooks to this file */
908     struct list_head        f_ep_links;
909     struct list_head        f_tfile_llink;
910 #endif /* #ifdef CONFIG_EPOLL */

```

<http://lxr.free-electrons.com/source/include/linux/fs.h>

File System: from syscall to driver

```
460 ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
461 {
462     ssize_t ret;
463
464     if (!(file->f_mode & FMODE_READ))
465         return -EBADF;
466     if (!(file->f_mode & FMODE_CAN_READ))
467         return -EINVAL;
468     if (unlikely(!access_ok(VERIFY_WRITE, buf, count)))
469         return -EFAULT;
470
471     ret = rw_verify_area(READ, file, pos, count);
472     if (!ret) {
473         if (count > MAX_RW_COUNT)
474             count = MAX_RW_COUNT;
475         ret = __vfs_read(file, buf, count, pos);
476         if (ret > 0) {
477             fsnotify_access(file);
478             add_rchar(current, ret);
479         }
480         inc_syscr(current);
481     }
482
483     return ret;
484 }
```

http://lxr.free-electrons.com/source/fs/read_write.c

File System: from syscall to driver

```
448 ssize_t __vfs_read(struct file *file, char __user *buf, size_t count,
449                    loff_t *pos)
450 {
451     if (file->f_op->read)
452         return file->f_op->read(file, buf, count, pos);
453     else if (file->f_op->read_iter)
454         return new_sync_read(file, buf, count, pos);
455     else
456         return -EINVAL;
457 }
```

http://lxr.free-electrons.com/source/fs/read_write.c

Low Level Driver



<http://lxr.free-electrons.com/source/include/linux/fs.h#L1679>

```
1679 struct file_operations {
1680     struct module *owner;
1681     loff_t (*llseek) (struct file *, loff_t, int);
1682     ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
1683     ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
1684     ssize_t (*read_iter) (struct kiocb *, struct iov_iter *);
1685     ssize_t (*write_iter) (struct kiocb *, struct iov_iter *);
1686     int (*iterate) (struct file *, struct dir_context *);
1687     int (*iterate_shared) (struct file *, struct dir_context *);
1688     unsigned int (*poll) (struct file *, struct poll_table_struct *);
1689     long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);
1690     long (*compat_ioctl) (struct file *, unsigned int, unsigned long);
1691     int (*mmap) (struct file *, struct vm_area_struct *);
1692     int (*open) (struct inode *, struct file *);
1693     int (*flush) (struct file *, fl_owner_t id);
1694     int (*release) (struct inode *, struct file *);
1695     int (*fsync) (struct file *, loff_t, loff_t, int datasync);
1696     int (*aio_fsync) (struct kiocb *, int datasync);
1697     int (*fasync) (int, struct file *, int);
1698     int (*lock) (struct file *, int, struct file_lock *);
1699     ssize_t (*sendpage) (struct file *, struct page *, int, size_t, loff_t *, int);
1700     unsigned long (*get_unmapped_area) (struct file *, unsigned long, unsigned long,
1701     int (*check_flags) (int);
1702     int (*flock) (struct file *, int, struct file_lock *);
1703     ssize_t (*splice_write) (struct pipe_inode_info *, struct file *, loff_t *, size
1704     ssize_t (*splice_read) (struct file *, loff_t *, struct pipe_inode_info *, size
1705     int (*setlease) (struct file *, long, struct file_lock **, void **);
1706     long (*fallocate) (struct file *file, int mode, loff_t offset,
1707                        loff_t len);
```


Device Drivers



Device Driver: Device-specific code in the kernel that interacts directly with the device hardware

- Supports a standard, internal interface
- Same kernel I/O system can interact easily with different device drivers
- Special device-specific configuration supported with the `ioctl()` system call

Device Drivers typically divided into two pieces:

- **Top half:** accessed in call path from system calls
 - implements a set of **standard, cross-device calls** like `open()`, `close()`, `read()`, `write()`, `ioctl()`, `strategy()`
 - This is the kernel's interface to the device driver
 - Top half will *start* I/O to device, may put thread to **sleep** until finished
- **Bottom half:** run as interrupt routine
 - Gets input or transfers next block of output
 - May **wake** sleeping threads if I/O now complete

Low Level Driver



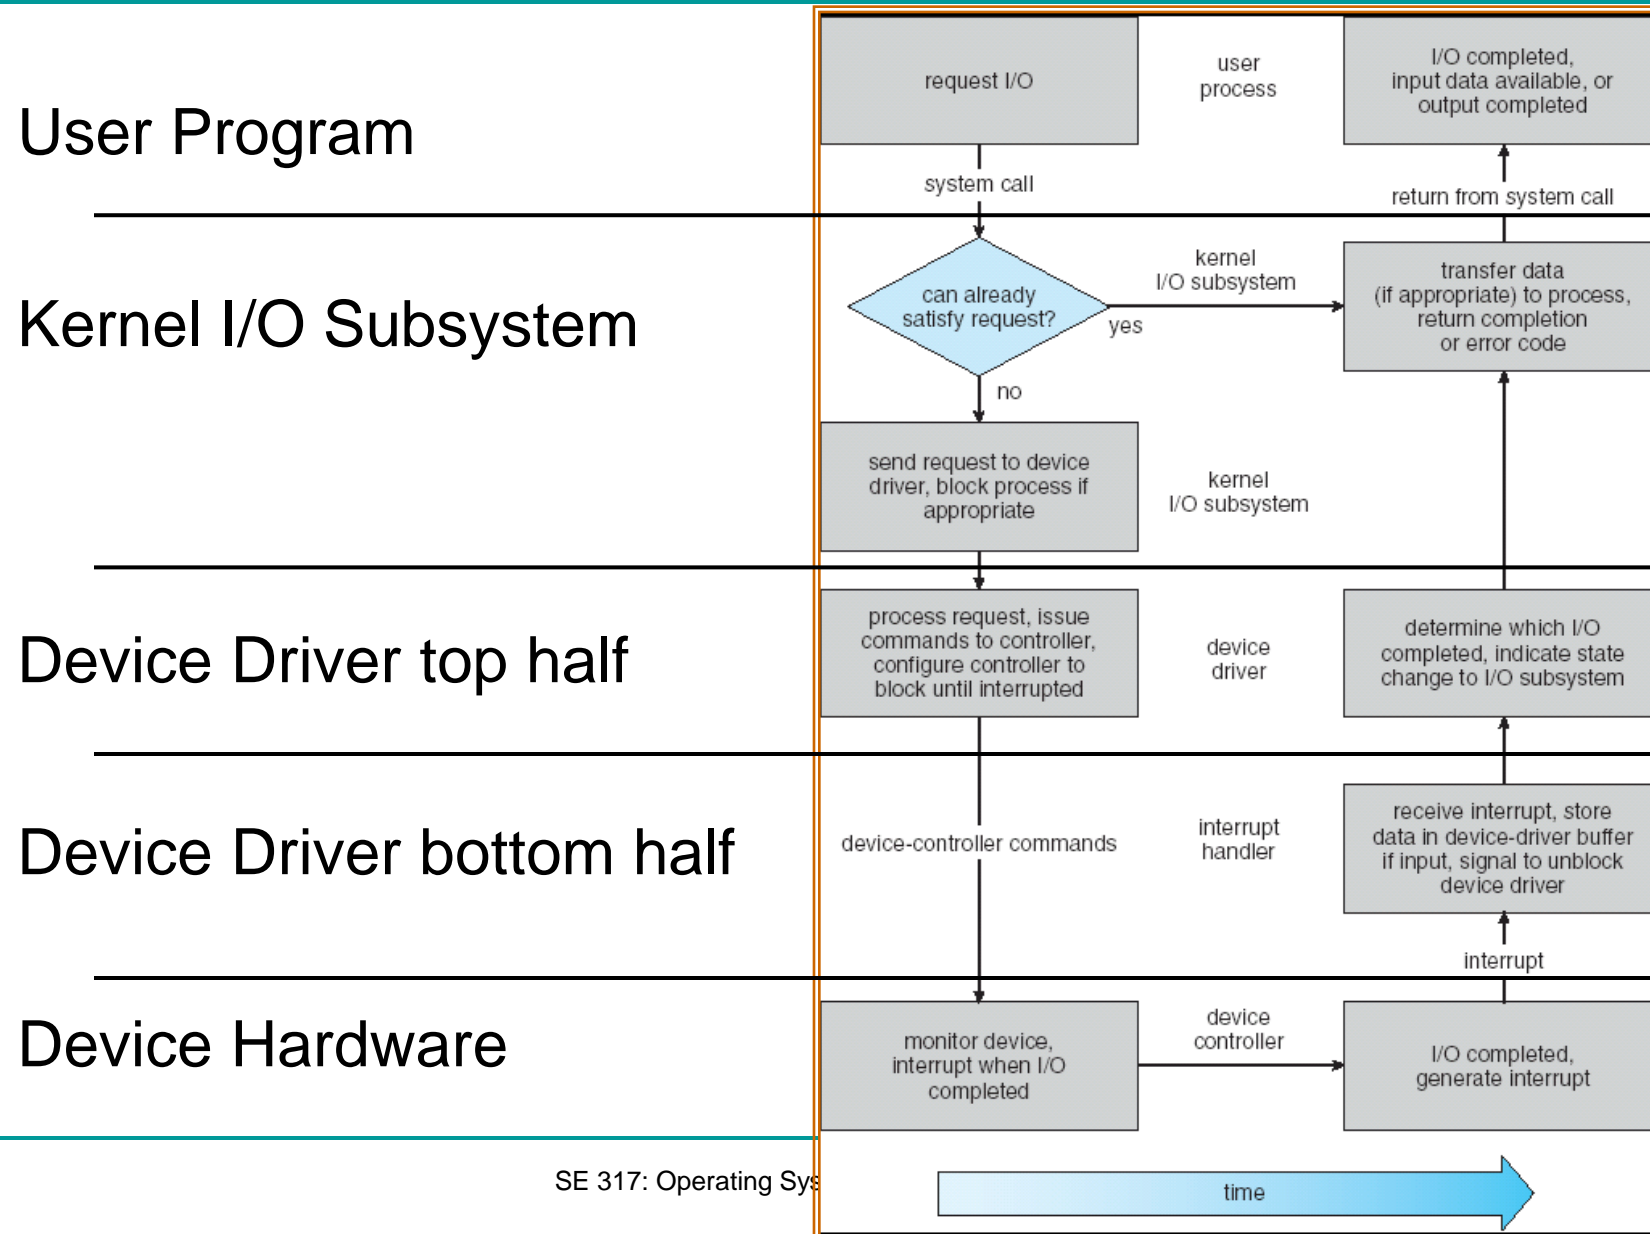
Photo by [Lewis J Goetz](#) on [Unsplash](#)

Associated with
particular
hardware device

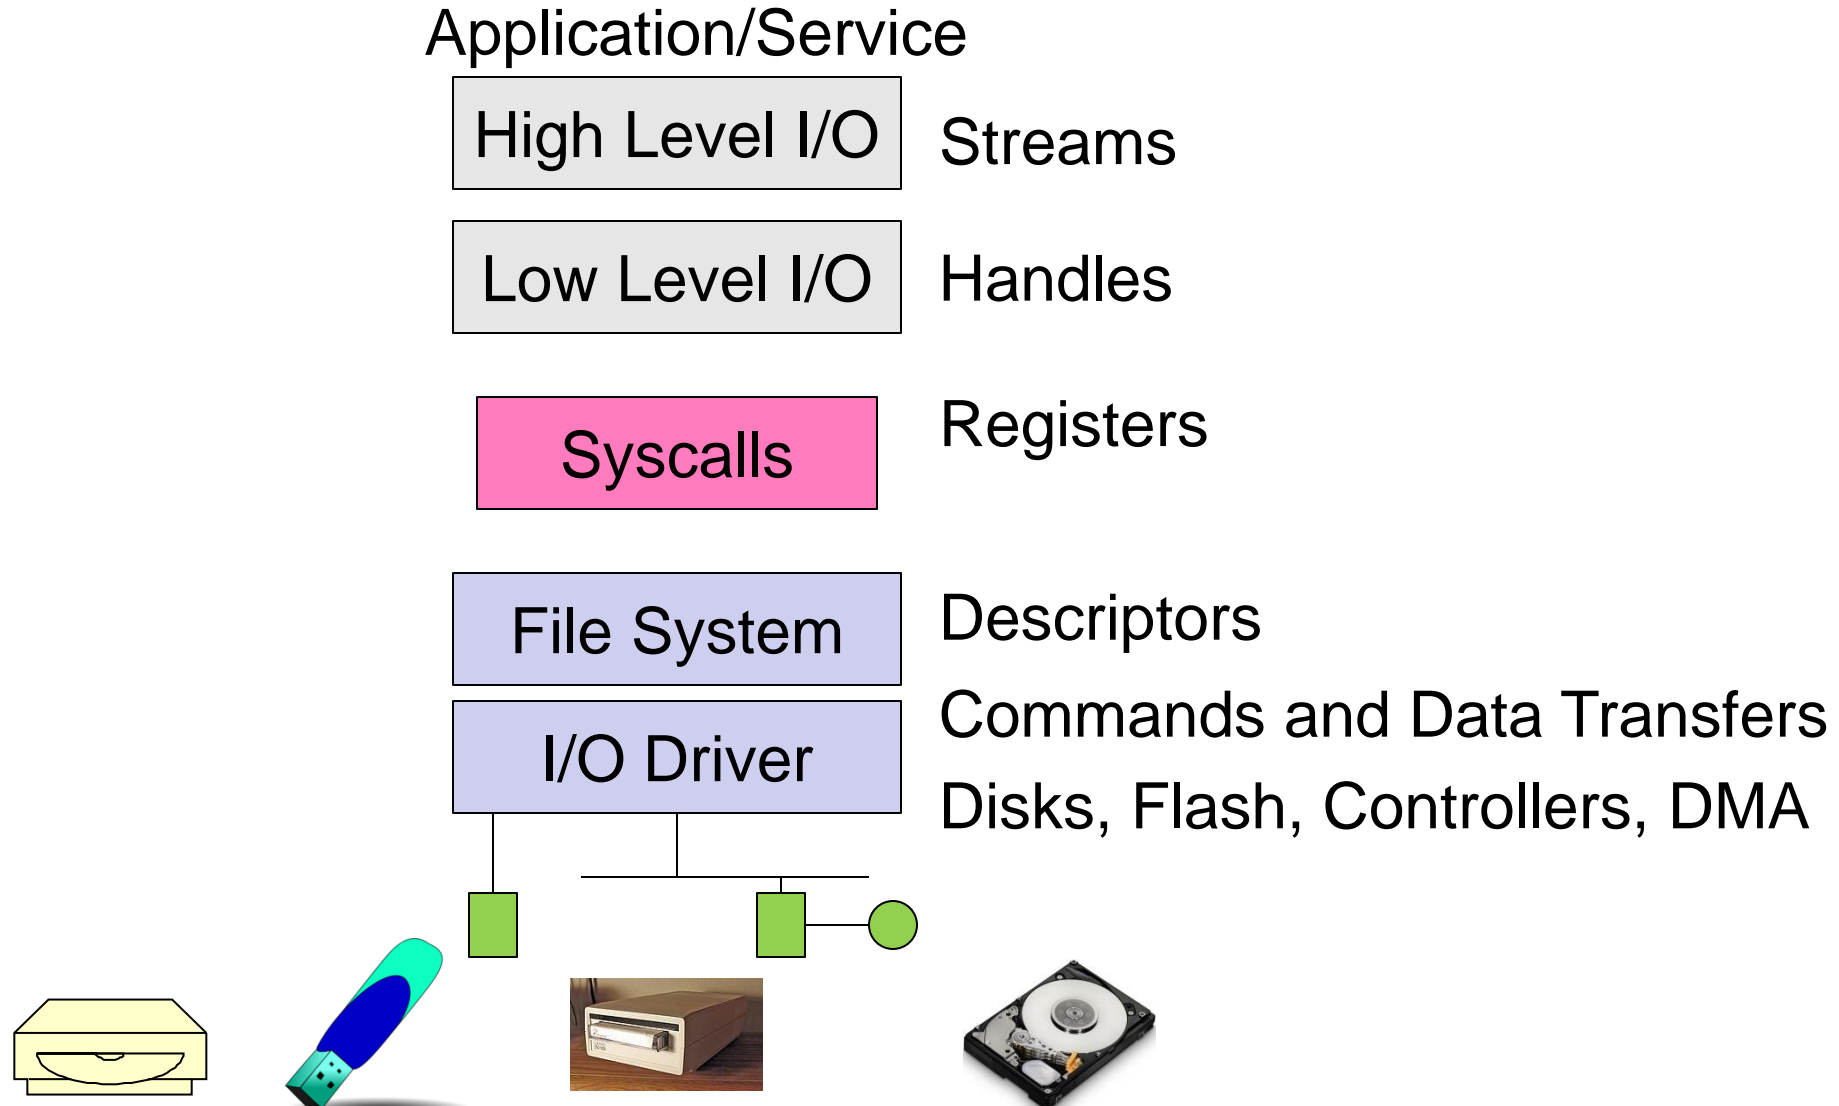
Registers /
Unregisters
itself with the
kernel

Handler functions for each of the file
operations

Life Cycle of An I/O Request



So what happens when you fgetc?




So Far

- Signals
- Syscalls
- Basic Support for I/O (drivers, etc.)
 - Files and Streams
 - Low Level
- I/O and Drivers
- Sockets and networks

Communication Between Processes

- Can we view files as communication channels?

```
write(wfd, wbuf, wlen);
```




```
n=read(rfd, rbuf, rmax);
```

- **Producer** and **Consumer** of a file may be different processes
 - May be separated in time (or not)
- However, what if data written once and consumed once?
 - Would be more like a queue, but still look like **File I/O**!

Communication across the world

```
write(wfd, wbuf, wlen);
```



- Connected queues over the internet
 - What's the analog of **open**?
 - What is the **namespace**?
 - How are they **connected** in time?

Request Response Protocol

Client (issues requests)

```
write(wfd, wbuf, wlen);
```

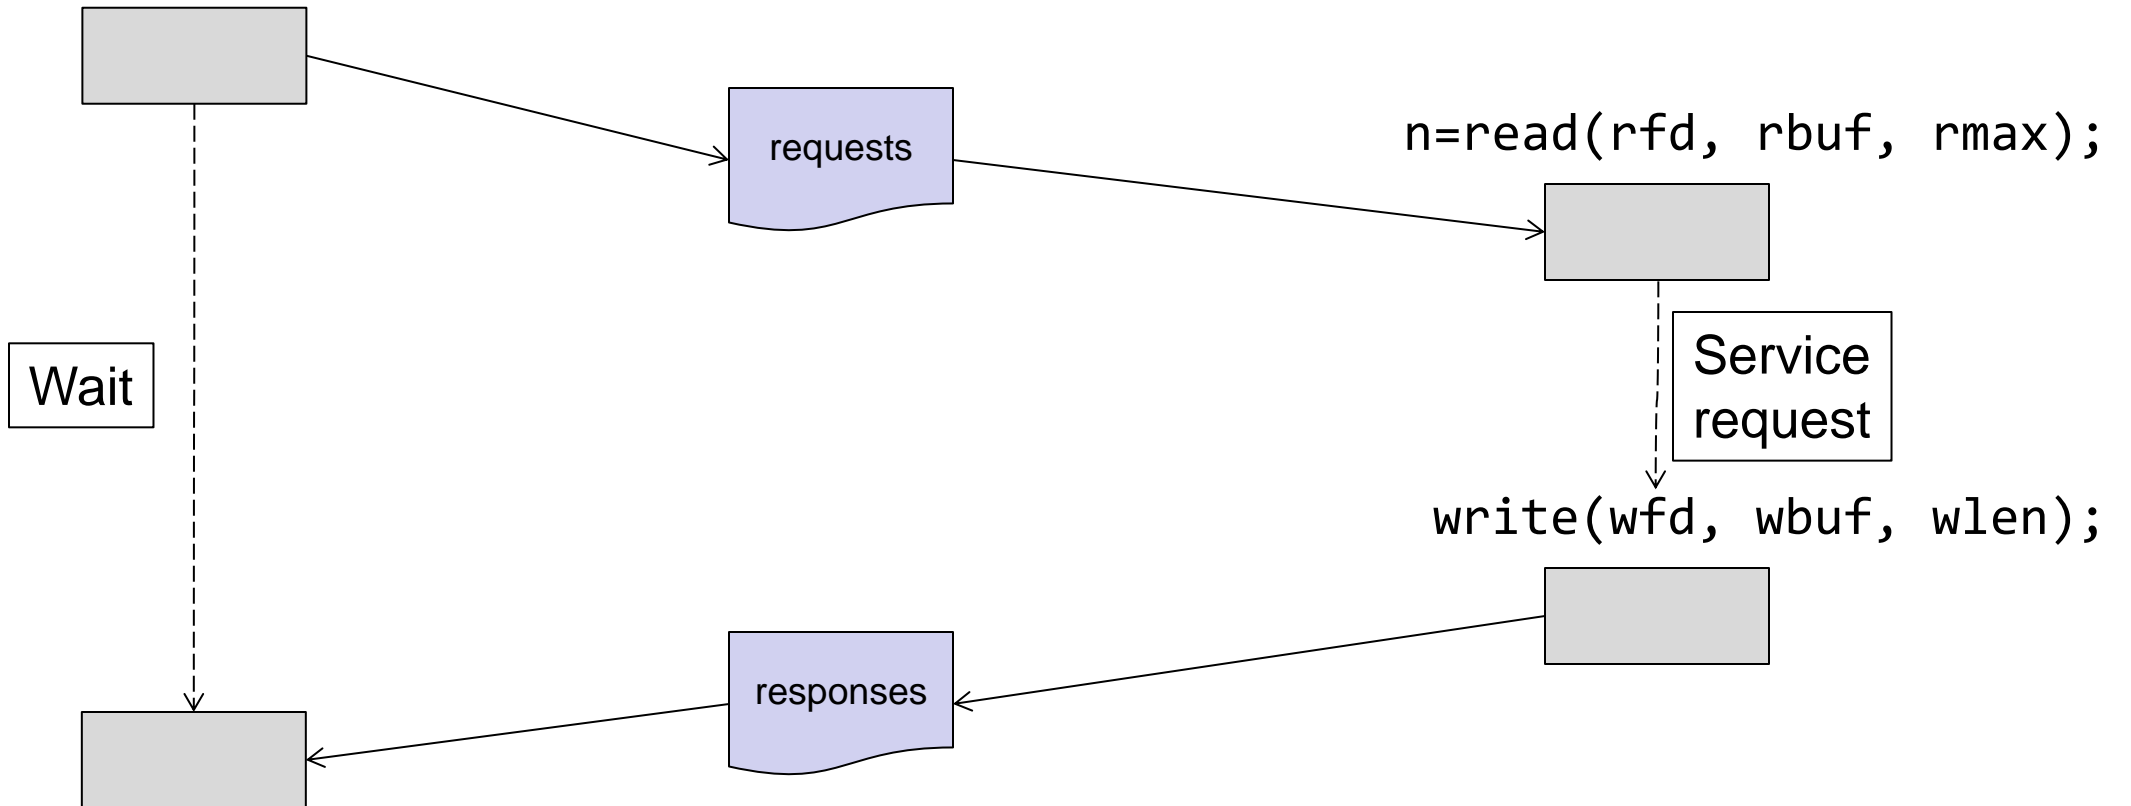
Server (performs operations)

```
n=read(rfd, rbuf, rmax);
```

Service request

```
write(wfd, wbuf, wlen);
```

```
n=read(rfd, rbuf, rmax);
```

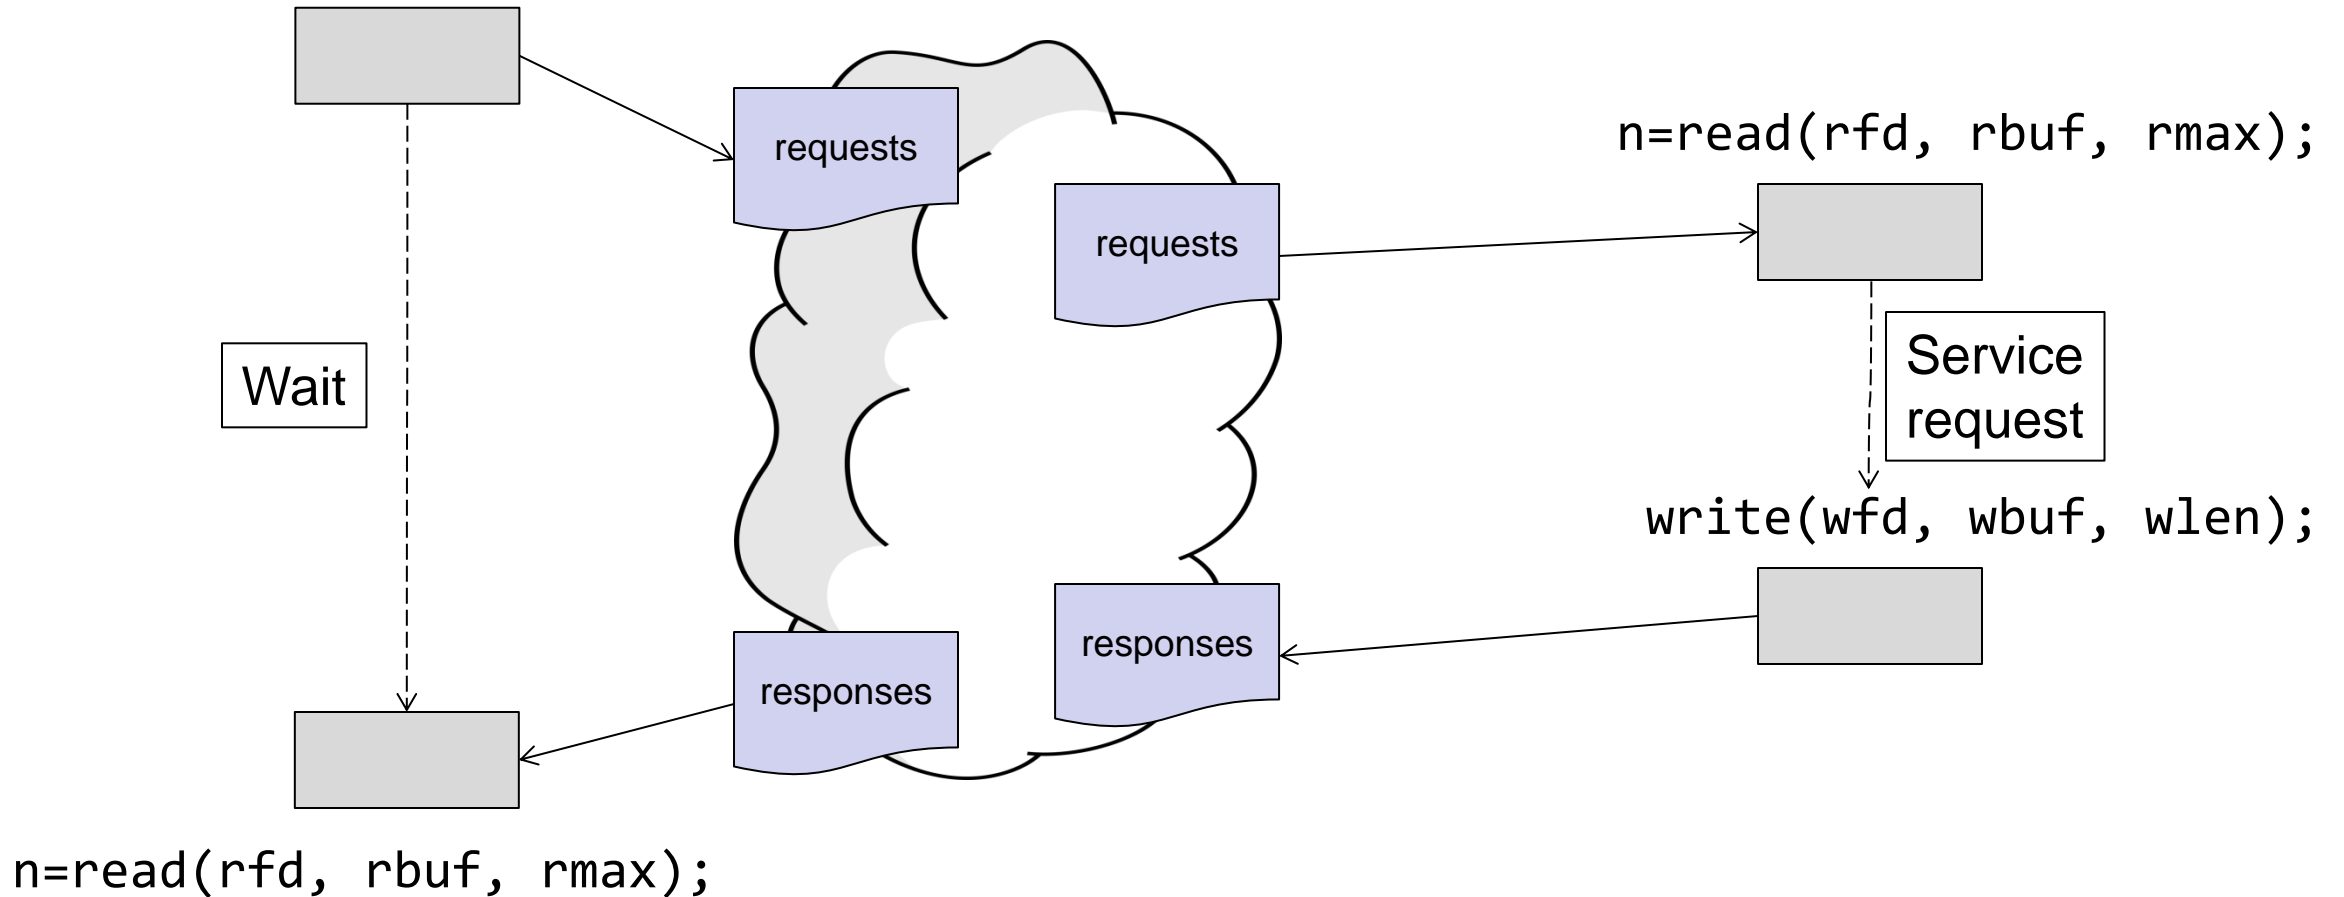


Request Response Protocol


Client (issues requests)

```
write(wfd, wbuf, wlen);
```

Server (performs operations)



Client-Server Models



- File servers, Web servers, FTP servers, Databases
- Many clients access a common server

Sockets



Socket: an abstraction of a network I/O queue

- Mechanism for inter-process communication
- Embodies **one side** of a communication channel
 - Same interface regardless of location of other end
 - Could be local machine (“UNIX socket”) or **remote machine** (“network socket”)
- First introduced in 4.2 BSD UNIX: big innovation at time
 - Now most operating systems provide some notion of socket

Data transfer like files

- Read / Write against a descriptor

Over **any kind** of network

- Local to a machine
- Over the internet (TCP/IP, UDP/IP)
- OSI, Appletalk, SNA, IPX, SIP, NS, ...

Socket Creation and Connection

File systems

- Provide a collection of permanent objects in structured name space
- Processes open, read/write/close them
- Files exist independent of the processes

Sockets

- Provide a means for processes to communicate (transfer data) to other processes.
 - Creation and connection is more complex
 - Form 2-way pipes between processes, possibly worlds away

Conclusion

- Signals
- Syscalls
- Basic Support for I/O (drivers, etc.)
 - Files and Streams
 - Low Level
- I/O and Drivers
- Sockets and networks