

Cooperating Threads, Synchronization Mutual Exclusion, Semaphores

19 December 2024
Lecture 7

Slides adapted from John Kubiawicz (UC Berkeley)

Concept Review

Thread
lifecycle

Thread join

Kernel
supported
threads

User
supported
threads

Scheduler
activation

yield()

switch()

Cooperating
threads

Topics for Today

- Concurrency challenge
- Motivation for Synchronization and Locks
- Atomic Read-Modify-Write Operations
- Higher Level Synchronization Atoms
 - Semaphores
 - Monitors

[illegible]

ATOMIC Operations

- To understand a **concurrent program**, we need to know what the underlying **indivisible operations** are!
- **Atomic Operation**: an operation that always runs to completion or not at all
 - It is *indivisible*: it cannot be stopped in the middle and state cannot be modified by someone else in the middle
 - Fundamental building block – **if no atomic operations, then have no way for threads to work together**



(1966-1973)

“As always, should you or any of your IM force be caught or killed, the Secretary will disavow any knowledge of your actions.”

“Good luck, Jim. This tape will self-destruct in five seconds.”

ATOMIC Operations

- On most machines, memory references and assignments (i.e. loads and stores) of words are atomic
 - Consequently – weird example that produces “3” on previous slide **can’t happen**
- Many instructions are not atomic
 - Double-precision floating point store often not atomic
 - VAX and IBM 360 had an instruction to **copy a whole array**

Correctness Requirements

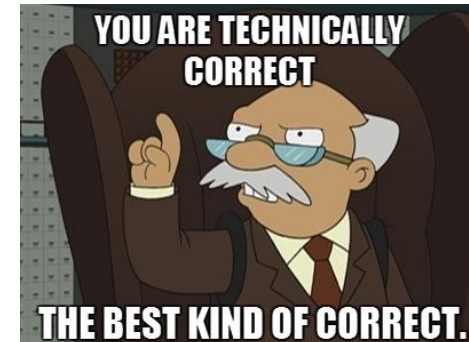
Threaded programs
must work for all
interleavings of
thread instruction
sequences

Cooperating threads
inherently non-
deterministic and
non-reproducible

Image source: <http://knowyourmeme.com/photos/909991-futurama>



Really hard to
debug unless
carefully designed!



Example: Therac-25

- Machine for radiation therapy
 - Software control of electron accelerator and electron beam/
X-Ray production
 - Software control of dosage
- Software errors caused the death of several patients
 - A series of race conditions on shared variables and poor software design
- “They determined that data entry speed during editing was the key factor in producing the error condition: If the prescription data was edited at a fast pace, the overdose occurred.”

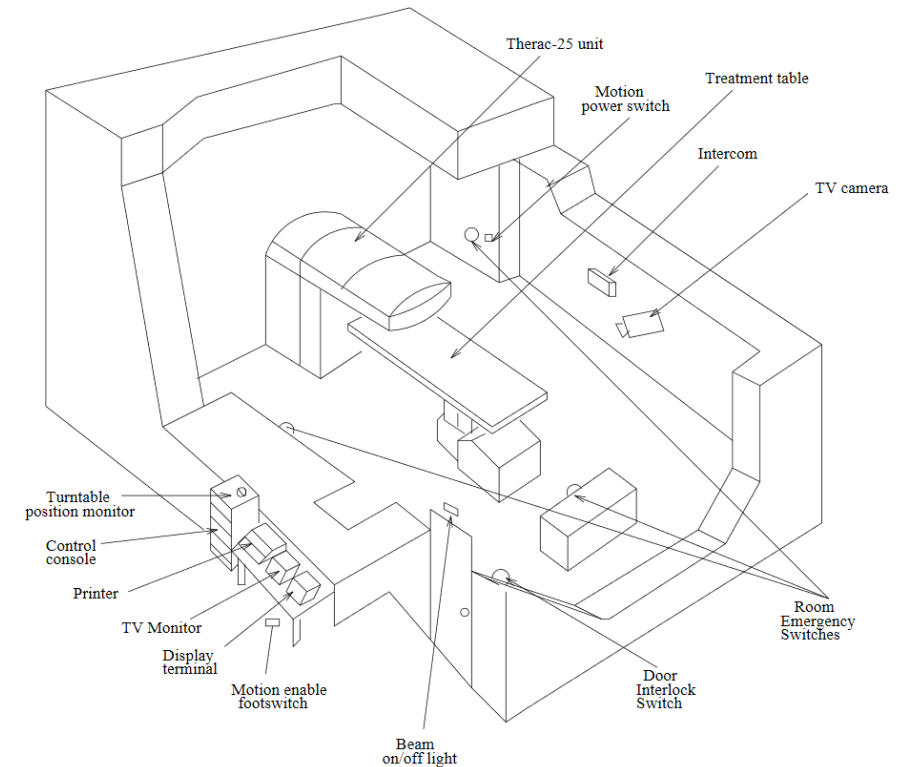


Figure 5: A typical Therac-25 facility after the final CAP.

Another Concurrent Program Example

- Two threads, A and B, compete with each other
 - One tries to increment a shared counter
 - The other tries to decrement the counter



Thread A

```
i = 0;
while (i < 10)
    i = i + 1;
printf("A wins!");
```

Thread B

```
i = 0;
while (i > -10)
    i = i - 1;
printf("B wins!");
```

- Assume that memory loads and stores are atomic, but incrementing and decrementing are *not atomic*
- Who wins? Could be *either*
- Is it **guaranteed** that someone wins? Why or why not?
- What if both threads have their *own CPU* running at same speed? Is it guaranteed that it goes on forever?

Hand Simulation Multiprocessor Example

- Inner loop looks like this:

Thread A
r1=0 load r1, M[i]
r1=1 add r1, r1, 1
M[i]=1 store r1, M[i]

Thread B
r1=0 load r1, M[i]
r1=-1 sub r1, r1, 1
M[i]=-1 store r1, M[i]

- **Hand Simulation:**

- And we're off. A gets off to an early start
- B says "hmph, better go fast" and tries really hard
- A goes ahead and writes "1"
- B goes and writes "-1"
- A says "HUH??? I could have sworn I put a 1 there"

- Could this happen on a uniprocessor?

- Yes! Unlikely, but if you are depending on it not happening, it will and your system will break...



So Far

- Concurrency challenge
- Motivation for Synchronization and Locks
- Atomic Read-Modify-Write Operations
- Higher Level Synchronization Atoms
 - Semaphores
 - Monitors

Motivation: “Too much humus”

- Great thing about OS’s – analogy between problems in OS and problems in real life
 - Help you understand real life problems better
 - But, computers are much stupider than people
- Example: People need to coordinate:



Time	Alice	Bob
3:00	Look in Fridge. Out of humus.	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of humus.
3:15	Buy humus	Leave for store
3:20	Arrive at home, put humus away	Arrive at store
3:25		Buy humus
3:30		Arrive at home, put humus away

Definitions

Synchronization: using atomic operations to ensure cooperation between threads

- For now, only loads and stores are atomic
- We are going to show that its hard to build anything useful with only reads and writes

Mutual Exclusion: ensuring that only one thread does a particular thing at a time

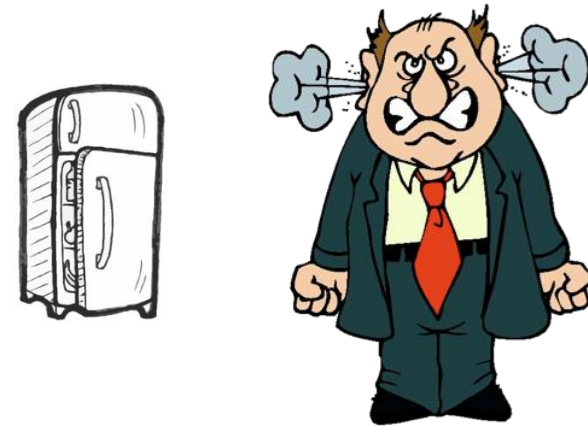
- One thread *excludes* the other while doing its task

Critical Section: piece of code that only one thread can execute at once. Only one thread at a time will get into this section of code.

- Critical section is the result of mutual exclusion
- Critical section and mutual exclusion are two ways of describing the same thing.

More Definitions

- **Lock:** prevents someone from doing something
 - Lock before entering critical section and before accessing shared data
 - Unlock when leaving, after accessing shared data
 - Wait if locked
 - Important idea: all synchronization involves waiting
- For example: fix the humus problem by putting a **key** on the refrigerator
 - Lock it and take key if you are going to go buy humus
 - Fixes too much: roommate angry if only wants OJ
- Of Course – We don't know how to make a lock yet



Too Much Humus: Correctness Properties

- Need to be careful about correctness of concurrent programs, since non-deterministic
 - Always **write down behavior first**
 - Impulse is to start coding first, then when it doesn't work, pull hair out
 - Instead, think first, then code
- What are the correctness properties for the “Too much humus” problem???
- Never more than one person buys
- Someone buys if needed
- Restrict ourselves to use only atomic load and store operations as building blocks

Too Much Humus: Solution #1

- Use a note to avoid buying too much humus:
 - Leave a note before buying (kind of “**lock**”)
 - Remove note after buying (kind of “**unlock**”)
 - Don’t buy if there’s a note (**wait**)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noHumus) {  
    if (noNote) {  
        leave Note;  
        buy humus;  
        remove note;  
    }  
}
```



- Result?
 - Still too much humus **but only occasionally!**
 - Thread can get context switched after checking humus and note but before buying humus!
- Solution makes problem worse since fails **intermittently**
 - Makes it really hard to debug...
 - Must work despite what the dispatcher does!

Too Much Humus: Solution #1½

- Clearly the Note is not quite blocking enough
 - Let's try to fix this by placing note first
- Another try at previous solution:

```
leave Note;  
if (noHumus) {  
    if (noNote) {  
        leave Note;  
        buy humus;  
    }  
}  
remove note;
```

- What happens here?
 - Well, with human, probably nothing bad
 - With computer: no one ever buys humus



Too Much Humus Solution #2

- How about labeled notes?
 - Now we can leave note before checking
- Algorithm looks like this:

Thread A

```
leave note A;  
if (noNote B) {  
    if (noHumus) {  
        buy Humus;  
    }  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noHumus) {  
        buy Humus;  
    }  
}  
remove note B;
```

Too Much Humus Solution #2

Does this work?

Possible for neither thread to buy humus

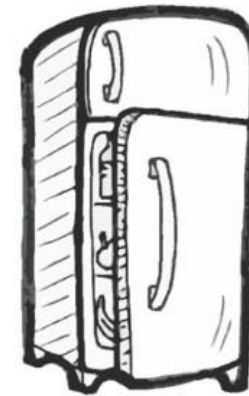
- Context switches at exactly the wrong time can lead each to think that the other is going to buy

Really insidious:

- **Extremely unlikely** that this would happen, but will at worst possible time
- Probably something like this in UNIX

Too Much Humus Solution #2 Problem

- *I'm* not getting humus, *You're* getting humus
- This kind of lockup is called “starvation!”



Too Much Humus Solution #3

- Here is a possible **two-note** solution:

Thread A

```
leave note A;
while (note B) { //X
    do nothing;
}
if (noHumus) {
    buy Humus;
}
remove note A;
```

Thread B

```
leave note B;
if (noNote A) { //Y
    if (noHumus) {
        buy Humus;
    }
}
remove note B;
```

Too Much Humus Solution #3

- Does this work? **Yes**. Both can guarantee that:
 - It is safe to buy, or
 - Other will buy, ok to quit
- At **X**:
 - if no note B, safe for A to buy,
 - otherwise wait to find out what will happen
- At **Y**:
 - if no note A, safe for B to buy
 - Otherwise, A is either buying or waiting for B to quit

Solution #3 discussion

- Our solution protects a single “Critical-Section” piece of code for each thread:

```
if (noHumus) {  
    buy humus;  
}
```

- Solution #3 works, but it's unsatisfactory
 - Really complex – even for this simple an example
 - Hard to convince yourself that this really works
 - A's code is different from B's – what if you have many threads?
 - Code would have to be slightly different for each thread
 - While A is waiting, it is consuming CPU time
 - This is “busy-waiting”

Solution #3 discussion

There's a better way:

- Have hardware provide better (**higher-level**) primitives than atomic load and store
- Build even **higher-level programming abstractions** on this new hardware support

Too Much Humus: Solution #4

- Let's make an implementation of a lock (more later).
 - `Lock.Acquire()` – wait until lock is free, then grab
 - `Lock.Release()` – Unlock, waking up anyone waiting
 - Must be **atomic operations** – if **two** threads are waiting for the lock and both see it's free, only **one** succeeds in grabbing the lock
- Then, our humus problem is easy:

```
humuslock.Acquire();  
if (noHumus)  
    buy humus;  
humuslock.Release();
```
- Section of code between `Acquire()` and `Release()` is a **“Critical Section”**
- You can make this even simpler: suppose you are out of ice cream instead of humus
 - Skip the test since you always need more ice cream.



Where are we going with synchronization?

Programs	Shared Programs
Higher Level API	Locks, Semaphores, Monitors, Send/Receive
Hardware	Load/Store, Disable Interrupts, Test & Set, Compare & Swap

- We are going to implement various **higher-level** synchronization primitives using **atomic operations**
 - Everything is pretty painful if the only atomic primitives are load and store
 - Need to provide primitives which are useful at user-level

So Far

- Concurrency challenge
- Motivation for Synchronization and Locks
- Atomic Read-Modify-Write Operations
- Higher Level Synchronization Atoms
 - Semaphores
 - Monitors

Atomic Read-Modify-Write instructions

- Problems with interrupts only based solution:
 1. Can't give lock implementation to **users**
 2. Doesn't work well on **multiprocessor**
 - Disabling interrupts on **all processors** requires **messages** and would be very time consuming
- Alternative: **Atomic Instruction Sequences**

Instructions that read a value from memory and write a new value atomically

Hardware is responsible for implementing this correctly

- On uniprocessors (not too hard)
- On multiprocessors (requires help from **cache coherence protocol**)

Unlike disabling interrupts, can be used on **both uniprocessors and multiprocessors**

Examples of Read-Modify-Write

```
test&set (&address) {      /* most architectures */
    result = M[address];
    M[address] = 1;
    return result;
}
```

```
swap (&address, register) { /* x86 */
    temp = M[address];
    M[address] = register;
    register = temp;
}
```

Examples of Read-Modify-Write

```
compare&swap (&address, reg1, reg2) { /* 68000 */  
    if (reg1 == M[address]) {  
        M[address] = reg2;  
        return success;  
    } else {  
        return failure;  
    }  
}
```

Examples of Read-Modify-Write

```
load-linked&store conditional(&address) {  
    /* R4000, alpha */  
    loop:  
        ll r1, M[address];  
        movi r2, 1;          /* Can do arbitrary comp */  
        sc r2, M[address];  
        beqz r2, loop;  
}
```


Implementing Locks with test&set

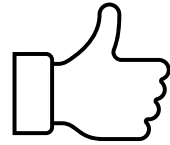
- Another flawed, but simple solution:

```
int value = 0; // Free
Acquire() {
    while (test&set(value)); // while busy
}
Release() {
    value = 0;
}
```

- Simple explanation:
 - If lock is **free**, test&set reads **0** and sets value=1, so lock is now busy. It returns 0 so while exits.
 - If lock is **busy**, test&set reads **1** and sets value=1 (no change). It returns 1, so while loop continues
 - When we set **value = 0**, someone else can get lock
- **Busy-Waiting**: thread consumes cycles while waiting

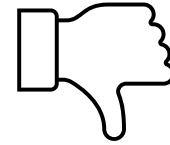
Problem: Busy-Waiting for Lock

Positives



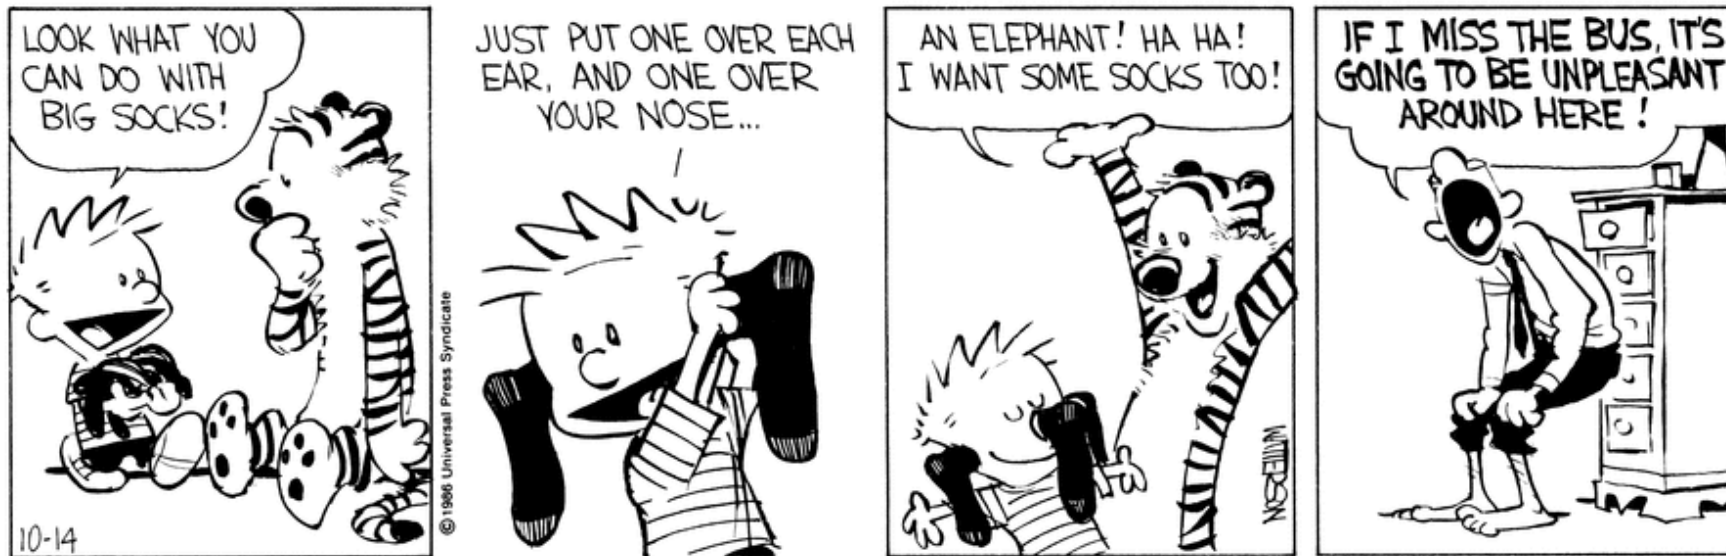
- Machine can receive interrupts
- User code can use the lock
- Works on a multiprocessor

Negatives



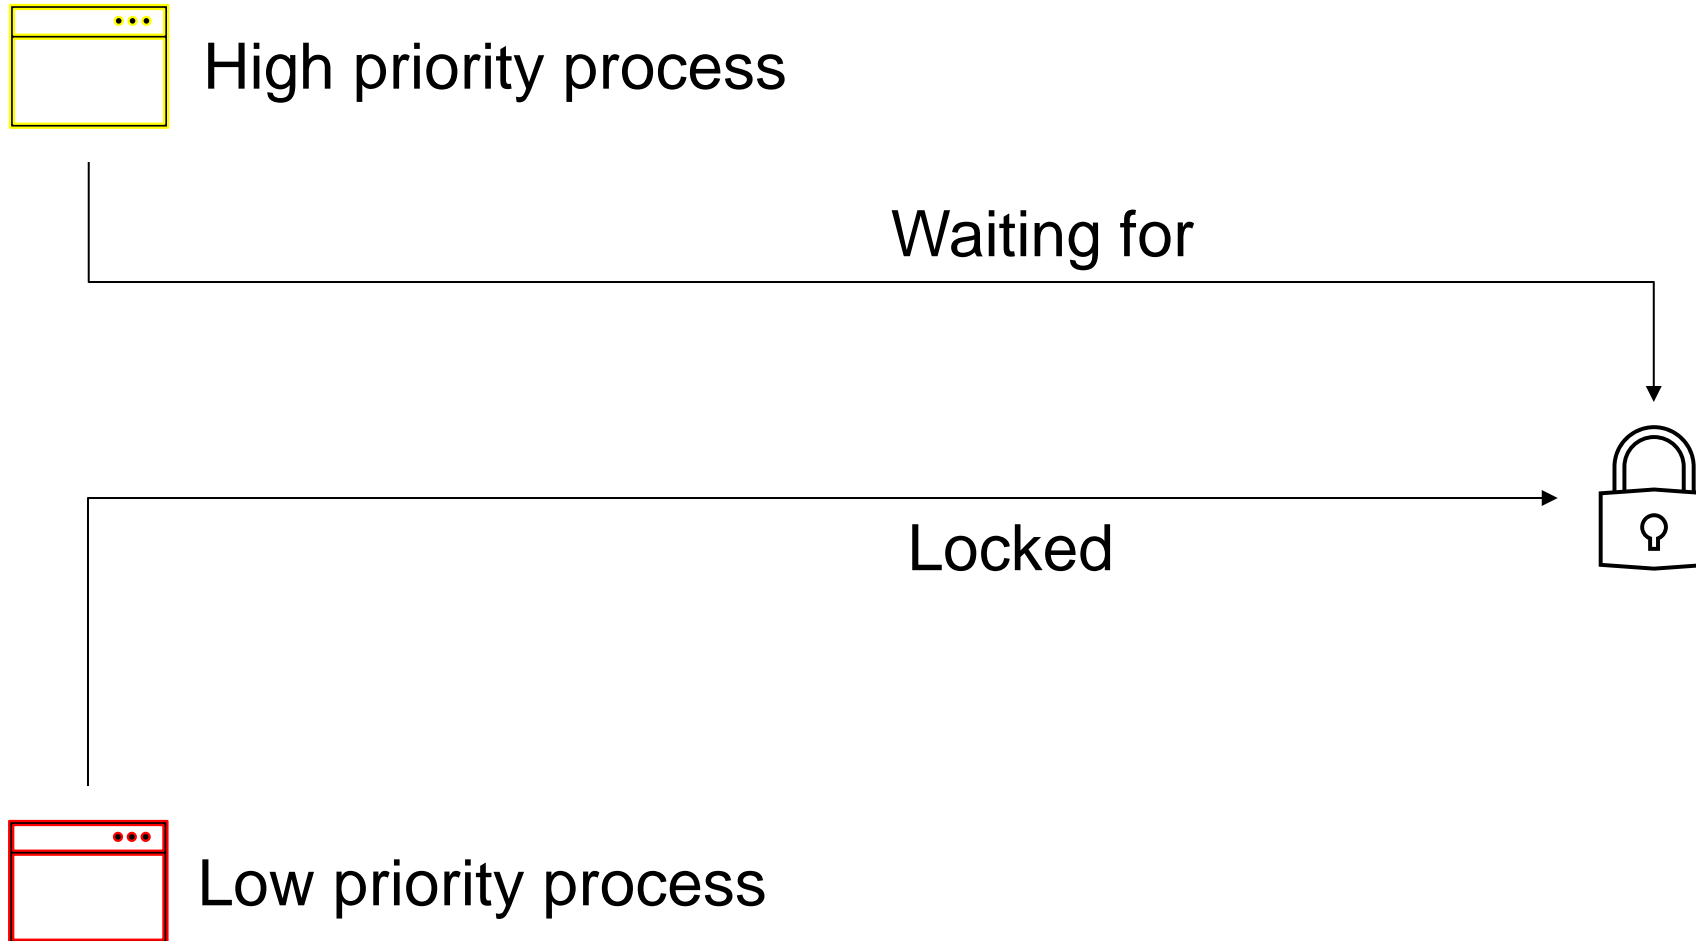
- **Very inefficient** because the busy-waiting thread consumes cycles waiting
- Waiting thread may take cycles away from thread holding lock (no one wins!)
- **Priority Inversion**: If busy-waiting thread has higher priority than thread holding lock \Rightarrow no progress!
 - Priority Inversion problem with original **Martian rover**

Priority inversion

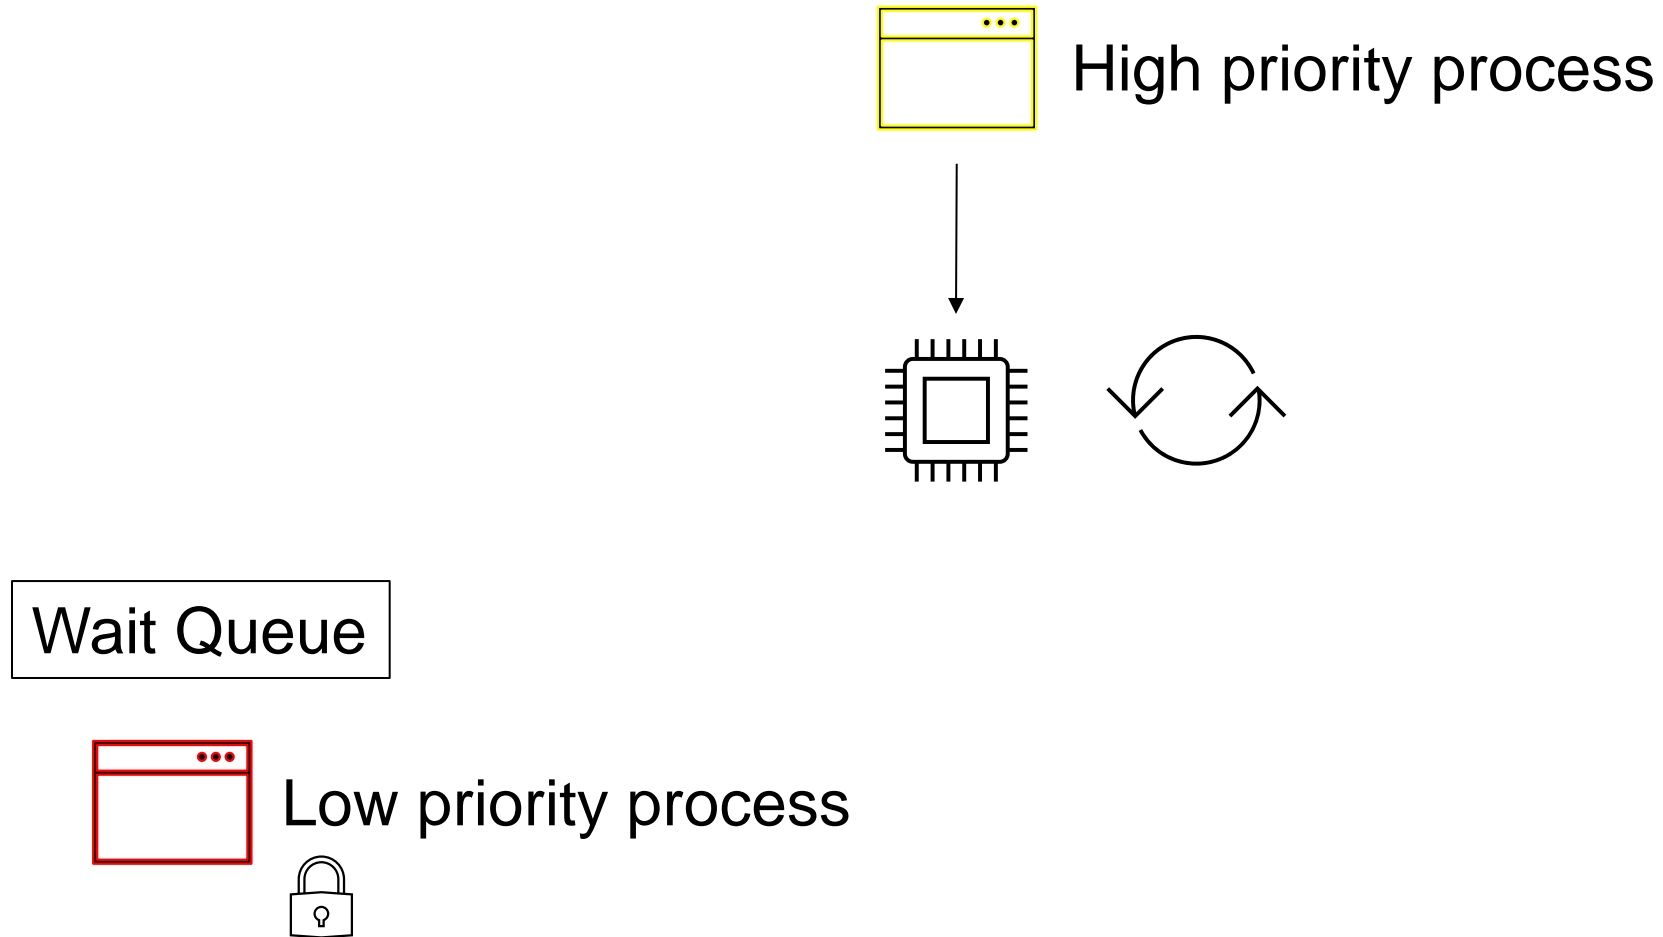


Calvin and Hobbes by Bill Watterson for October 14, 1986
<https://www.gocomics.com/calvinandhobbes/1986/10/14>

Priority Inversion



Priority Inversion



Problem: Busy-Waiting for Lock

For **semaphores and monitors**,
waiting thread may wait for an
arbitrary length of time!

- Even if busy-waiting OK for **locks**
definitely not ok for **other primitives**
- Homework/exam solutions should
not have busy-waiting!

Multiprocessor Spin Locks: test&test&set

- A better solution for multiprocessors:

```
int mylock = 0; // Free
Acquire() {
    do {
        while(mylock); // Wait until might be free
    } while(test&set(&mylock)); // exit if get lock
}

Release() {
    mylock = 0;
}
```

- Simple explanation:
 - Wait until lock **might be** free (only reading – stays in cache)
 - Then, **try** to grab lock with test&set
 - Repeat if **fail** to actually get lock
- Issues with this solution:
 - **Busy-Waiting**: thread still consumes cycles while waiting
 - However, it does not impact other processors!

Better Locks using test&set

Can we build test&set locks **without** busy-waiting?

- Can't entirely, but can minimize!
- Idea: only busy-wait to atomically check lock value

```
int guard = 0;
int value = FREE;
Acquire() {
    // Short busy-wait time
    while (test&set(guard));
    if (value == BUSY) {
        put thread on wait queue;
        go to sleep() & guard = 0;
    } else {
        value = BUSY;
        guard = 0;
    }
}
```

```
Release() {
    // Short busy-wait time
    while (test&set(guard));
    if anyone on wait queue {
        take thread off wait queue
        Place on ready queue;
    } else {
        value = FREE;
    }
    guard = 0;
}
```

- **Note:** sleep has to be sure to reset the guard variable
 - Why can't we do it just before or just after the sleep?

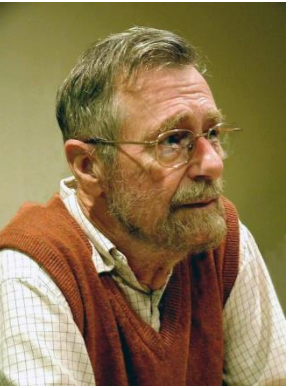
So Far

- Concurrency challenge
- Motivation for Synchronization and Locks
- Atomic Read-Modify-Write Operations
- Higher Level Synchronization Atoms
 - Semaphores
 - Monitors

Higher-level Primitives than Locks

- Goal so far:
 - What is the right abstraction for synchronizing threads that share memory?
 - Want as high a level primitive as possible
- Good **primitives and practices** important!
 - Since execution is **not entirely sequential**, really hard to find bugs, since they happen rarely
 - UNIX is **stable** now, but up until mid-80s (10 years after started), systems running UNIX would **crash every week or so** – concurrency bugs
- **Synchronization** is a way of coordinating multiple concurrent activities that are using **shared state**
 - We need **paradigms** to structure the sharing

Semaphores

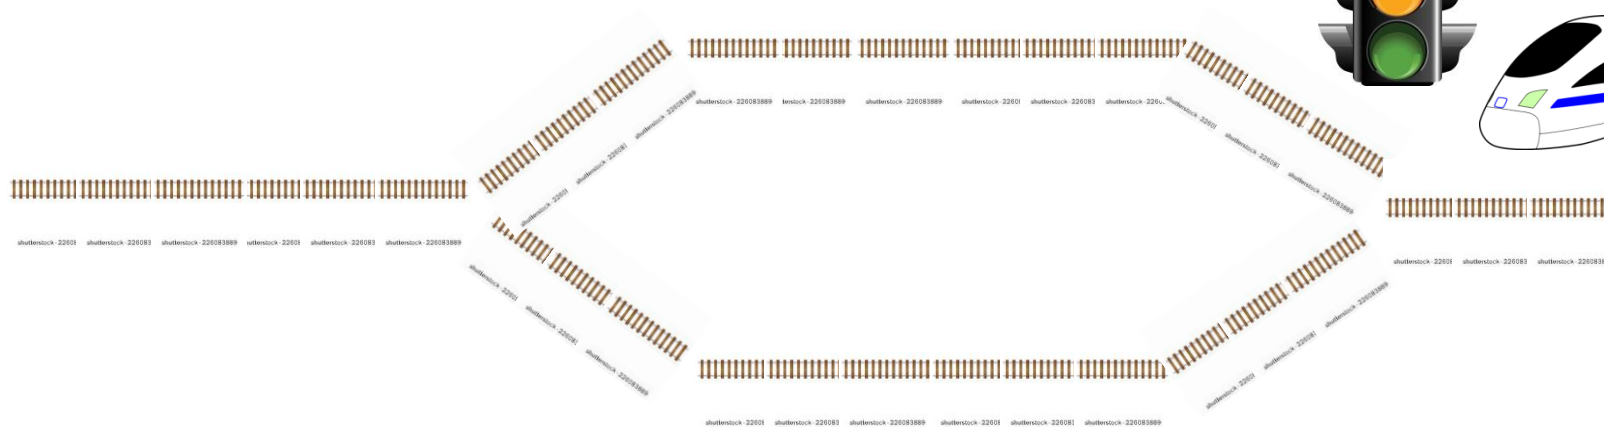


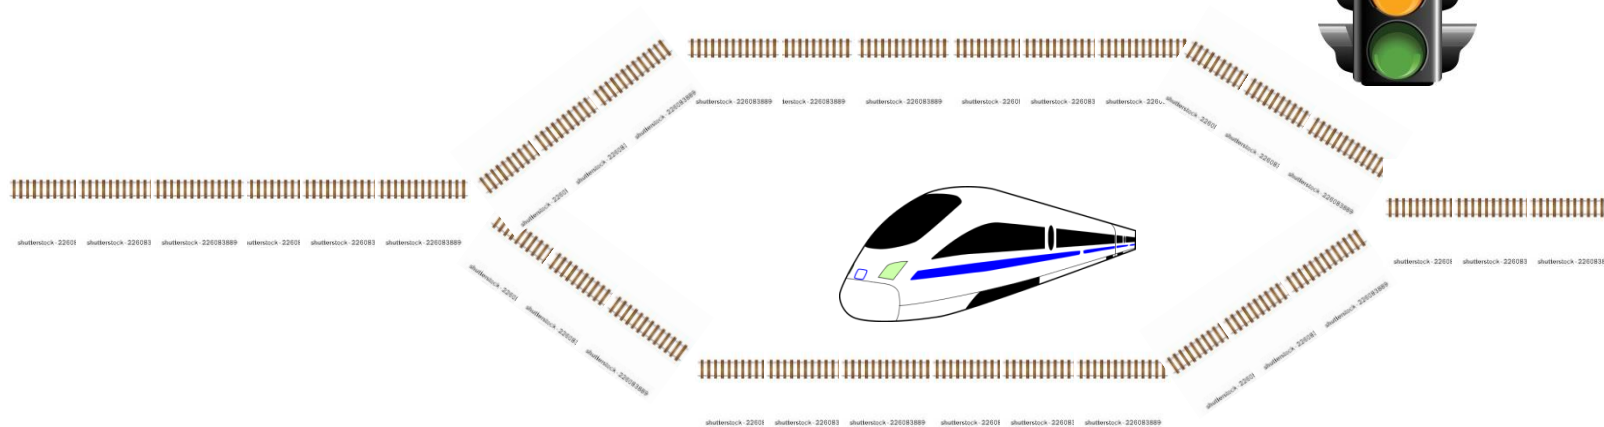
- Semaphores are a kind of generalized lock
 - First defined by Dijkstra in late 60s
 - Main synchronization primitive used in original UNIX
- Definition: *A Semaphore has a non-negative integer value and supports the following two operations:*
 - $P()$: an atomic operation that waits for semaphore to become positive, then decrements it by 1
 - Think of this as the `wait()` operation
 - $V()$: an atomic operation that increments the semaphore by 1, waking up a waiting P , if any
 - Think of this as the `signal()` operation
 - Note that $P()$ stands for “*proberen*” (to test) and $V()$ stands for “*verhogen*” (to increment) in Dutch

<https://www.youtube.com/watch?v=LKQpy107yUY>

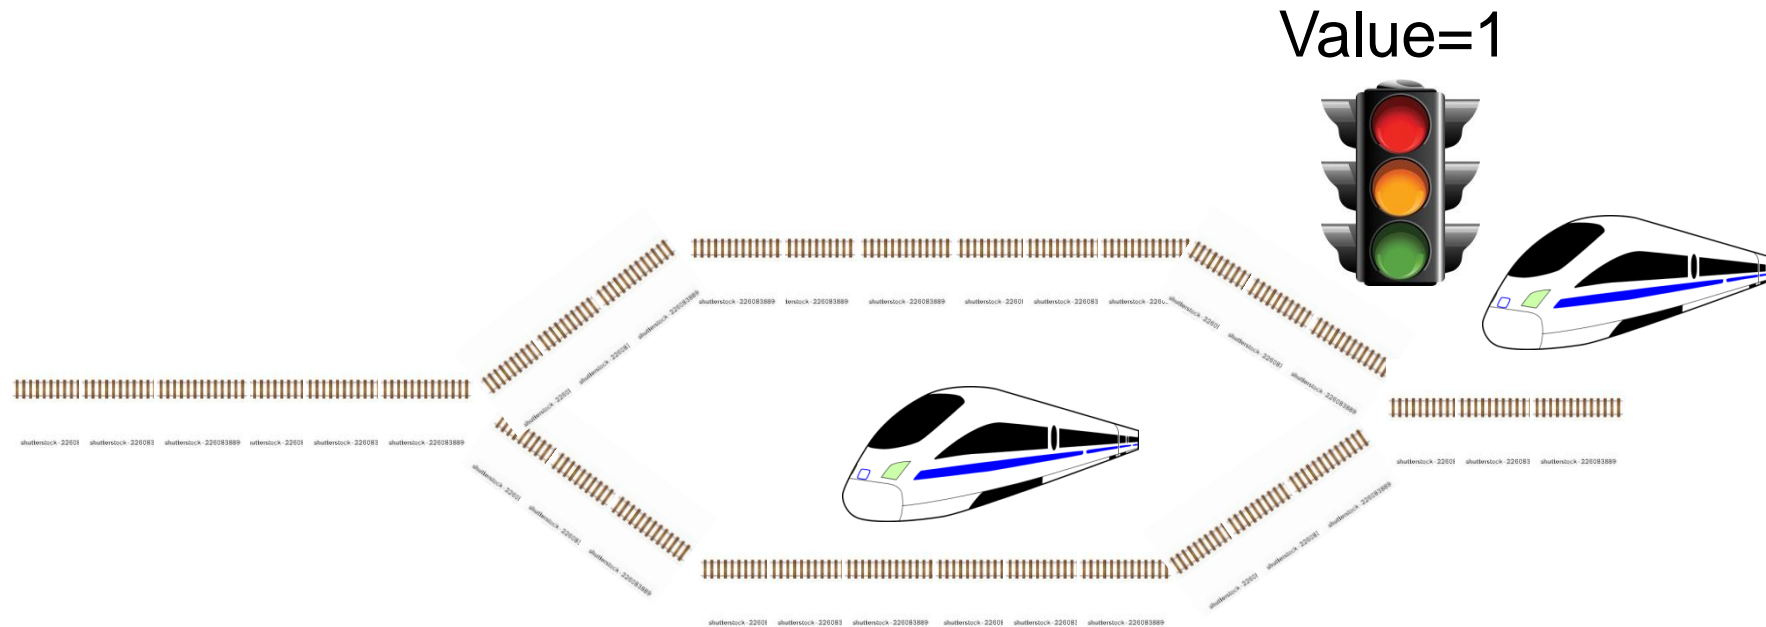
Semaphores Like Integers Except

- Semaphores are like integers, except
 - No negative values
 - Only operations allowed are P and V – can't read or write value, except to set it initially
 - Operations must be **atomic**
 - Two P's together **can't decrement** value below zero
 - Similarly, thread going to sleep in P won't miss wakeup from V – even if they **both happen at same time**
- Semaphore from **railway analogy**
 - Here is a semaphore initialized to 2 for resource control:

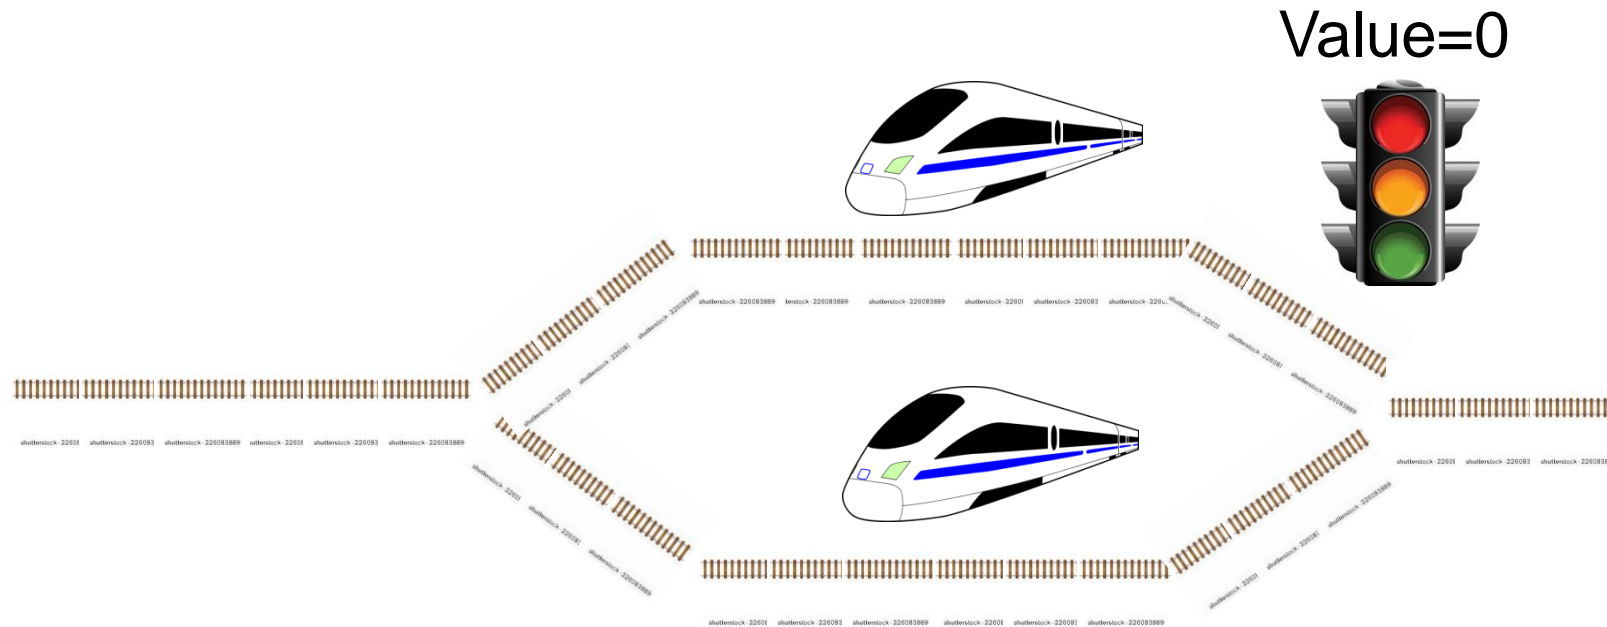


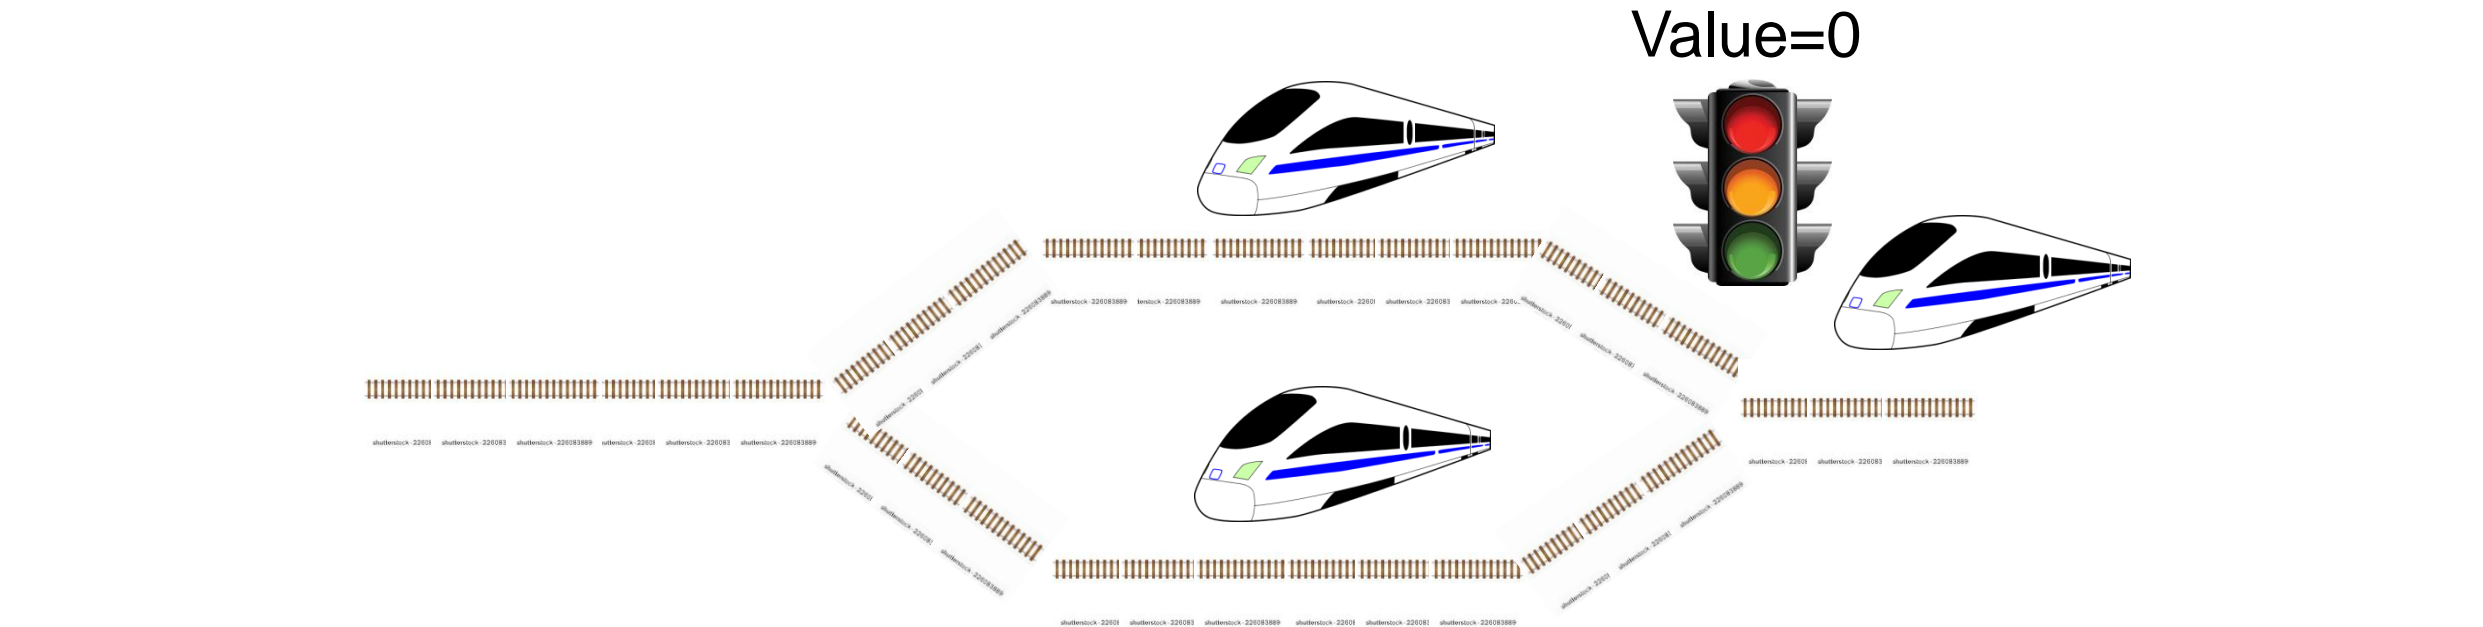


Railroad Analogy Step 3

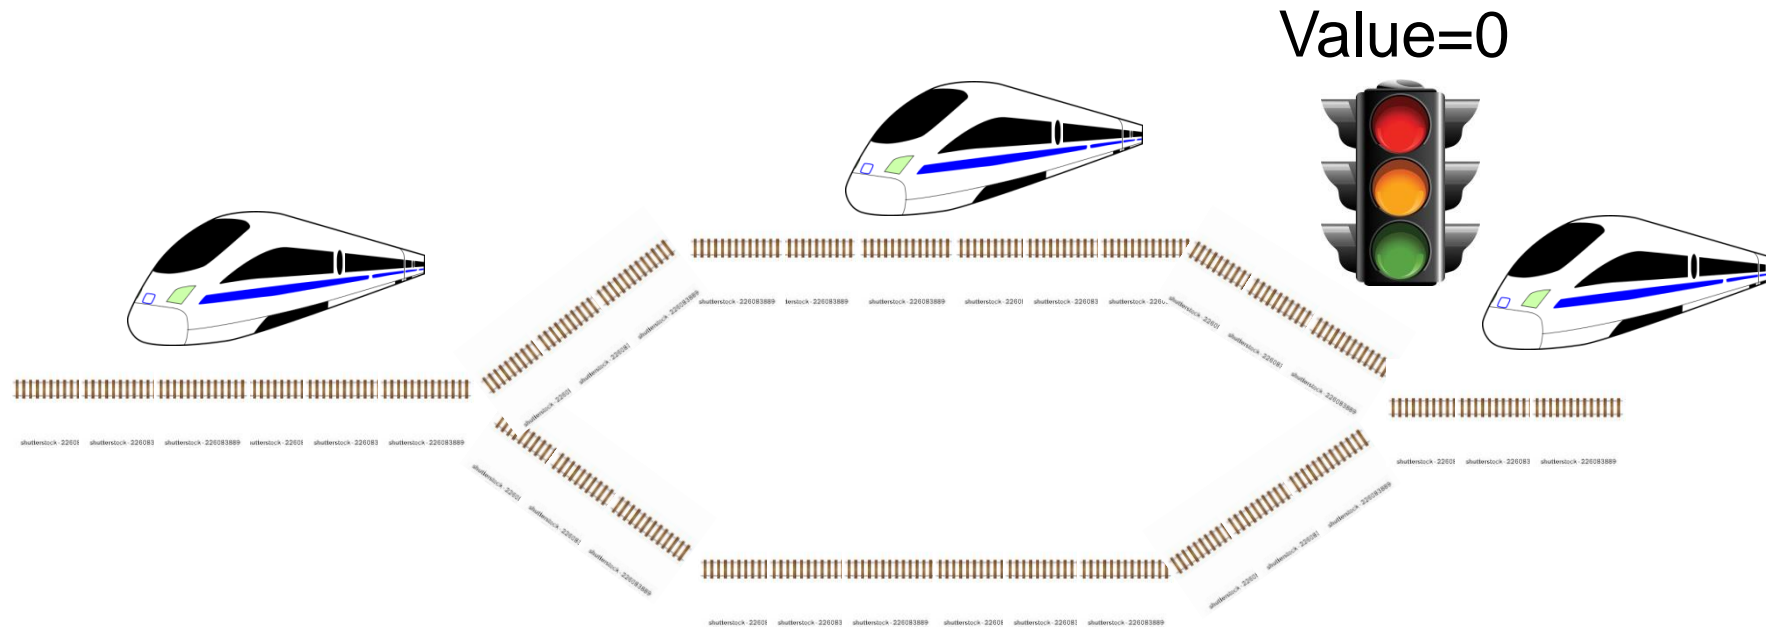


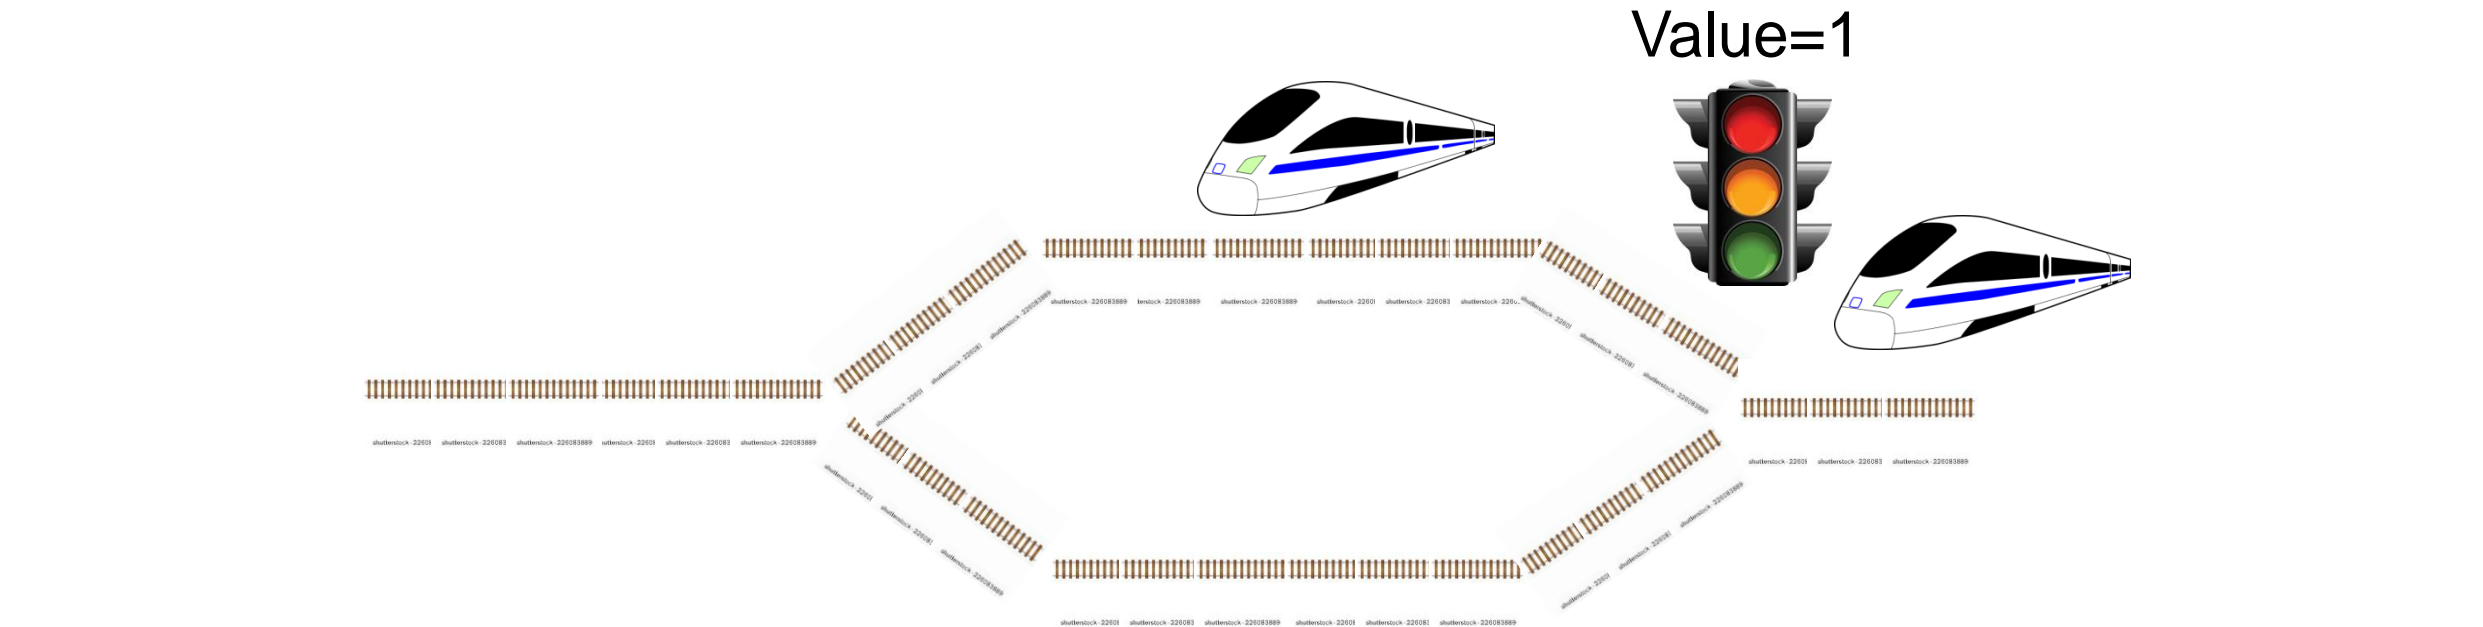
Railroad Analogy Step 4



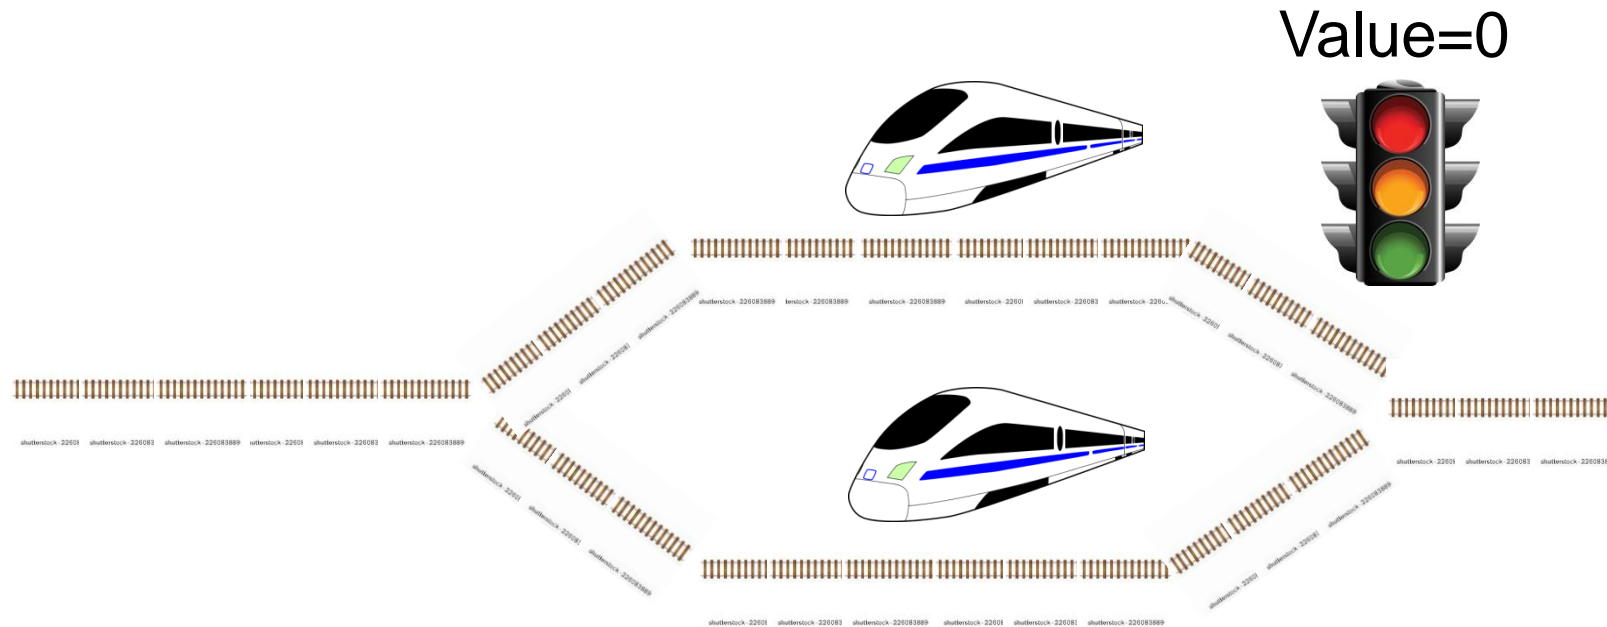


Railroad Analogy Step 6





Railroad Analogy Step 7



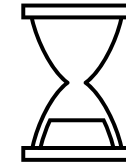
Two Uses of Semaphores

Mutual Exclusion



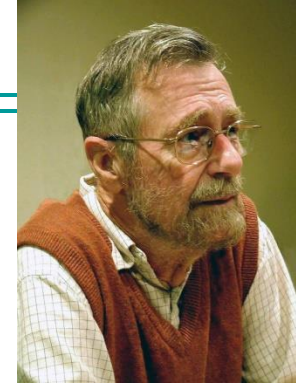
- Initial value = 1
- Also called “Binary Semaphore”.
- Can be used for mutual exclusion:
`semaphore.P();`
`// Critical section // goes`
`here`
`semaphore.V();`

Scheduling Constraints



- Initial value = 0
- What if you want a thread to wait for something?
- Example: Implement ThreadJoin (wait for a thread to terminate):
Initial value of semaphore = 0
`ThreadJoin {`
 `semaphore.P();`
`}`
`ThreadFinish {`
 `semaphore.V();`
`}`

Semaphores



- Semaphores are a kind of generalized lock
 - First defined by Dijkstra in late 60s
 - Main synchronization primitive used in original UNIX
- Definition: *A Semaphore has a non-negative integer value and supports the following two operations:*
 - $P()$: an atomic operation that waits for semaphore to become positive, then decrements it by 1
 - Think of this as the `wait()` operation
 - $V()$: an atomic operation that increments the semaphore by 1, waking up a waiting P , if any
 - Think of this as the `signal()` operation
 - Note that $P()$ stands for “*proberen*” (to test) and $V()$ stands for “*verhogen*” (to increment) in Dutch

<https://www.youtube.com/watch?v=LKQpy107yUY>

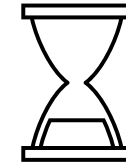
Two Uses of Semaphores

Mutual Exclusion



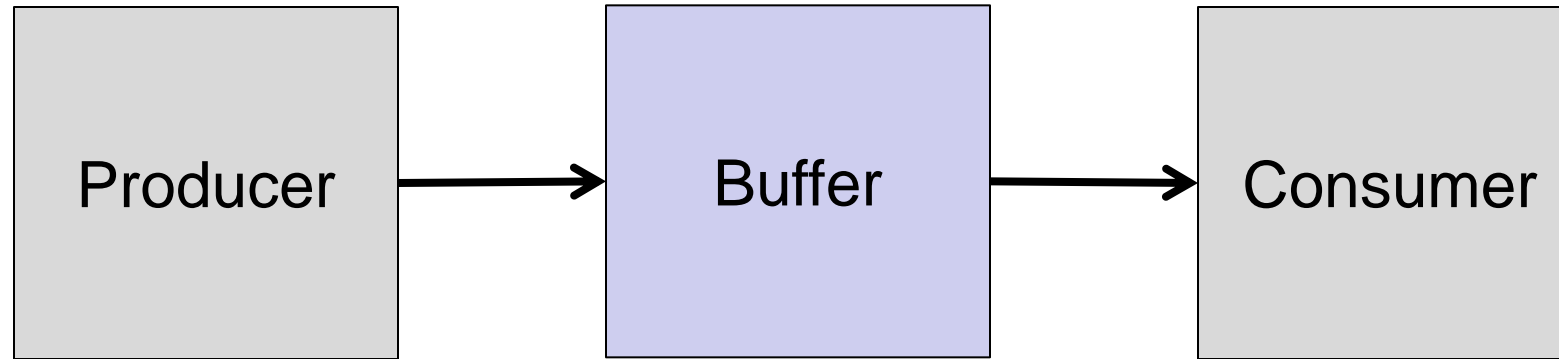
- Initial value = 1
- Also called “Binary Semaphore”.
- Can be used for mutual exclusion:
`semaphore.P();`
`// Critical section // goes`
`here`
`semaphore.V();`

Scheduling Constraints



- Initial value = 0
- What if you want a thread to wait for something?
- Example: Implement ThreadJoin (wait for a thread to terminate):
Initial value of semaphore = 0
`ThreadJoin {`
 `semaphore.P();`
`}`
`ThreadFinish {`
 `semaphore.V();`
`}`

A Bounded Buffer



Producer-consumer with a bounded buffer

Problem Definition

- Producer puts things into a **shared buffer**
- Consumer takes them out
- Need **synchronization** to coordinate producer/consumer

Don't want producer and consumer to have to work in **lockstep**, so put a **fixed-size buffer** between them

- Synchronize access to the buffer
- **Producer** needs to wait if buffer is **full**
- **Consumer** needs to wait if buffer is **empty**

Example: Drink machine

- Producer can put limited number of bottles in machine
- Consumer can't take bottles out if machine is empty

Example: GCC

- `cpp | cc1 | cc2 | as | ld`

Correctness constraints for solution

- Correctness Constraints:
 - Consumer must wait for producer to fill buffers, if all are empty (scheduling constraint)
 - Producer must wait for consumer to empty buffers, if all are full (scheduling constraint)
 - Only one thread can manipulate buffer queue at a time (mutual exclusion)
- Remember why we need mutual exclusion
 - Computers are stupid
 - Imagine if in real life: the delivery person is filling the machine and somebody comes up and tries to stick their money into the machine



Correctness constraints for solution

- General rule of thumb:

Use a separate semaphore for each constraint

- Semaphore fullBuffers; // consumer's constraint
- Semaphore emptyBuffers; // producer's constraint
- Semaphore mutex; // mutual exclusion

Full Solution to Bounded Buffer

```
Semaphore fullBuffer = 0;    // Initially, no pop
Semaphore emptyBuffers = numBuffers;
                                // Initially, num empty slots
Semaphore mutex = 1;         // No one using machine

Producer(item) {
    emptyBuffers.P();         // Wait until space
    mutex.P();               // Wait until buffer free
    Enqueue(item);
    mutex.V();
    fullBuffers.V();         // Tell consumers there is
                                // more pop
}

Consumer() {
    fullBuffers.P();         // Check if there's a pop
    mutex.P();               // Wait until machine free
    item = Dequeue();
    mutex.V();
    emptyBuffers.V();        // tell producer need more
    return item;
}
```

Discussion about Solution

- Why asymmetry?
 - Producer does: `emptyBuffer.P()`, `fullBuffer.V()`
 - Consumer does: `fullBuffer.P()`, `emptyBuffer.V()`
- Is order of P's important?
 - Yes! Can cause deadlock
- Is order of V's important?
 - No, except that it might affect scheduling efficiency
- What if we have 2 producers or 2 consumers?
 - Do we need to change anything?

So Far

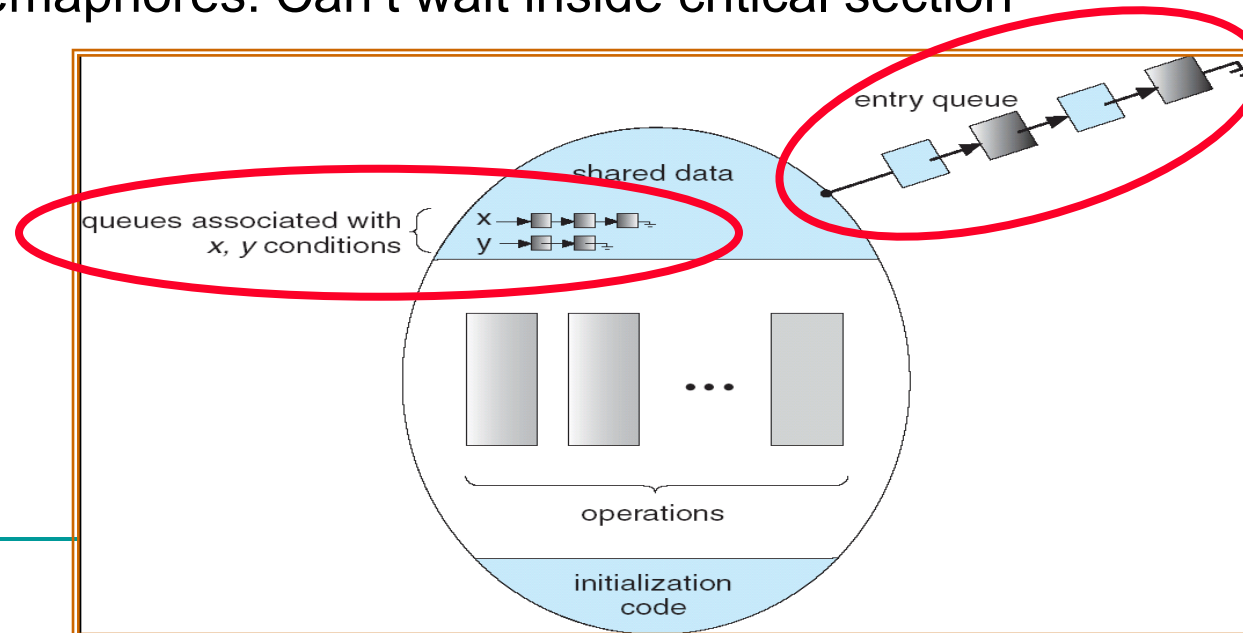
- Concurrency challenge
- Motivation for Synchronization and Locks
- Atomic Read-Modify-Write Operations
- Higher Level Synchronization Atoms
 - Semaphores
 - **Monitors**

Motivation for Monitors and Condition Variables

- Semaphores are a huge step up; just think of trying to do the bounded buffer with only loads and stores
 - Problem is that semaphores are dual purpose:
 - They are used for both mutex and scheduling constraints
 - Example: the fact that flipping of P's in bounded buffer gives deadlock is not immediately obvious. How do you prove correctness to someone?
- Cleaner idea: Use *locks* for mutual exclusion and *condition variables* for scheduling constraints
- Definition: Monitor: a lock and zero or more condition variables for managing concurrent access to shared data
 - Some languages like Java provide this natively
 - Most others use actual locks and condition variables

Monitor with Condition Variables

- Lock: the lock provides mutual exclusion to shared data
 - Always acquire before accessing shared data structure
 - Always release after finishing with shared data
 - Lock initially free
- Condition Variable: a queue of threads waiting for something *inside* a critical section
 - Key idea: make it possible to go to sleep inside critical section by atomically releasing lock at time we go to sleep
 - Contrast to semaphores: Can't wait inside critical section





Simple Monitor Example (version 1)

- Here is an (infinite) synchronized queue

```
Lock lock;  
Queue queue;
```

```
AddToQueue(item) {  
    lock.Acquire();           // Lock shared data  
    queue.enqueue(item);      // Add item  
    lock.Release();           // Release Lock  
}
```

```
RemoveFromQueue() {  
    lock.Acquire();           // Lock shared data  
    item = queue.dequeue();    // Get next item or null  
    lock.Release();           // Release Lock  
    return(item);              // Might return null  
}
```

- Not very interesting use of “Monitor”
 - It only uses a lock with no condition variables
 - Cannot put consumer to sleep if no work!

Condition Variables

- How do we change the RemoveFromQueue() routine to wait until something is on the queue?
 - Could do this by keeping a count of the number of things on the queue (with semaphores), but error prone
- Condition Variable: a queue of threads waiting for something *inside* a critical section
 - Key idea: allow sleeping inside critical section by atomically releasing lock at time we go to sleep
 - Contrast to semaphores: Can't wait inside critical section

Condition Variables

- Operations:
 - `wait(&lock)`: Atomically release lock and go to sleep. Reacquire lock later, before returning.
 - `Signal()`: Wake up one waiter, if any
 - `Broadcast()`: Wake up all waiters
- Rule: Must hold lock when doing condition variable ops!

Complete Monitor Example (with condition variable)

- Here is an (infinite) synchronized queue

```
Lock lock;
Condition dataready;
Queue queue;
AddToQueue(item) {
    lock.Acquire();           // Get Lock
    queue.enqueue(item);      // Add item
    dataready.signal();       // Signal any waiters
    lock.Release();           // Release Lock
}
RemoveFromQueue() {
    lock.Acquire();           // Get Lock
    while (queue.isEmpty()) {
        dataready.wait(&lock); // If nothing, sleep
    }
    item = queue.dequeue();    // Get next item
    lock.Release();           // Release Lock
    return(item);
}
```

Mesa vs. Hoare monitors

- Need to be careful about precise definition of signal and wait. Consider a piece of our dequeue code:

```
while (queue.isEmpty()) {  
    dataready.wait(&lock); // If nothing, sleep  
}  
item = queue.dequeue(); // Get next item
```

- Why didn't we do this?

```
if (queue.isEmpty()) {  
    dataready.wait(&lock); // If nothing, sleep  
}  
item = queue.dequeue(); // Get next item
```

Mesa vs. Hoare scheduling

Hoare-style (textbooks):

- Signaler gives lock and CPU to waiter
- Waiter runs immediately
- Waiter gives lock and CPU back to signaler when it exits critical section or waits again

Mesa-style (most real OS):

- Signaler keeps lock and processor
- Waiter placed on ready queue with no special priority
- Practically, need to check condition again after wait

Conclusion

- Concurrency challenge
- Motivation for Synchronization and Locks
- Atomic Read-Modify-Write Operations
- Higher Level Synchronization Atoms
 - Semaphores
 - Monitors