

---

---

# Barriers, Readers/Writers, Scheduling Intro

26 December 2024  
Lecture 8

Slides adapted from John Kubiawicz (UC Berkeley)

# Concept Review

---

Atomic  
Reads and  
Writes

Starvation

Race  
condition

Mutual  
exclusion

Critical  
section

Lock

- Spin lock
- In-Kernel lock

Busy waiting

Semaphores

Condition  
Variables

Monitors

# Topics for Today


---

- Higher Level Synchronization Atoms
  - Barrier Synchronization
  - Example: Readers and Writers
- Mutual Exclusion
  - Mutual Exclusion in High Level Language
- Scheduling
  - FIFO
  - Round Robin




# Barrier Synchronization - 3

---




# Barrier Synchronization - 3

---




# Barrier Synchronization - 3

---



# Barrier Synchronization - 3


---






# Barrier Synchronization - 3

---




# Barrier Synchronization - 3

---





# Barrier Synchronization - 3

---




# Reusable Barrier Synchronization - 3

---




# Reusable Barrier Synchronization - 3

---




# Reusable Barrier Synchronization - 3

---




# Reusable Barrier Synchronization - 3

---



# Reusable Barrier Synchronization - 3


---






# Reusable Barrier Synchronization - 3

---




# Reusable Barrier Synchronization - 3

---




# Reusable Barrier Synchronization - 3

---



# Reusable Barrier Synchronization - 3

---




# So Far

---

- Higher Level Synchronization Atoms
  - Barrier Synchronization
  - Example: Readers and Writers
- Mutual Exclusion
  - Mutual Exclusion in High Level Language
- Scheduling
  - FIFO
  - Round Robin

# Extended example: Readers/Writers Problem




Motivation: Consider a shared database

- Two classes of users:
  - Readers – never modify database
  - Writers – read and modify database
- Is using a single lock on the whole database sufficient?
  - Allow many readers at the same time
  - Only one writer at a time

# Basic Readers/Writers Solution

---


- Correctness Constraints:
  - Readers can access database when no writers
  - Writers can access database when no readers or writers
  - Only one thread manipulates state variables at a time
- Basic structure of a solution:
  - Reader()
    - Wait until no writers
    - Access data base
    - Check out – wake up a waiting writer
  - Writer()
    - Wait until no active readers or writers
    - Access database
    - Check out – wake up waiting readers or writer



# Basic Readers/Writers Solution

---

- State variables (Protected by a lock called “lock”):
  - int AR: Number of active readers; initially = 0
  - int WR: Number of waiting readers; initially = 0
  - int AW: Number of active writers; initially = 0
  - int WW: Number of waiting writers; initially = 0
  - Condition okToRead = NIL
  - Condition okToWrite = NIL





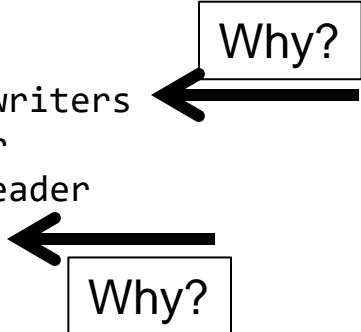
# Code for a Reader

```
Reader() {  
    // First check self into system  
    lock.Acquire();  
    while ((AW + WW) > 0) {        // Is it safe to read?  
        WR++;                     // No. Writers exist  
        okToRead.wait(&lock);     // Sleep on cond var  
        WR--;                     // No longer waiting  
    }  
    AR++;                         // Now we are active!  
    lock.release(); ←  
    // Perform actual read-only access  
    AccessDatabase(ReadOnly);  
    // Now, check out of system  
    lock.Acquire();  
    AR--;                         // No longer active  
    if (AR == 0 && WW > 0)        // No other active readers  
        okToWrite.signal();      // Wake up one writer  
    lock.Release();  
}
```

Why?

# Code for a Writer

```
Writer() {  
    // First check self into system  
    lock.Acquire();  
    while ((AW + AR) > 0) {                // Is it safe to write?  
        WW++; // No. Active users exist  
        okToWrite.wait(&lock);            // Sleep on cond var  
        WW--; // No longer waiting  
    }  
    AW++; // Now we are active!  
    lock.release();  
    // Perform actual read/write access  
    AccessDatabase(ReadWrite);  
    // Now, check out of system  
    lock.Acquire();  
    AW--; // No longer active  
    if (WW > 0){                          // Give priority to writers  
        okToWrite.signal();               // Wake up one writer  
    } else if (WR > 0) {                   // Otherwise, wake reader  
        okToRead.broadcast();             // Wake all readers  
    }  
    lock.Release();  
}
```



# Simulation R/W Step 1

---

- Consider the following sequence of operators:
  - R1, R2, W1, R3
- On entry, each reader checks the following:

---

```
while ((AW + WW) > 0) {           // Is it safe to read?
    WR++;                          // No. Writers exist
    okToRead.wait(&lock);          // Sleep on cond var
    WR--;                          // No longer waiting
}
AR++;                             // Now we are active!
```

---

- First, R1 comes along:  
AR = 1, WR = 0, AW = 0, WW = 0
- Second, R2 comes along:  
AR = 2, WR = 0, AW = 0, WW = 0
- Now, readers make take a while to access database
  - Situation: Locks released
  - Only AR is non-zero

# Simulation R/W Step 2

---

- Next, W1 comes along:

```
while ((AW + AR) > 0) {    // Is it safe to write?
    WW++;                 // No. Active users exist
    okToWrite.wait(&lock); // Sleep on cond var
    WW--;                 // No longer waiting
}
AW++;
```

- Can't start because of readers, so go to sleep:

AR = 2, WR = 0, AW = 0, WW = 1

- Finally, R3 comes along:

AR = 2, WR = 1, AW = 0, WW = 1

- Now, say that R2 finishes before R1:

AR = 1, WR = 1, AW = 0, WW = 1

- Finally, last of first two readers (R1) finishes and wakes up a writer:

```
if (AR == 0 && WW > 0)    // No other active readers
    okToWrite.signal();    // Wake up one writer
```

# Simulation R/W Step 3

---

- When the writer wakes up, get:

$AR = 0, WR = 1, AW = 1, WW = 0$

- Then, when writer finishes:

```
if (WW > 0){                // Give priority to writers
    okToWrite.signal();      // Wake up one writer
} else if (WR > 0) {         // Otherwise, wake reader
    okToRead.broadcast();    // Wake all readers
}
```

- Writer wakes up reader, so get:

$AR = 1, WR = 0, AW = 0, WW = 0$

- When reader completes, we are finished

# Questions about R/W

- Can readers starve? Consider Reader() entry code:

```
while ((AW + WW) > 0) {           // Is it safe to read?
    WR++;                          // No. Writers exist
    okToRead.wait(&lock);          // Sleep on cond var
    WR--;                          // No longer waiting
}
AR++;                             // Now we are active!
```

- What if we erase the condition check in Reader exit?

```
AR--;                             // No longer active
 if (AR == 0 && WW > 0)           // No other active readers
    okToWrite.signal();           // Wake up one writer
```

- Further, what if we turn the signal() into broadcast()

```
AR--;                             // No longer active
okToWrite.broadcast();            // Wake up one writer
```

- Finally, what if we use only one condition variable (call it “okToContinue”) instead of two separate ones?

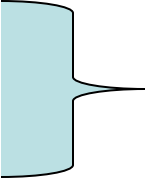
- Both readers and writers sleep on this variable
- Must use broadcast() instead of signal()

# Monitors Conclusion

---

- Monitors represent the logic of the program
  - Wait if necessary
  - Signal when you change something so any waiting threads can proceed
- Basic structure of monitor-based program:

```
lock
while (need to wait) {
    condvar.wait();
}
unlock
```



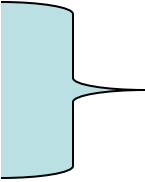
Check or update state variables.  
Wait if necessary

do something so no need to wait

```
lock

condvar.signal();

unlock
```



Check or update state variables.

# So Far

---

- Higher Level Synchronization Atoms
  - Barrier Synchronization
  - Example: Readers and Writers
- Mutual Exclusion
  - Mutual Exclusion in High Level Language
- Scheduling
  - FIFO
  - Round Robin



# C-Language Support for Synchronization


- C language: Straightforward synchronization

- Just make sure you know *all* the code paths out of a critical section

```
int Rtn() {  
    lock.acquire();  
    ...  
    if (exception) {  
        lock.release();  
        return errReturnCode;  
    }  
    ...  
    lock.release();  
    return OK;  
}
```

- Watch out for setjmp/longjmp!

- Can cause a non-local jump out of procedure
- In example, procedure E calls longjmp, popping stack back to procedure B
- If Procedure C had lock.acquire, problem!



# C++ Language Support for Synchronization

---

- Languages with exceptions like C++
  - Languages that support exceptions are problematic (easy to make a non-local exit without releasing lock)
  - Consider:

```
void Rtn() {  
    lock.acquire();  
    ...  
    DoFoo();  
    ...  
    lock.release();  
}  
void DoFoo() {  
    ...  
    if (exception) throw errException;  
    ...  
}
```
  - Notice that an exception in DoFoo() will exit without releasing the lock!

# C++ Language Support for Synchronization

---

- Must catch all exceptions in critical sections
  - Catch exceptions, release lock, and re-throw exception:

```
void Rtn() {  
    lock.acquire();  
    try {  
        ...  
        DoFoo();  
        ...  
    } catch (...) {           // catch exception  
        lock.release();       // release lock  
        throw;                 // re-throw the exception  
    }  
    lock.release();  
}  
void DoFoo() {  
    ...  
    if (exception) throw errException;  
    ...  
}
```
  - Even Better: `auto_ptr<T>` facility. See C++ Spec.
    - Can deallocate/free lock regardless of exit method

# Java Language Support for Synchronization

---

- Java has explicit support for threads and thread synchronization

- Bank Account example:

```
class Account {  
    private int balance;  
    // object constructor  
    public Account (int initialBalance) {  
        balance = initialBalance;  
    }  
    public synchronized int getBalance() {  
        return balance;  
    }  
    public synchronized void deposit(int amount) {  
        balance += amount;  
    }  
}
```

- Every object has an associated lock which gets automatically acquired and released on entry and exit from a *synchronized* method.

# Java Language Support for Synchronization

---

Java also has *synchronized* blocks:

```
int i, j;
void foo() {
    Object locker = new Object();
    synchronized (locker) {
        i += j;
    }
}
```

- Since every Java object has one associated lock, the statement acquires and releases the object's lock on entry and exit of the block
- Problem is that the code here doesn't protect anything. Why?

# Java Language Support for Synchronization

---

A better form of the code:

```
Object locker = new Object();  
int i, j;  
void foo() {  
    synchronized (locker) {  
        i += j;  
    }  
}
```

- Now all threads will use the same lock and we'll get some mutual exclusion.

# Java Language Support for Synchronization

---

- Works properly even with exceptions:

```
synchronized (locker) {  
    ...  
    DoFoo();  
    ...  
}  
void DoFoo() {  
    throw errException;  
}
```

- Lock is released when the exception is thrown.

# Java Language Support for Synchronization

---

- Every object also has one condition variable associated with it
  - How to wait inside a synchronization method or block:
    - `void wait(long timeout);` // Wait for timeout
    - `void wait(long timeout, int nanoseconds);` //variant
    - `void wait();`
  - How to signal in a synchronized method or block:
    - `void notify();` // wakes up oldest waiter
    - `void notifyAll();` // like broadcast, wakes everyone



# Java Language Support for Synchronization

---

- Condition variables can wait for a bounded length of time. This is useful for handling exception cases:

```
t1 = time.now();
while (!ATMRequest()) {
    wait (CHECKPERIOD);
    t2 = time.now();
    if (t2 - t1 > LONG_TIME) checkMachine();
}
```

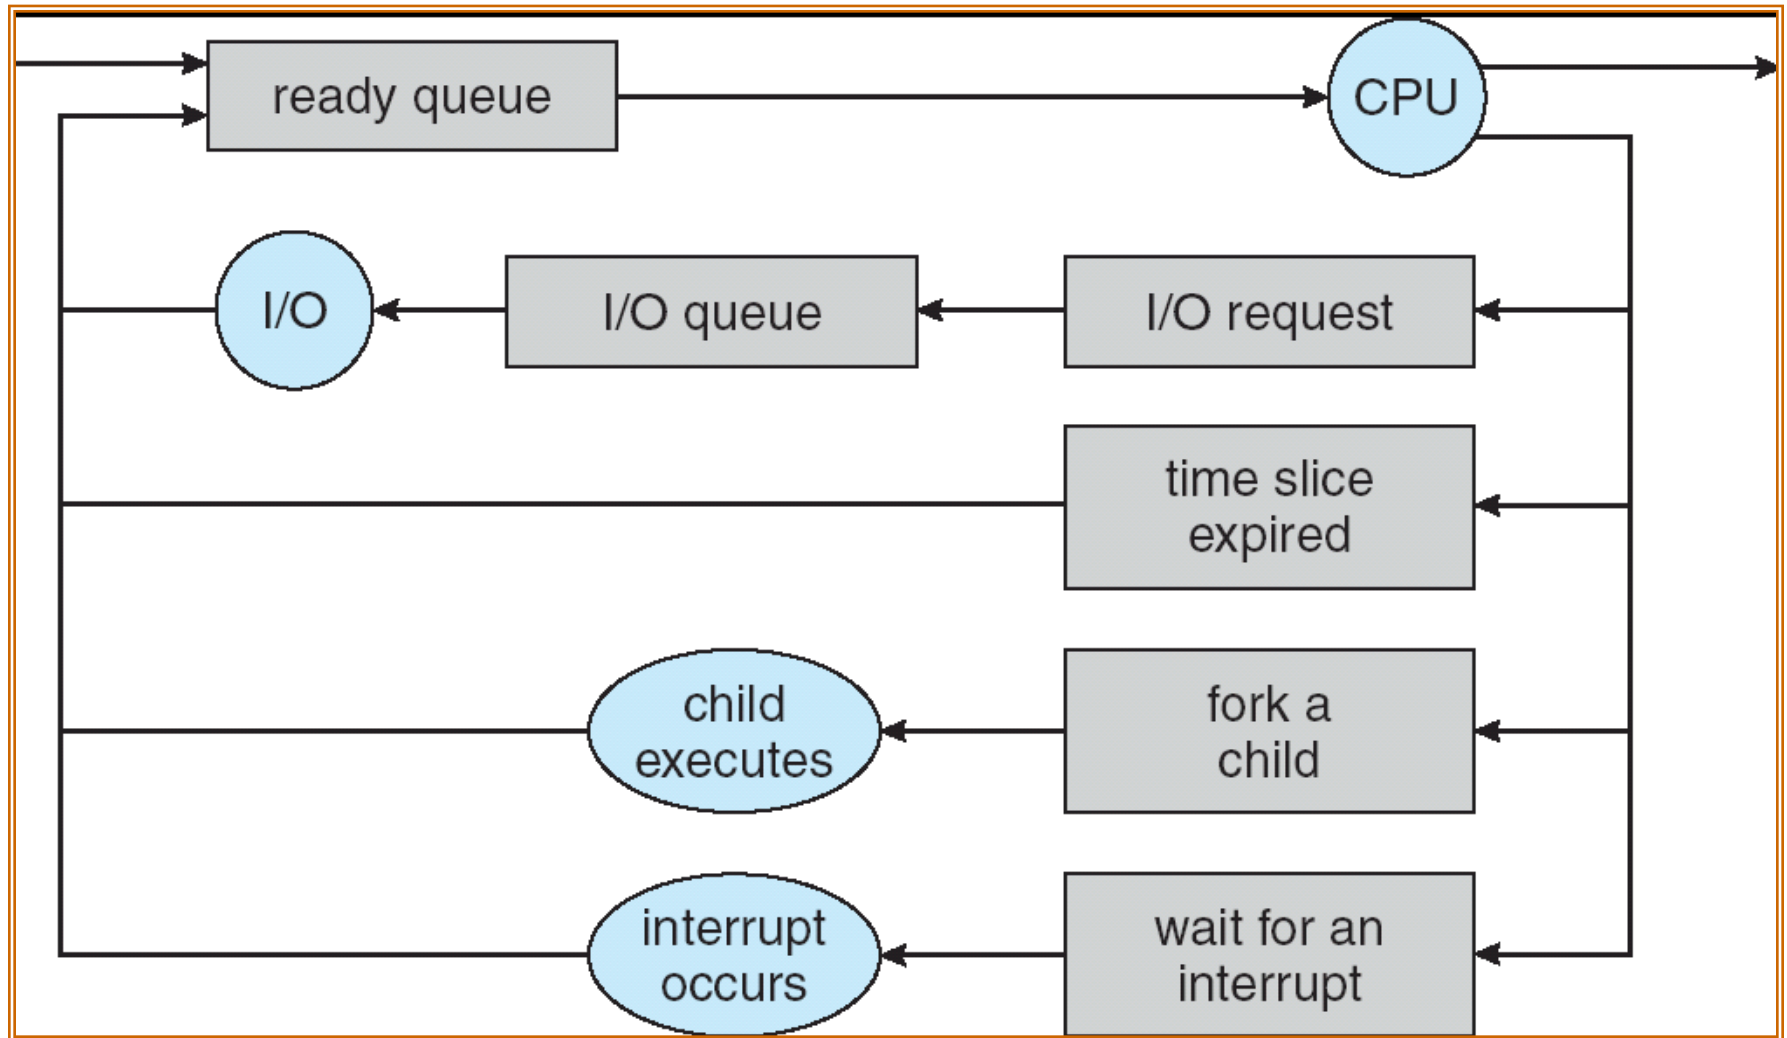
- Not all Java VMs equivalent!
  - Different scheduling policies, not necessarily preemptive!

# So Far

---


- Higher Level Synchronization Atoms
  - Barrier Synchronization
  - Example: Readers and Writers
- Mutual Exclusion
  - Mutual Exclusion in High Level Language
- Scheduling
  - FIFO
  - Round Robin

# Recall: CPU Scheduling



# Scheduling Assumptions (review)

- CPU scheduling big area of research in early 70's
- Many implicit assumptions for CPU scheduling:
  - One program per user
  - One thread per program
  - Programs are independent
- Clearly, these are **unrealistic** but they simplify the problem so it can be solved
  - For instance: is “**fair**” about fairness among users or programs?
    - If I run one compilation job and you run five, you get five times as much CPU on many operating systems
- **The high-level goal**: Dole out CPU time to optimize some desired parameters of system



# A snapshot

Processes	
^	
Name	Status
Apps (14)	
> Firefox (35)	
> Mail	
> Microsoft PowerPoint	
> Microsoft Word (2)	
> Notepad++	
> SumatraPDF	
> SumatraPDF	
> Task Manager	
> Terminal (3)	
> Thunderbird (5)	
> Toggl Track	
> WhatsApp (2)	
> Windows Explorer (2)	
> WinEdt 10.3 (2)	

## Background processes (117)

- > Adobe Acrobat Update Servic...
- Application Frame Host
- > Background Task Host (6)
- COM Surrogate
- COM Surrogate
- COM Surrogate
- crashpad\_handler
- crashpad\_handler
- crashpad\_handler
- crashpad\_handler
- CTF Loader
- DAX API
- > DAX API
- Device Association Framework...
- > Elan Service


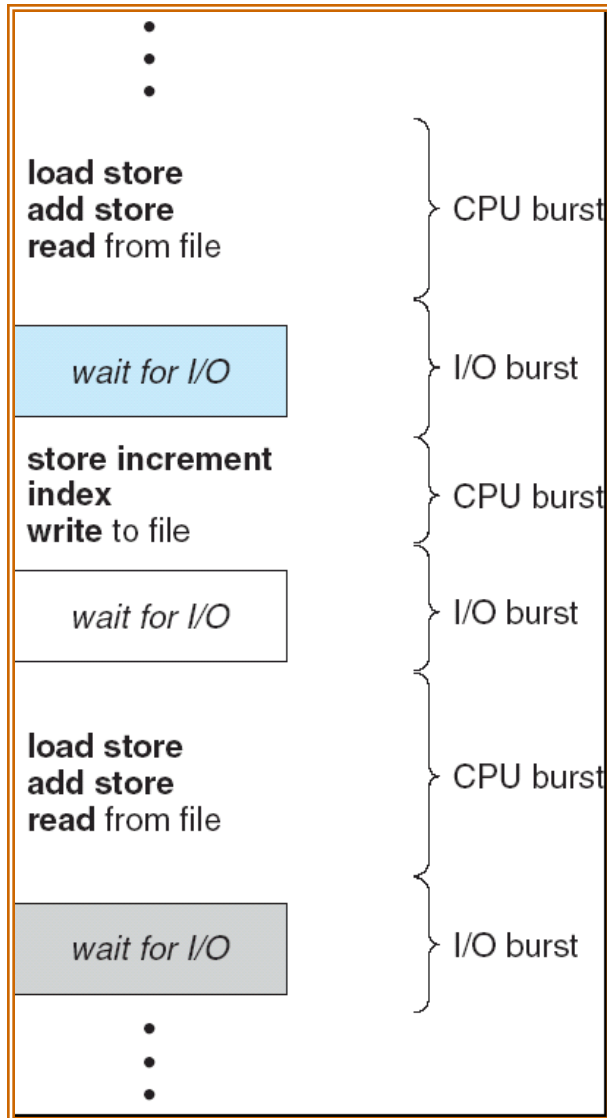
Then I click on



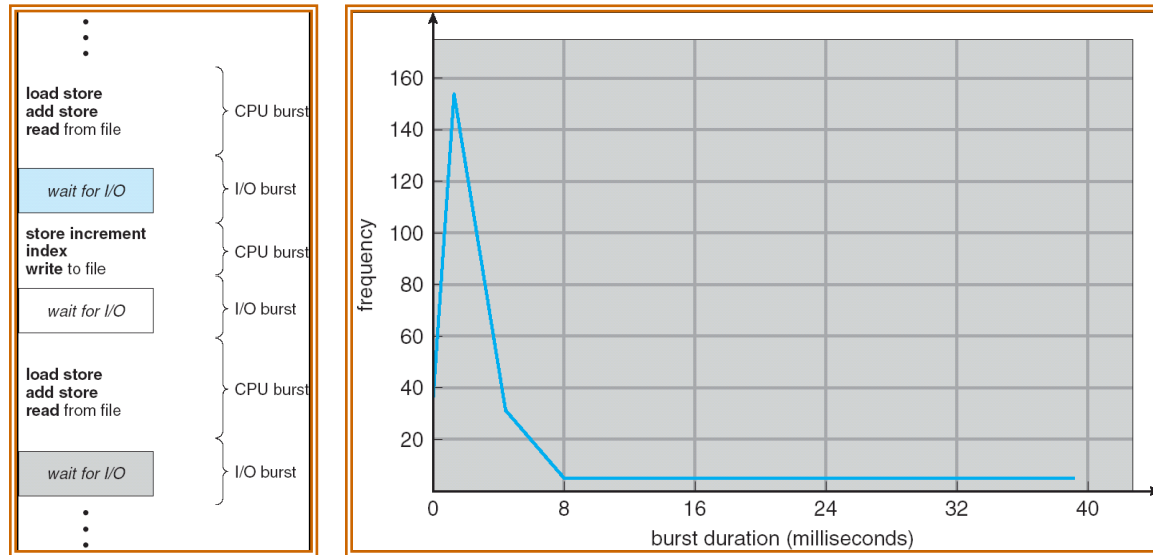
How much CPU resources does each process get?

The new one?

# Assumption: CPU Bursts



# Assumption: CPU Bursts




- Execution model: programs alternate between bursts of CPU and I/O
  - Program typically uses the CPU for some period of time, then does I/O, then uses CPU again
  - Each scheduling decision is about which job to give to the CPU for use by its next CPU burst
  - With time slicing, thread may be forced to give up CPU before finishing current CPU burst

# First-Come, First-Served (FCFS) Scheduling

- First-Come, First-Served (FCFS) or FIFO or “Run until done”
  - Used to mean one program scheduled until done (including I/O)
  - Now, means keep CPU until thread blocks


- Example:

Arrival Order ↓	Process	Burst Time
	$P_1$	24
	$P_2$	3
	$P_3$	3




- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
  - Average waiting time:  $(0 + 24 + 27)/3 = 17$
  - Average Completion time:  $(24 + 27 + 30)/3 = 27$
- *Convoy effect*: short process behind long process







רגע, אני רק  
שאלה...



# First-Come, First-Served (FCFS) Scheduling

- Suppose that processes arrive in order:  $P_2, P_3, P_1$   
Schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- Average Completion time:  $(3 + 6 + 30)/3 = 13$
- In **second** case:
  - Average waiting time is much better (before it was 17)
  - Average completion time is better (before it was 27)
- FIFO Pros and Cons:
  - Simple (👍)
  - Short jobs get stuck behind long ones (👎)
    - Rami Levy: Getting humus, always stuck behind cart full of small items.  
Upside: get to catch up on email and Facebook!

# Round Robin (RR)

---

- FCFS Scheme: **Potentially bad** for short jobs!
  - Depends on **submit order**
  - If you are first in line at supermarket with humus, you don't care who is behind you, on the other hand...
- **Round Robin Scheme**
  - Each process gets a small unit of CPU time (**time quantum**), usually 10-100 **milliseconds**
  - After quantum expires, the process is **preempted** and added to the end of the ready queue.
  - $n$  processes in ready queue and time quantum is  $q \Rightarrow$ 
    - Each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units
    - **No process waits more than  $(n - 1)q$  time units**


# Round Robin (RR)

---

- Performance
  - $q$  large  $\Rightarrow$  FCFS
  - $q$  small  $\Rightarrow$  Interleaved (really small  $\Rightarrow$  hyperthreading?)
  - $q$  must be **large** with respect to context switch, otherwise **overhead is too high** (all overhead)

# Round Robin Example

Process	Burst Time
$P_1$	53
$P_2$	8
$P_3$	68
$P_4$	24



- Waiting times:
  - $P_1 = (68 - 20) + (112 - 88) = 72$
  - $P_2 = (20 - 0) = 20$
  - $P_3 = (28 - 0) + (88 - 48) + (125 - 108) = 85$
  - $P_4 = (48 - 0) + (108 - 68) = 88$
- Average waiting time =  $(72 + 20 + 85 + 88)/4 = 66.25$
- Average completion time =  $\frac{125+28+153+112}{4} = 104.5$

# Round Robin

---

- Better for short jobs, Fair (👍)
- Context-switching time adds up for long jobs (👎)

# Round-Robin Discussion

---

- How do you **choose** time slice?
  - What if **too big**? Response time suffers
  - What if **infinite** ( $\infty$ )? Same as FCFS
  - What if time slice **too small**? Throughput suffers!
- Actual choices of time slice:
  - Initially, UNIX time slice was **one second**:
    - Worked ok when UNIX was used by **one or two people**.
    - What if **three compilations** going on? 3 seconds to echo a keystroke!
  - In practice, Need to balance **short-job performance** and **long-job throughput**:
    - Typical time slice today is between **10ms – 100ms**
    - Typical context-switching overhead is **0.1ms – 1ms**
    - Roughly **1%** overhead due to context-switching

# Comparing FCFS and RR

- Assuming zero-cost context-switching time, is RR always better than FCFS?
- Simple example:
  - 10 jobs, each take 100s of CPU time
  - RR scheduler quantum of 1s
  - All jobs arrive at the same time
- Completion Times:

Job #	FIFO	RR
1	100	991
2	200	992
...	...	...
9	900	999
10	1000	1000




# Comparing FCFS and RR

Job #	FIFO	RR
1	100	991
2	200	992
...	...	...
9	900	999
10	1000	1000

- Both RR and FCFS finish at the same time
- Average response time is much worse under RR!
  - Bad when all jobs same length
- Also: Cache state must be shared between all jobs with RR but can be devoted to each job with FIFO
  - Total time for RR longer even for zero-cost switch!

# Earlier Example with Different Time Quanta

Best FCFS:




	Quantum	$P_1$	$P_2$	$P_3$	$P_4$	AVG
Wait Time	Best FCFS	32	0	85	8	31.25
	$Q = 1$	84	22	85	57	62
	$Q = 5$	82	20	85	58	61.25
	$Q = 8$	80	8	85	56	57.25
	$Q = 10$	82	10	85	68	61.25
	$Q = 20$	72	20	85	88	66.25
	Worst FCFS	68	145	0	121	83.5

	Quantum	$P_1$	$P_2$	$P_3$	$P_4$	AVG
Completion Time	Best FCFS	85	8	153	32	69.5
	$Q = 1$	137	30	153	81	100.25
	$Q = 5$	135	28	153	82	99.5
	$Q = 8$	133	16	153	80	95.5
	$Q = 10$	135	18	153	92	99.5
	$Q = 20$	125	28	153	112	104.5
	Worst FCFS	121	153	68	145	121.75

# RR Quanta Graphed

Average Completion and Wait Times using Round Robin  
for sample task set



# Conclusion

---

- Higher Level Synchronization Atoms
  - Barrier Synchronization
  - Example: Readers and Writers
- Mutual Exclusion
  - Mutual Exclusion in High Level Language
- Scheduling
  - FIFO
  - Round Robin