

---

---

# ARQ, Sliding Window,

27 November 2024  
Lecture 4


# Topics for Today

---

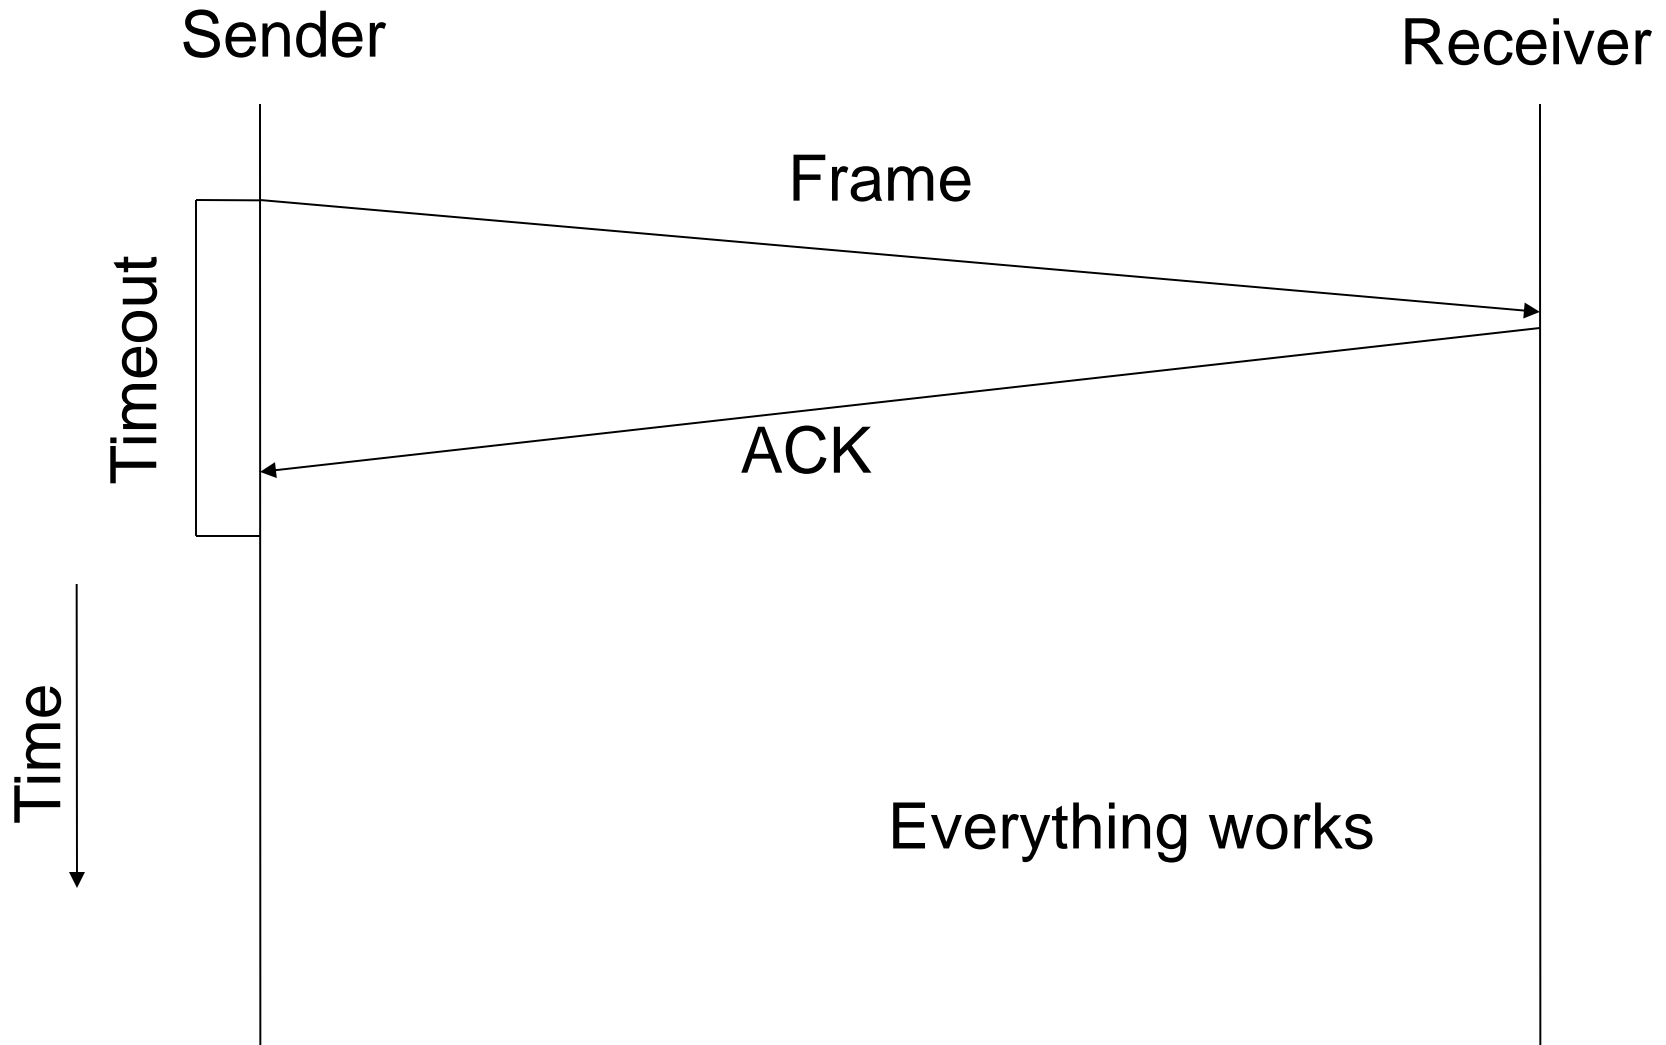
- ARQ
  - Stop and Wait
  - Sliding Window
- Ethernet
  - Topologies
  - Frame format
  - Media Access Control
  - Efficiency
- Source: Peterson and Davie 2.1-2.5, 2.6, Tanenbaum 4.3

# Stop-and-Wait - Simplest scheme

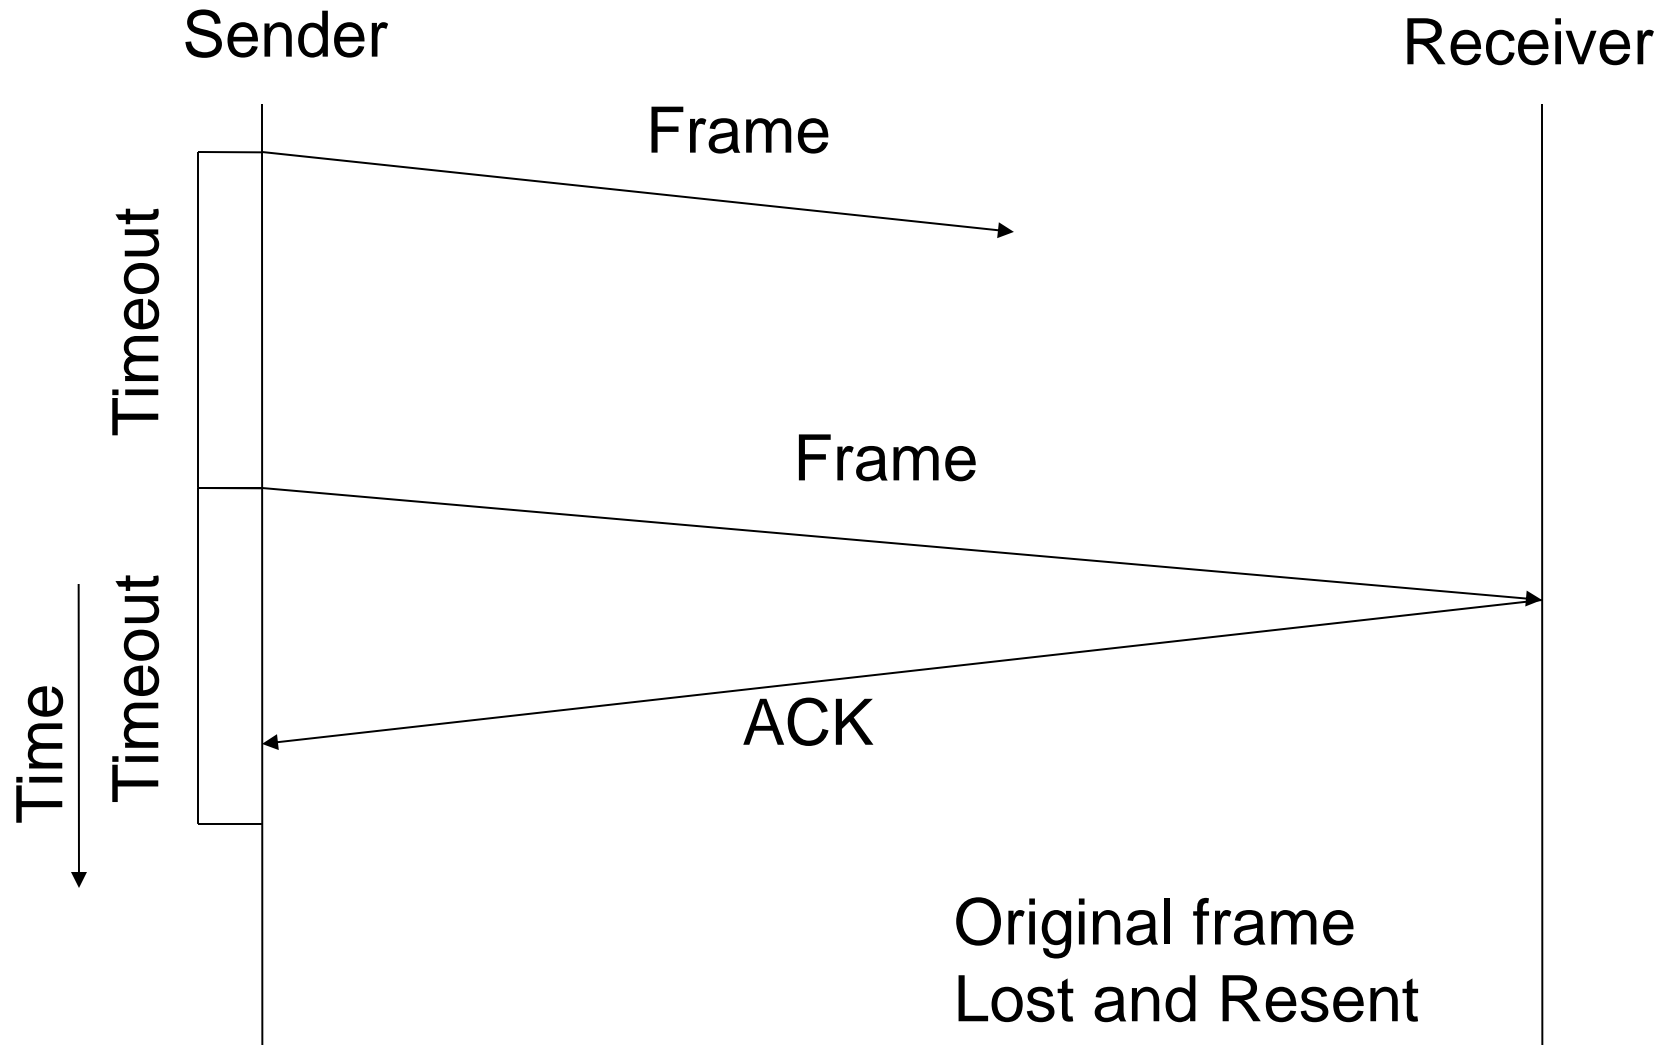
1. After transmitting one frame, sender waits for an ACK
2. If the ACK doesn't arrive, sender retransmits




# Stop-and-Wait scenario 1



# Stop-and-Wait scenario 2

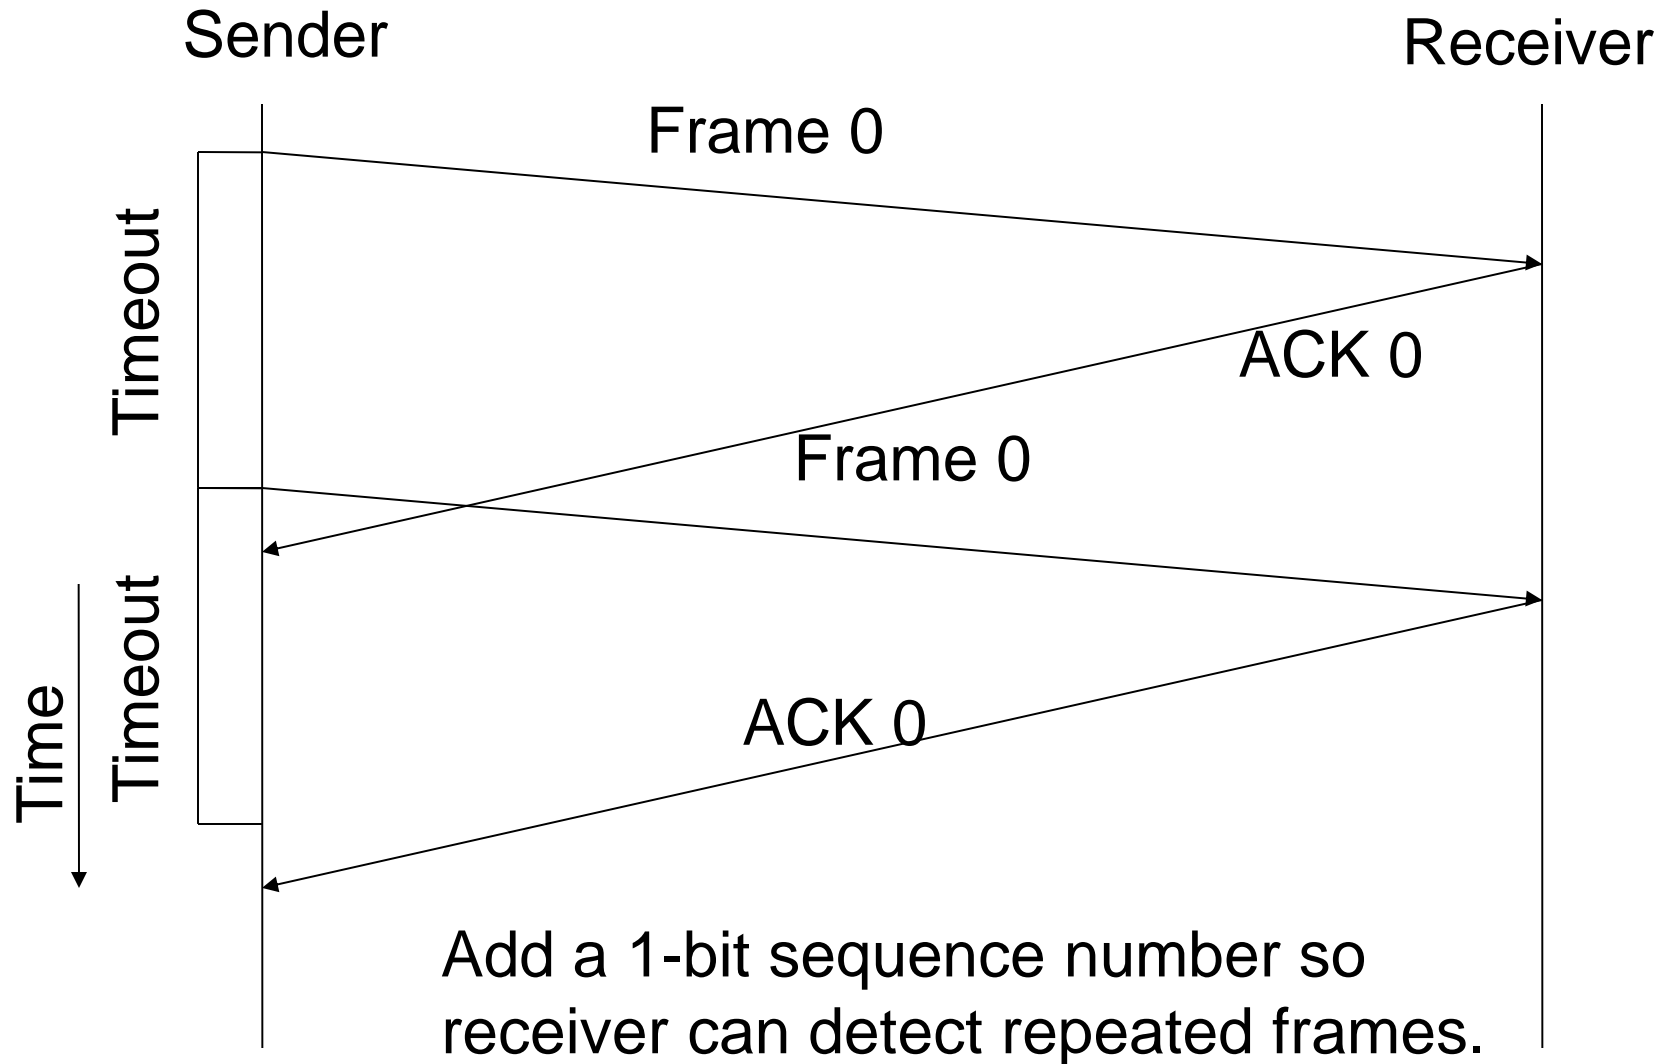


# Stop-and-Wait scenario 3





# Sequence numbers



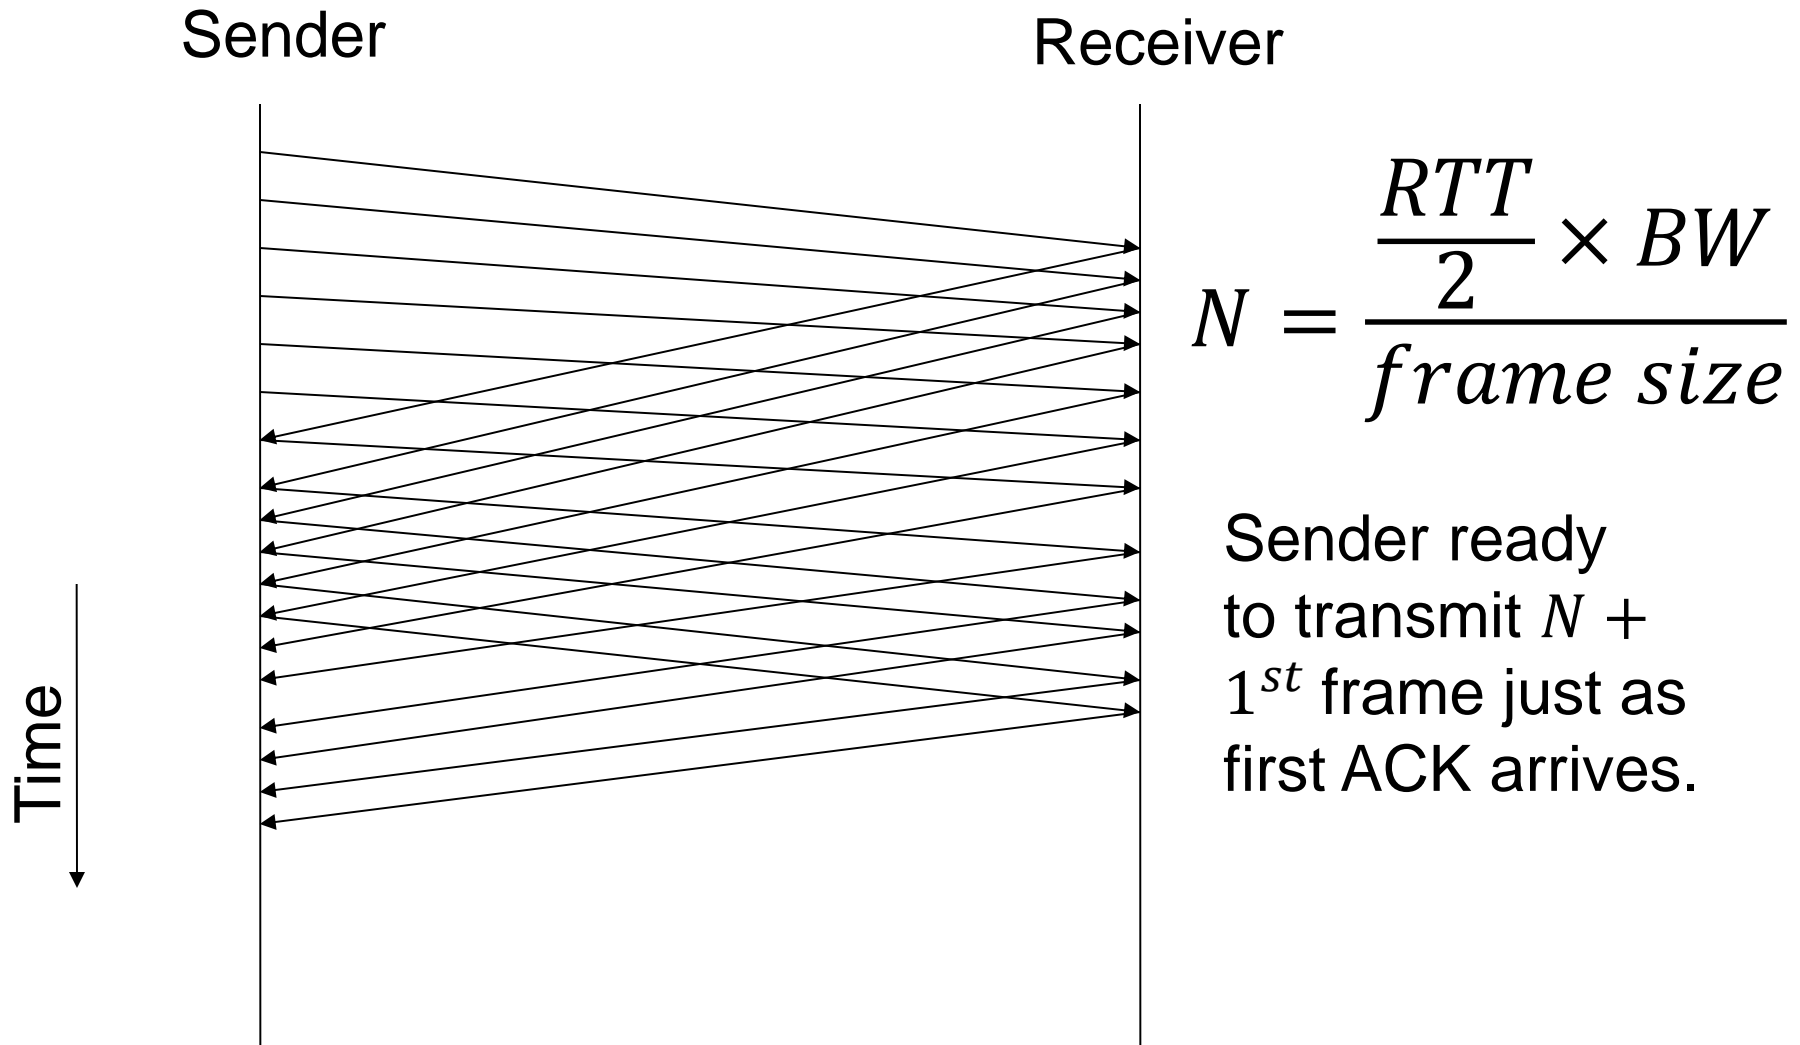


# Stop-and-Wait

---

- Inefficient use of link's capacity
- Sends 1 frame per RTT
- Example:
  - 10Mbps Link
  - 16ms RTT
  - Delay x Bandwidth product is about 20KB
  - Frame size of 1K yields about 5% link capacity

# More efficient solution



# So Far

---


- ARQ
  - Stop and Wait
  - Sliding Window

# Sliding Window Algorithm

- Sender assigns a *sequence number* to each frame:  
*SeqNum*
  - For now, assume *SeqNum* can grow infinitely
- Send Window Size (*SWS*)
  - Upper bound on # of unacknowledged frames sender will transmit
- Last ACK Received (*LAR*)
  - Sequence number of last ACK
- Last Frame Sent (*LFS*)




# Sender Invariant



- $LFS - LAR \leq SWS$
- Associates timeout with each frame sent
  - Retransmits if no ACK received before timeout
- When ACK arrives, increase  $LAR$ 
  - Means another frame can be sent

# Receiver Invariant



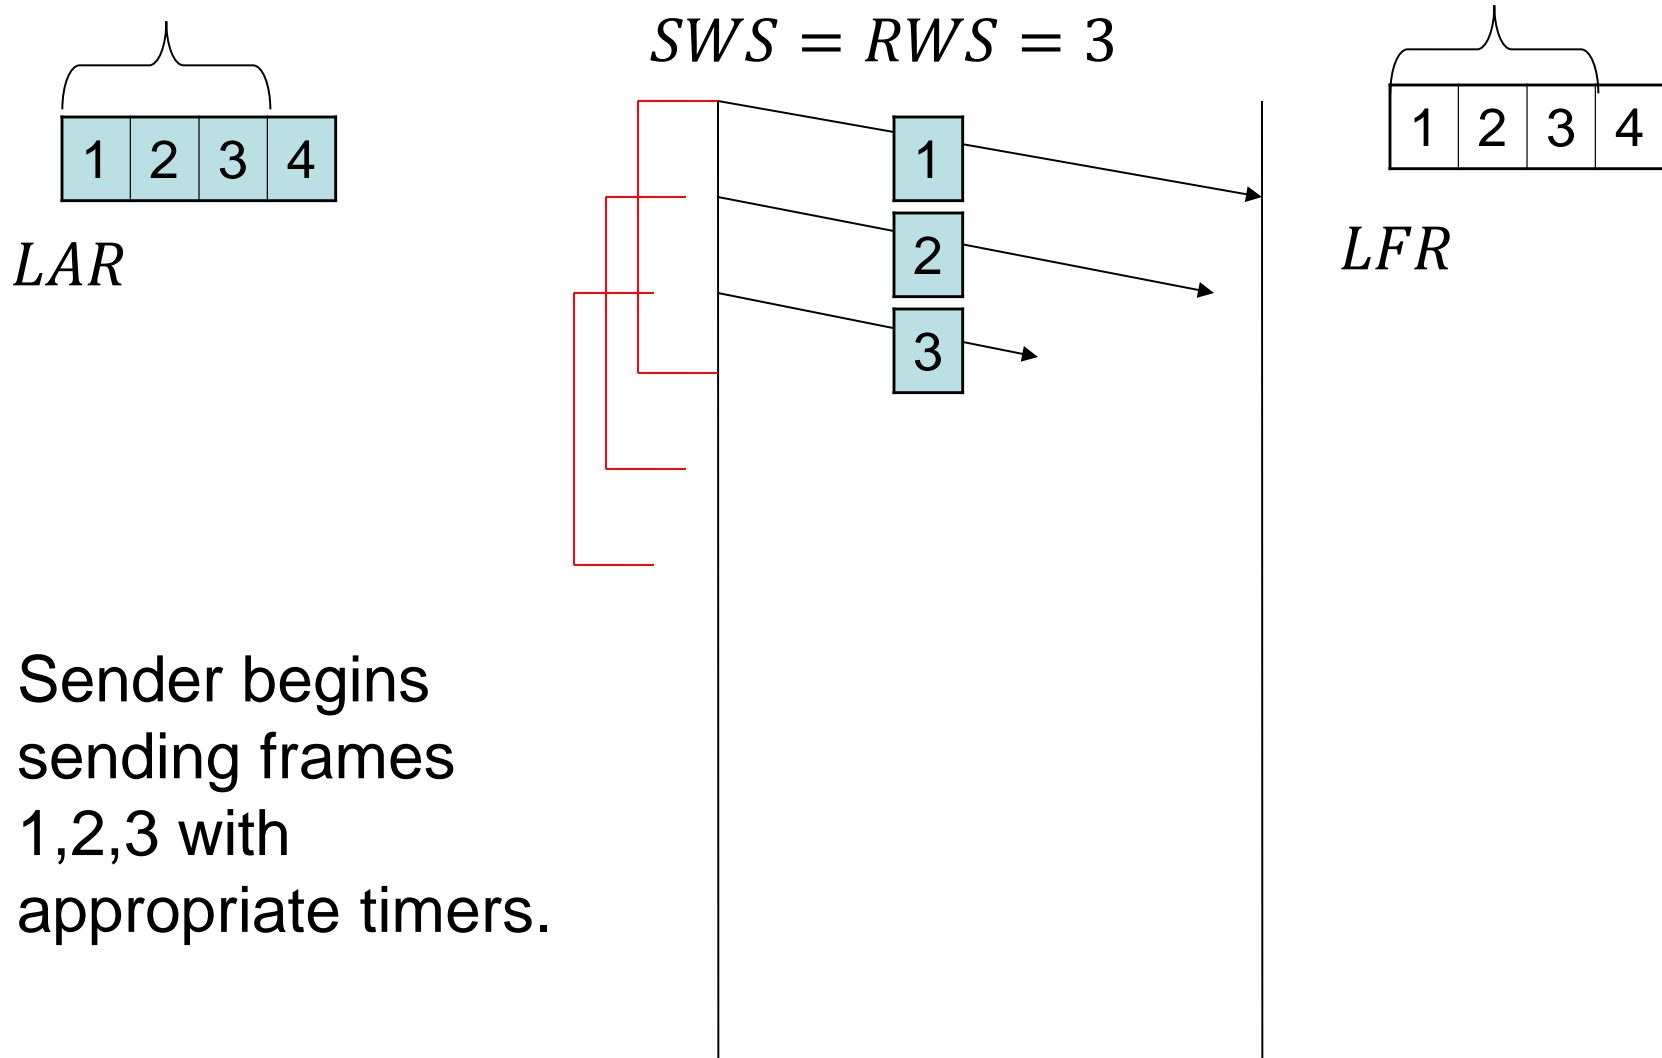
- Receive Window Size ( $RWS$ )
  - Number of out-of-order frames it will accept
- Largest Acceptable Frame ( $LAF$ )
  - Largest Frame Received ( $LFR$ )
- $LAF - LFR \leq RWS$

# Receiver Algorithm

---

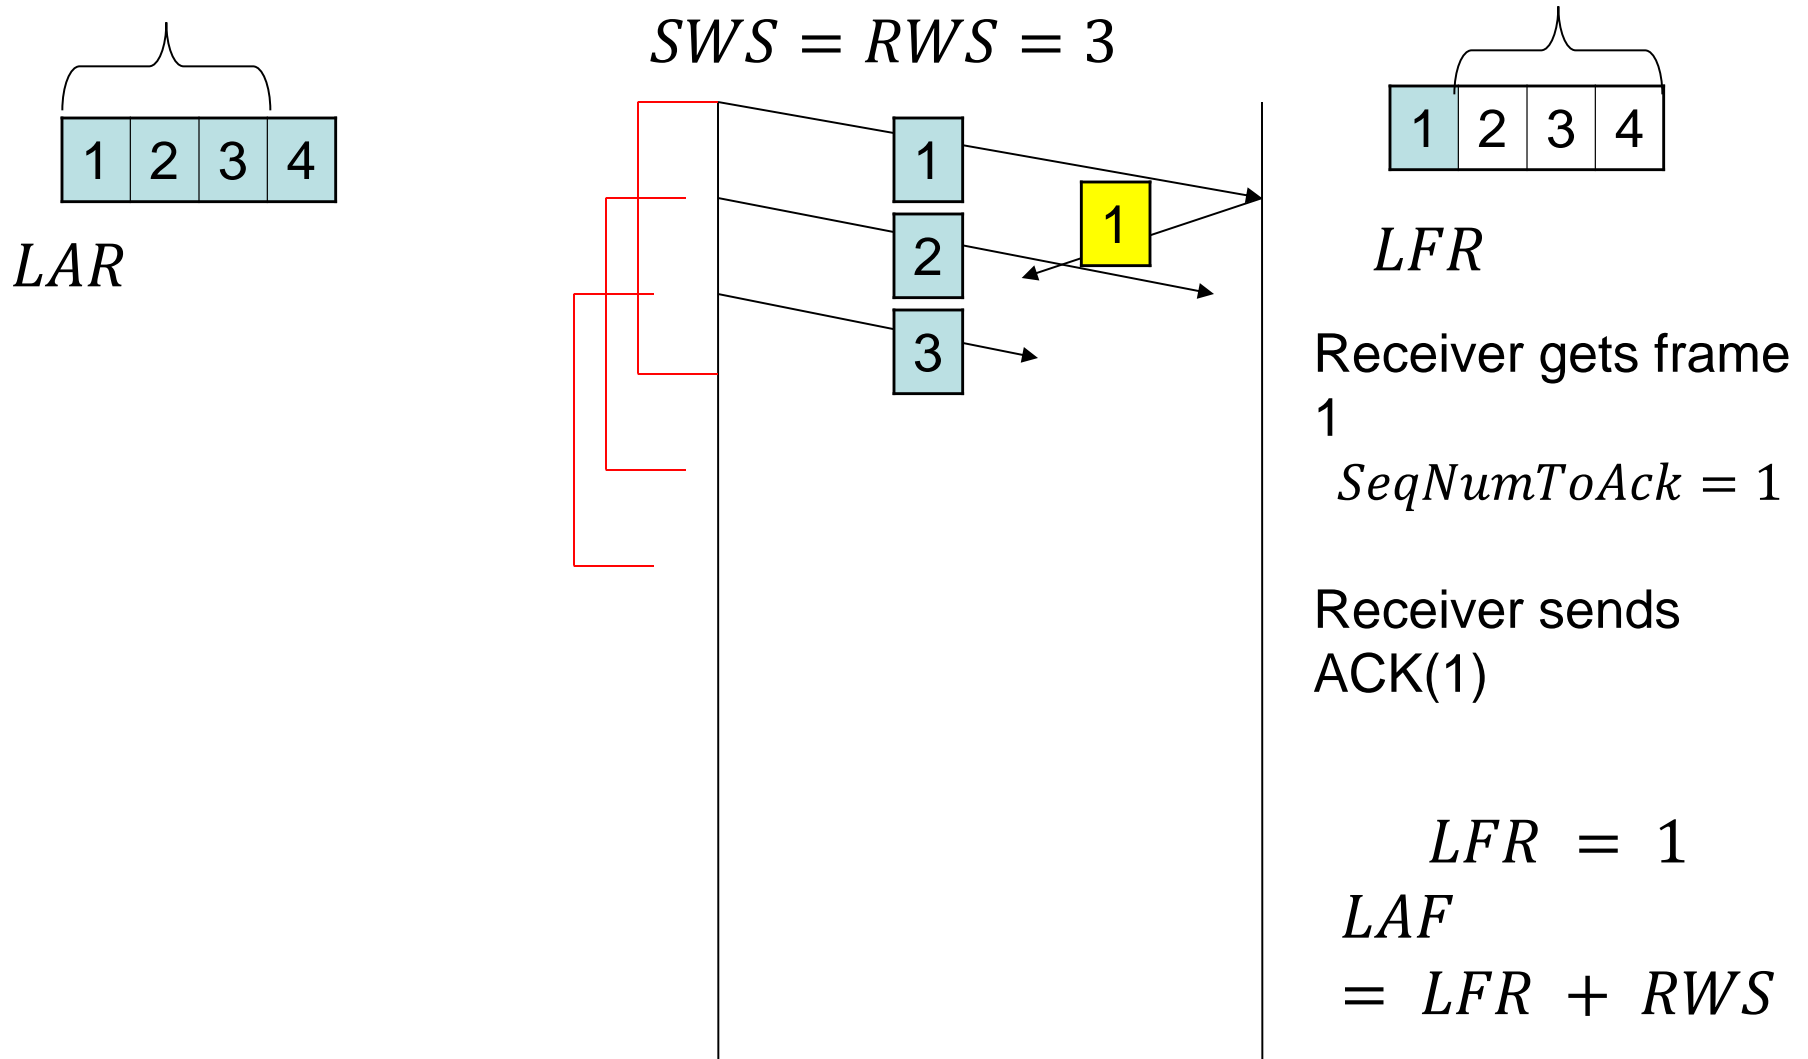
- When packet numbered  $SeqNum$  arrives
  - If  $(SeqNum \leq LFR)$  or  $(SeqNum > LAF)$  discard
  - Else accept the packet
- Define:  $SeqNumToAck$ 
  - Largest unACK'ed sequence number such that all earlier frames have been accepted
- Receiver sends  $ACK(SeqNumToAck)$
- $LFR = SeqNumToAck$
- $LAF = LFR + RWS$

# Example Sliding Window Protocol

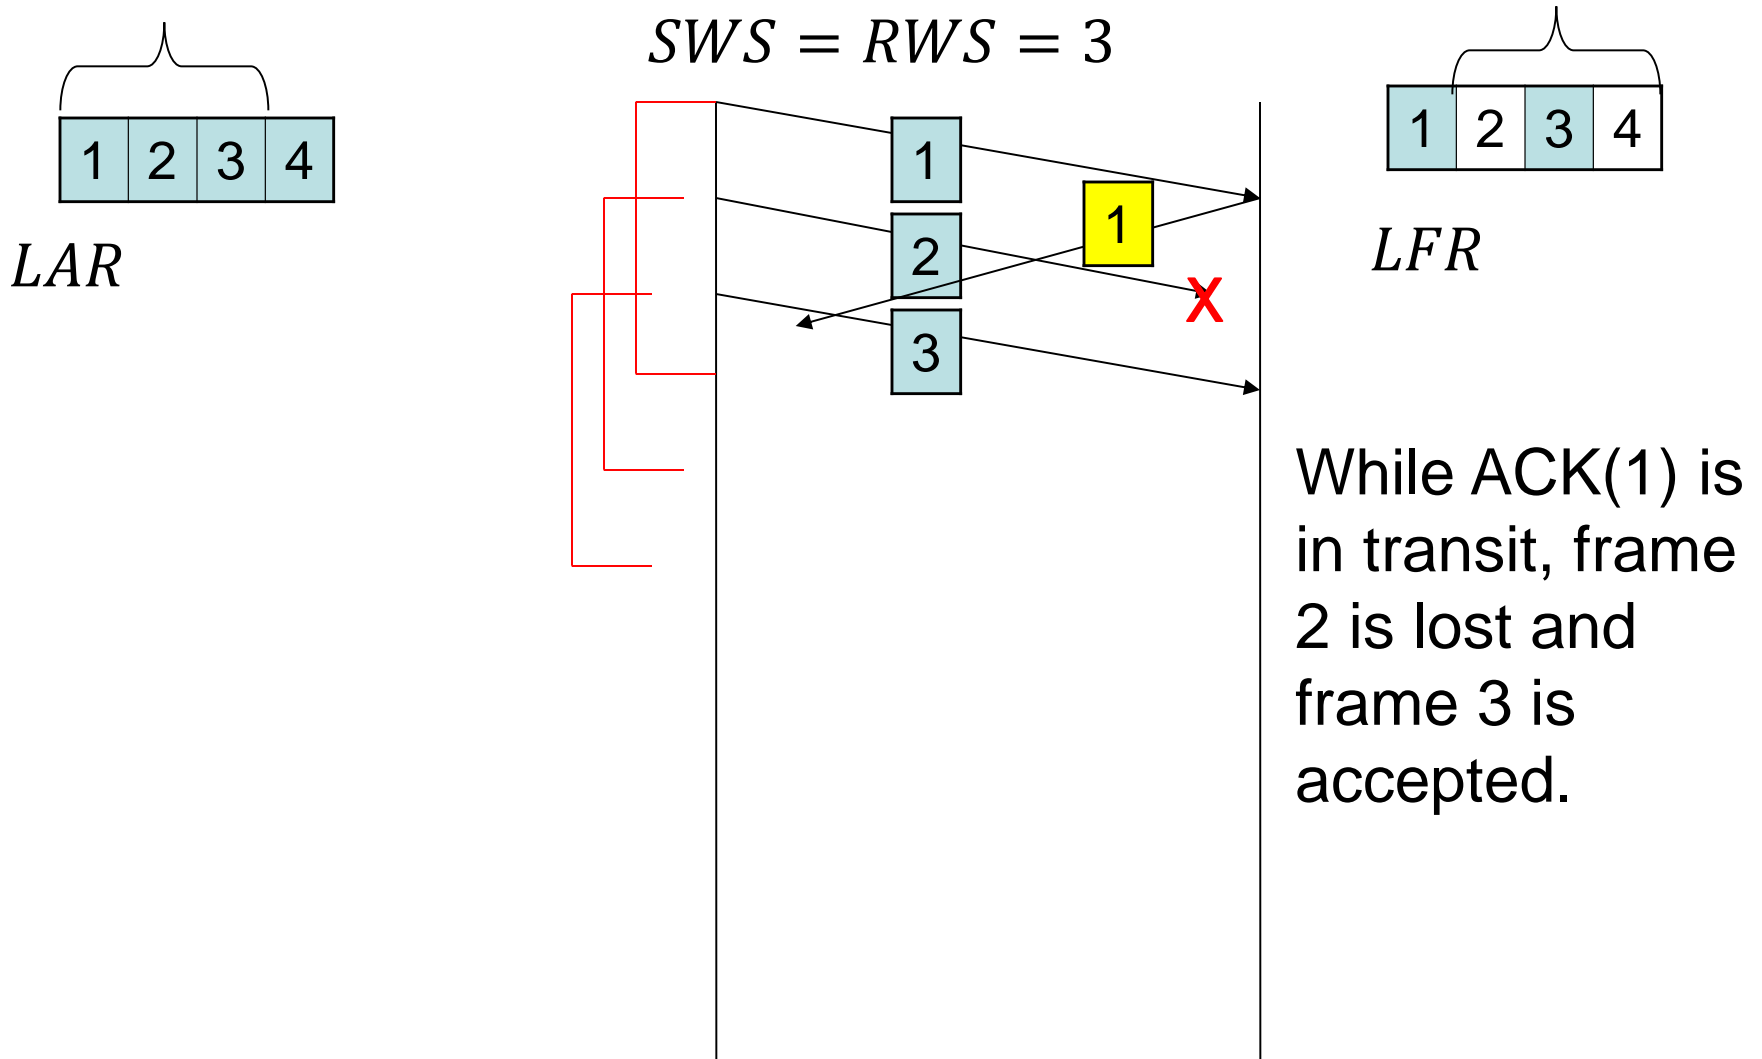




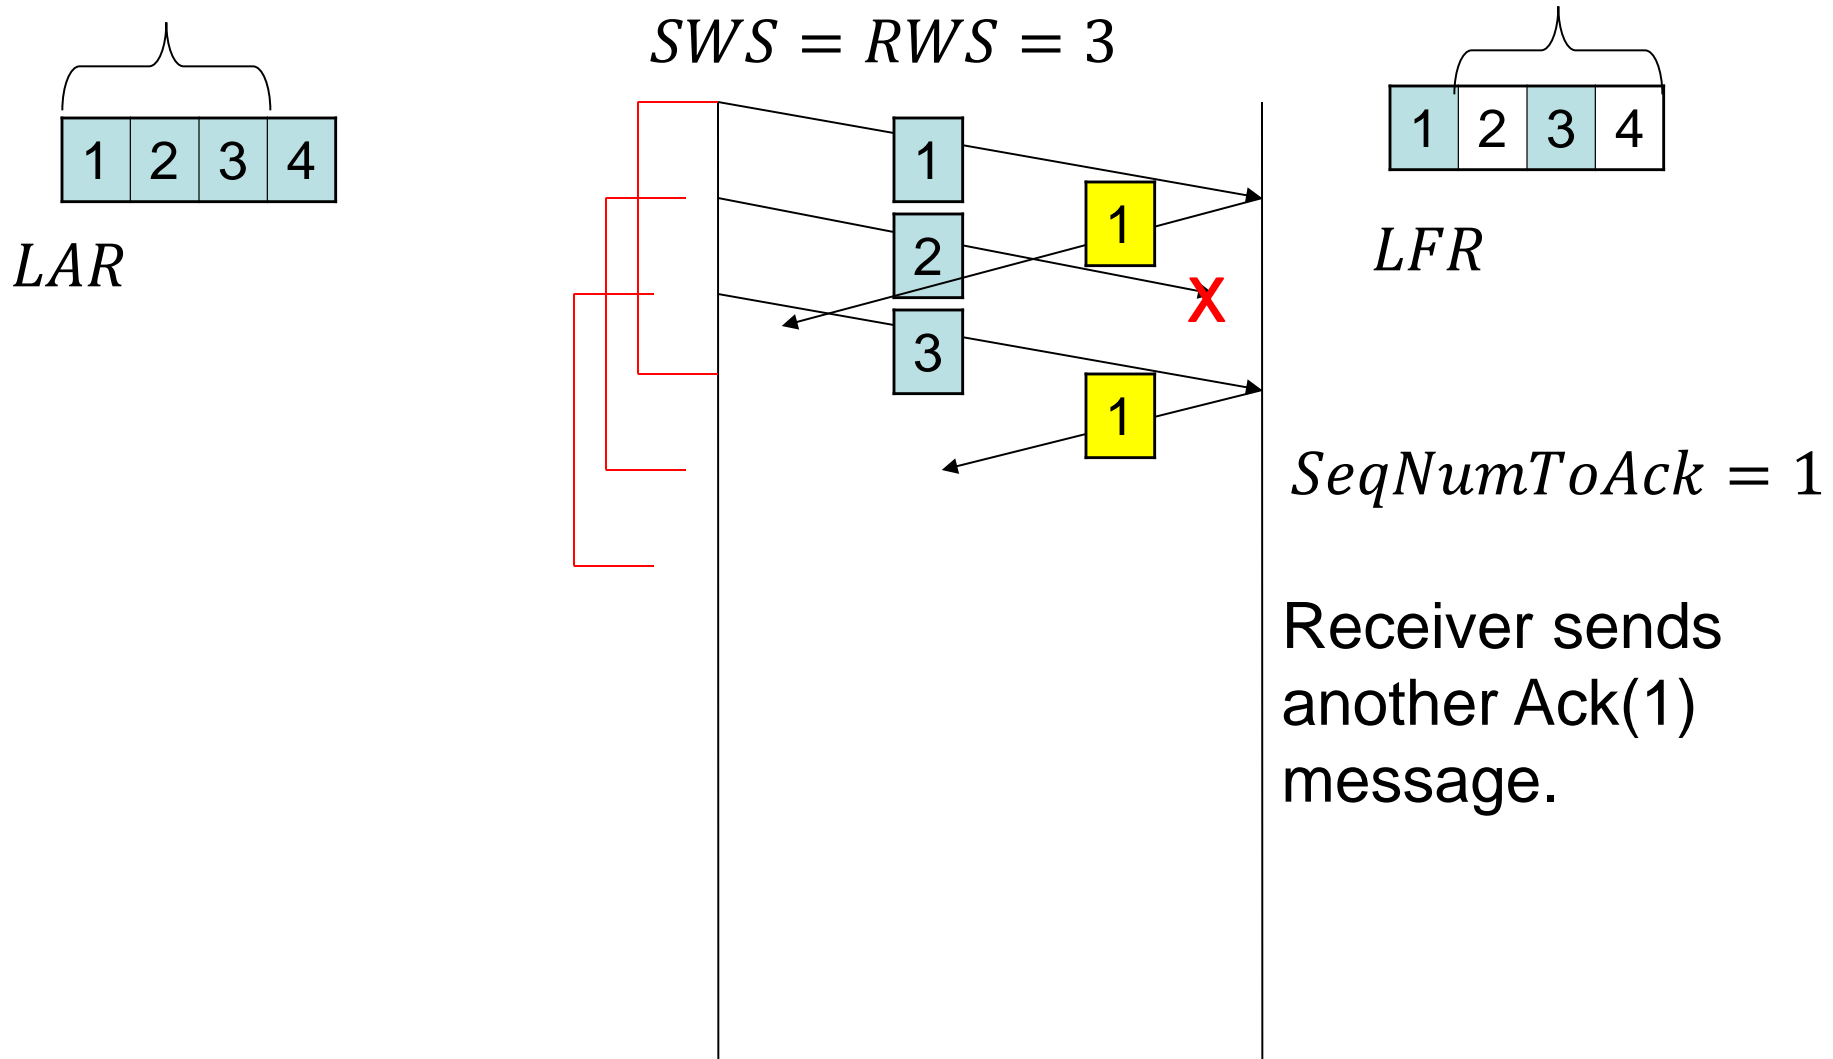
# Example Sliding Window Protocol



# Example Sliding Window Protocol

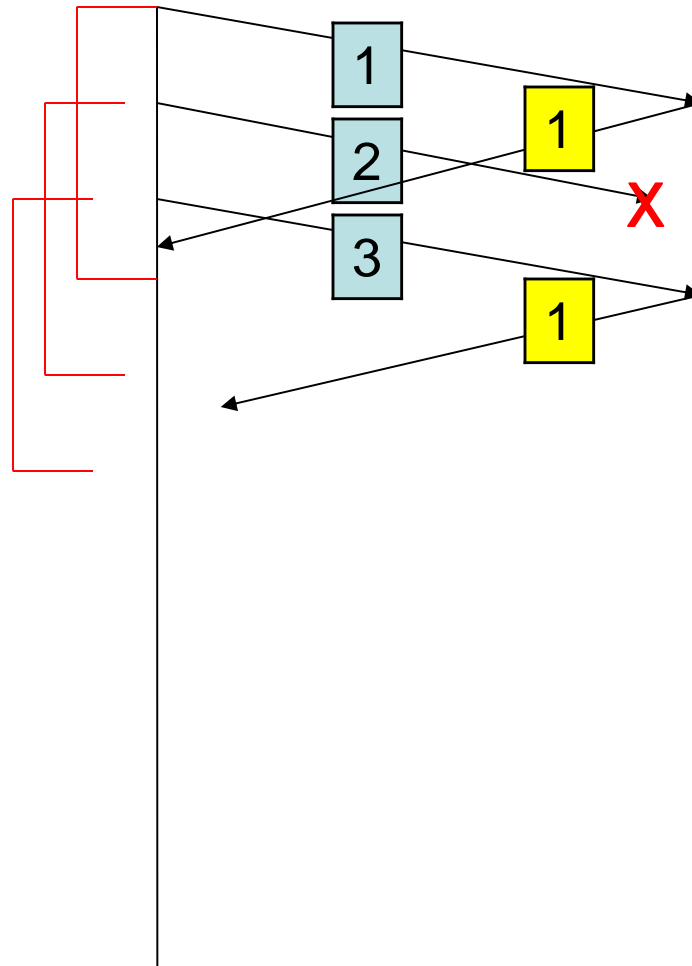




# Example Sliding Window Protocol




# Example Sliding Window Protocol

$$SWS = RWS = 3$$


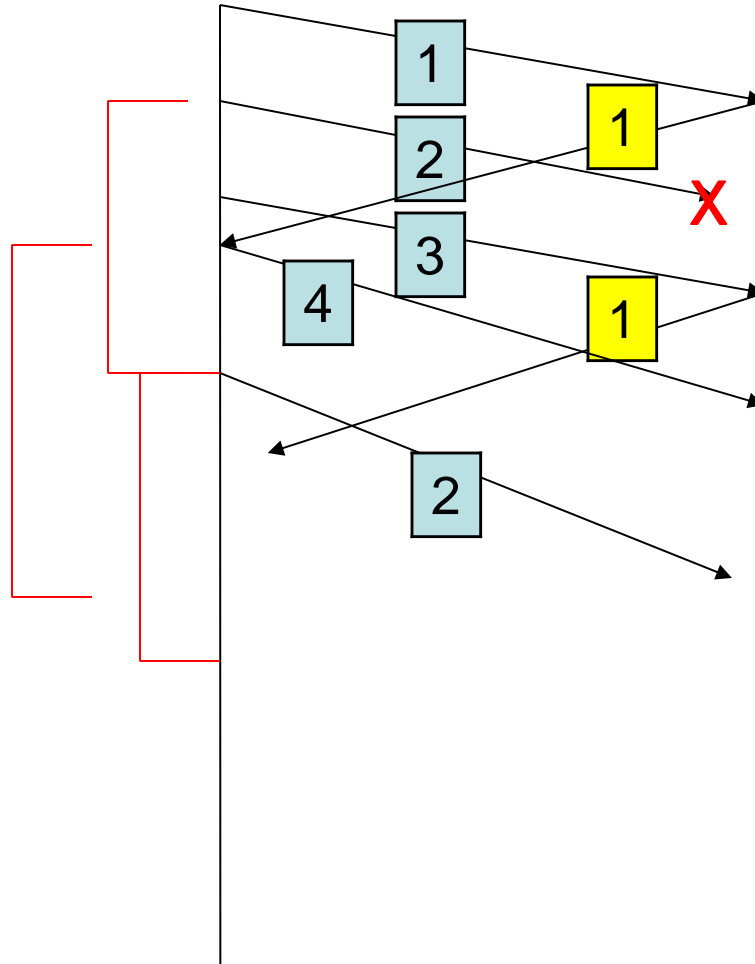


# Example Sliding Window Protocol

$$SWS = RWS = 3$$



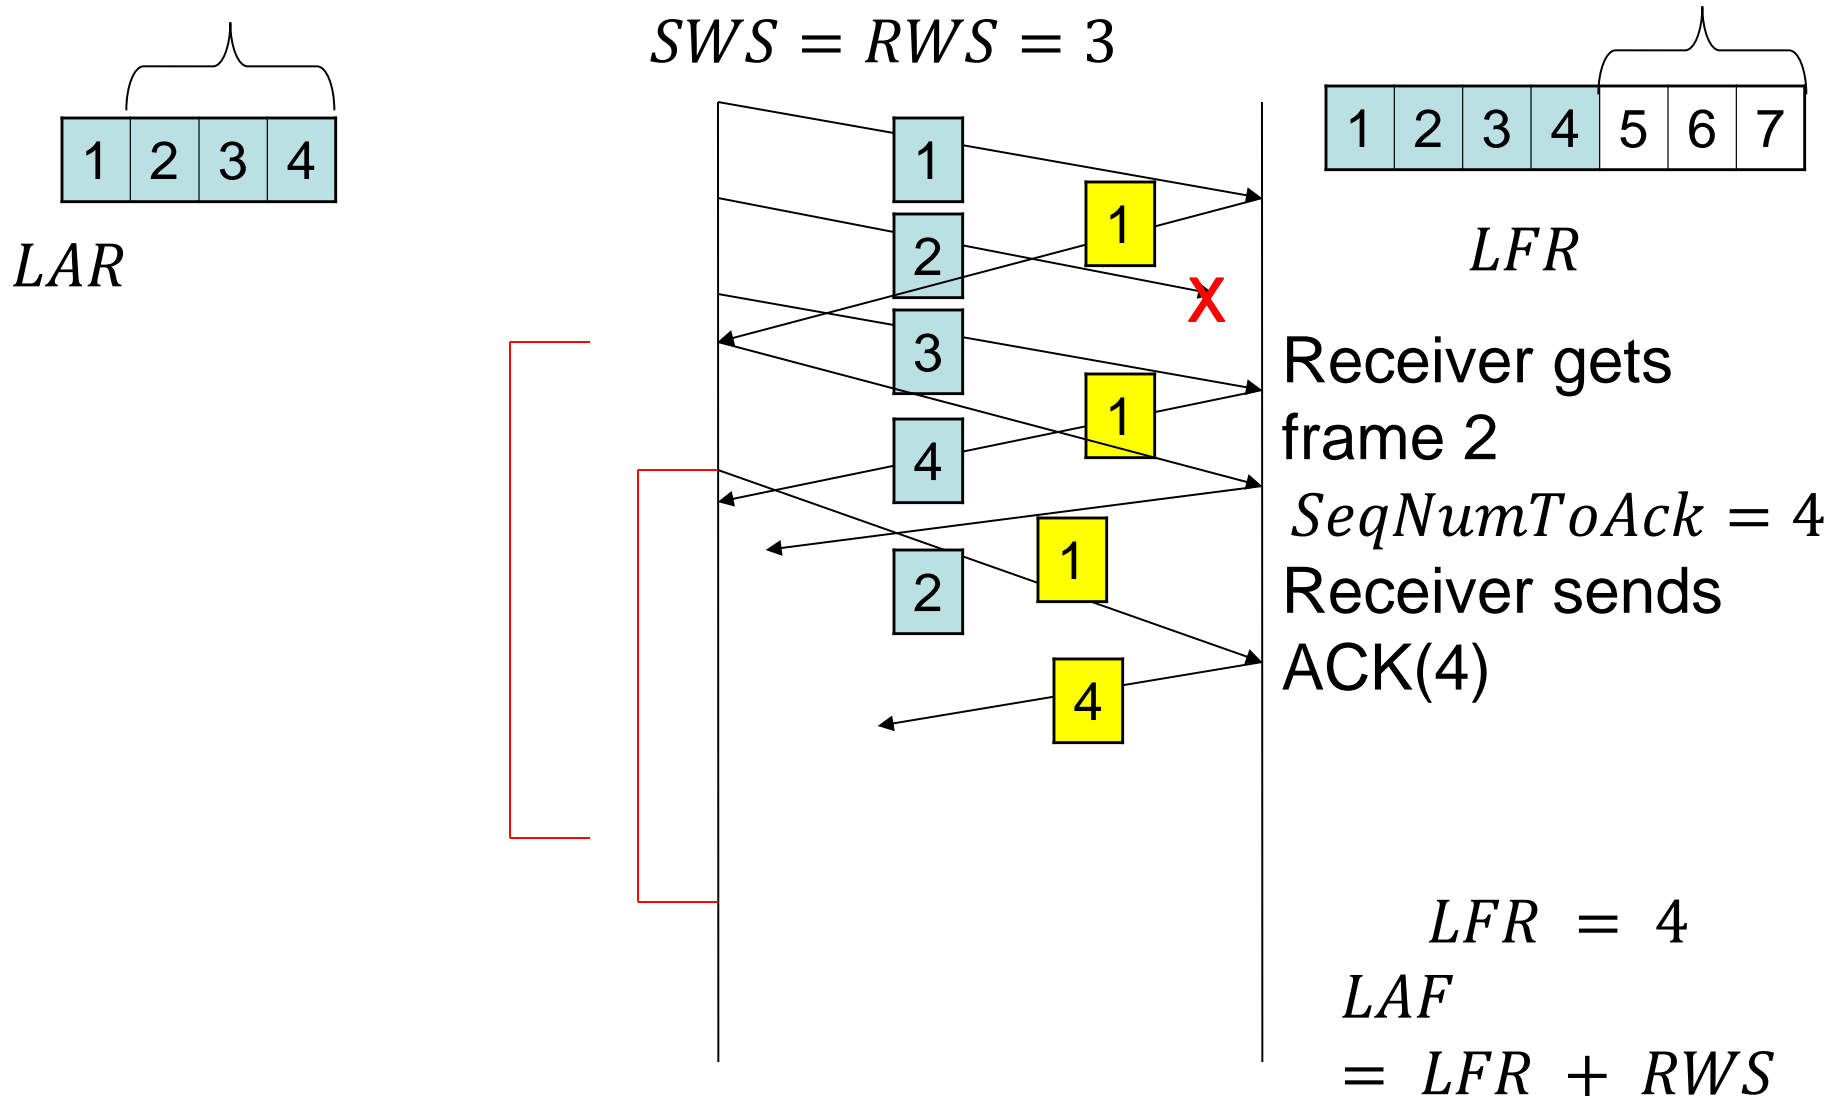
Sender transmits frame 4 and then the timer for frame 2 expires, so it resends.




---



# Example Sliding Window Protocol




# Example Sliding Window Protocol



*LAR*


$$SWS = RWS = 3$$



*LFR*

Sender gets  
ACK(1)  
again—ignores it.

Sender gets  
ACK(4)  
Sets *LAR* = 4  
Increases *LFS*





# Variants on Sliding Window

---

Receiver doesn't transmit redundant ACKs

- Just ignore out of order arrivals

Receiver transmits *selective ACKS*

- ACK indicates exactly which frames have been accepted

# What size for the window?

---

- If  $RTT \times Bandwidth$  product known, ideal is:

$$SWS = \frac{RTT}{2} \times \frac{Bandwidth}{Framesize}$$

- Common receive window size settings:

$$RWS = 1$$

- No buffering of out-of-order frames

$$RWS = SWS$$

- Buffer as many as can be in flight

Note:  $RWS > SWS$  is not sensible

# Finite Sequence Numbers

---

We've assumed infinite sequence numbers so far.

Real packets have finite size “Sequence Number” field

What do we do?



Image source: reddit.com

# What's a sufficient SeqNum Field size?

## Principle: Re-use sequence numbers

- 8-bit example
- They wrap: 0, 1, 2, ..., 254, 255, 0, 1, 2, ..., 254, 255, 0, 1, 2, ...

## Recall:

- For Stop-and-Wait we need 2 sequence values (0/1)
- 1 bit of space

What about for Sliding Window with  $X$  packets?

Suppose  $SWS = RWS$

- How many sequence numbers should there be?
- Is  $SWS + 1$  sufficient?

# Sufficient MaxSeqNum

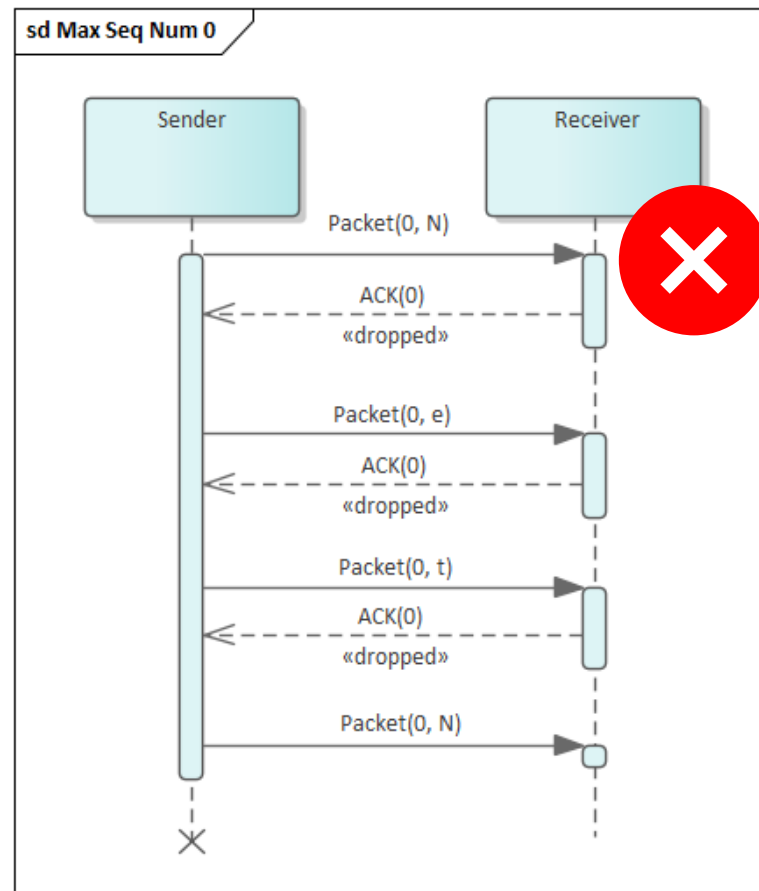
---

- Frame  $i$ 's sequence num is  $i \bmod \text{MaxSeqNum}$
- Assuming  $SWS = RWS$
- $SWS < \frac{\text{MaxSeqNum}+1}{2}$
- $\text{MaxSeqNum} > (2 \times SWS) - 1$
- Why?
  - Consider case where all the ACKS are lost.
  - Suppose  $SWS = RWS = 3$
  - $\text{MaxSeqNum} \geq 5$  since  $(0,1,2,3,4)$  are insufficient

# Sequence Number Worst Case

0	0	0	0	0	0	0	0
N	e	t	w	o	r	k	s


0	0	0	0	0	0	0	0



# Sequence Number Worst Case

0	1	0	1	0	1	0	1
N	e	t	w	o	r	k	s

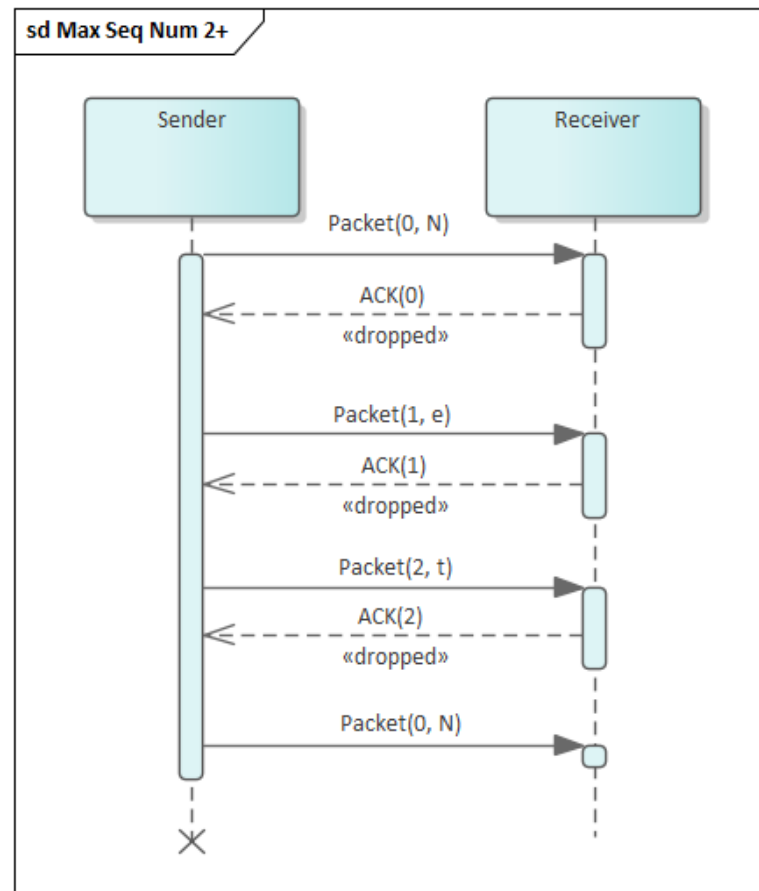
0	1	0	1	0	1	0	1



# Sequence Number Worst Case

0	1	2	0	1	2	0	1
N	e	t	w	o	r	k	s

0	1	2	0	1	2	0	1






# Sequence Number Worst Case

0	1	2	0	1	2	0	1
N	e	t	w	o	r	k	s


0	1	2	0	1	2	0	1
N							



# Sequence Number Worst Case

0	1	2	0	1	2	0	1
N	e	t	w	o	r	k	s


0	1	2	0	1	2	0	1
N	e						



# Sequence Number Worst Case

0	1	2	0	1	2	0	1
N	e	t	w	o	r	k	s

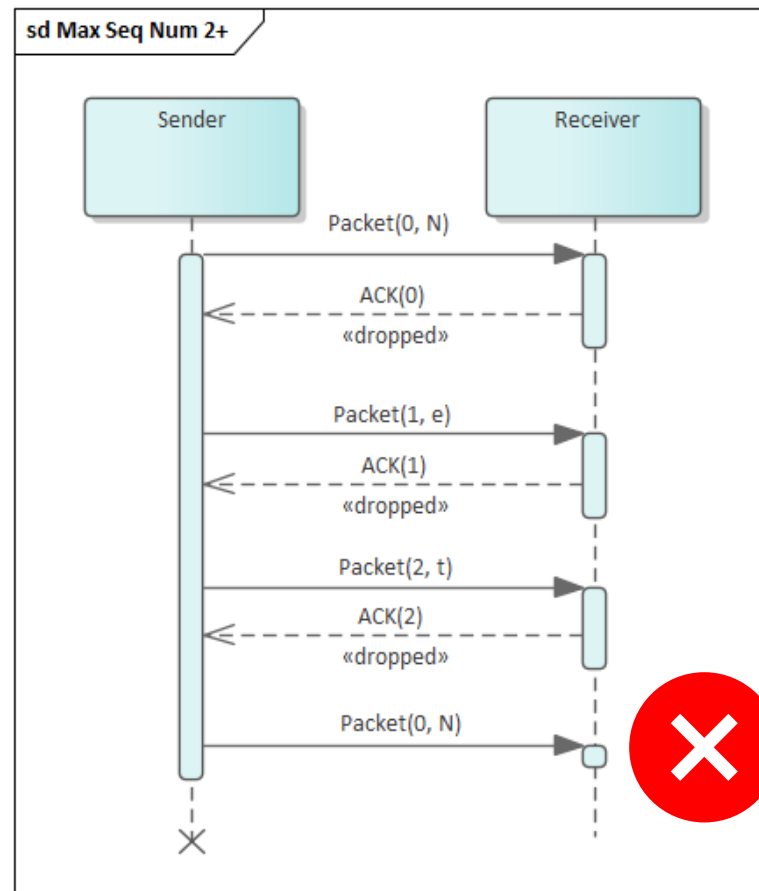
0	1	2	0	1	2	0	1
N	e	t					



# Sequence Number Worst Case

0	1	2	3	0	1	2	3
N	e	t	w	o	r	k	s


0	1	2	3	0	1	2	3
N	e	t					



# Sequence Number Worst Case

0	1	2	3	4	0	1	2
N	e	t	w	o	r	k	s


0	1	2	3	4	0	1	2
N	e	t					



# Sequence Number Worst Case

0	1	2	3	4	5	0	1
N	e	t	w	o	r	k	s

0	1	2	3	4	5	0	1
N	e	t					



# We're doing this backward

Max Sequence Number determined by the protocol design


- Fixed field size

Must select SWS and RWS accordingly

- $SWS \approx 0.5 * MaxSeqNum$
- $RWS \approx 0.5 * MaxSeqNum$
- $RWS \leq SWS$

Sequence numbers wrap fast

- Imagine a 100Gbps connection
- 12.5 GB per second
- 32 bit sequence number field ~4 billion possibilities
- Number wraps ~3 times per second





# Roles of Sliding Window Algorithm

---

## Reliable delivery

- It provides an efficient retransmission protocol for dealing with errors

## In-order delivery

- The receiver buffers frames and delivers them in sequence number order

## Flow control

- It sends ACKs back to give hints to sender
- More sophisticated version could give # of frames the receiver has room for → throttles the sender

# Sliding window in practice

---

## TCP (Transmission Control Protocol)

Uses sliding window algorithm


More complex because it's used in internet – not over a direct link

*Bandwidth* × *delay* not known

Dynamically changes timeouts

Larger buffers for in-order delivery

# Example: $SWS=RWS=3$ , 4 drops



# Conclusion

---

- ARQ
  - Stop and Wait
  - Sliding Window