

---

---

# Routing: RIP & OSPF, UDP

8 January 2025  
Lecture 9

Some Slides Credits: Steve Zdancewic (UPenn)

# Topics for Today

---

- Routing
  - Introduction and Goals
  - RIP
  - OSPF
- UDP
- Sources in PD:
  - RIP: 3.4.2
  - OSPF: 3.4.3
  - UDP: 5.1

# Routing



Images: <https://www.aaroads.com/blog/south-carolinas-new-highway-signs/#post/0>  
By de:User:Jutta234, CC BY-SA 2.5, <https://commons.wikimedia.org/w/index.php?curid=16078428>

# Routing protocols on the internet

---

## Routing information protocol (RIP)

- Distance vector routing
- Uses Bellman-Ford algorithm
- Outdated, suffers from count-to-infinity problem

## Open shortest path first (OSPF)

- Link state routing
- Runs over Layer 3 (over IP)
- Uses Dijkstra algorithm to determine shortest paths

## Intermediate System to Intermediate System IS-IS

- Link state routing – similar to OSPF
- Runs over Layer 2 (under IP)
- Uses Dijkstra's algorithm

## Border gateway protocol (BGP)

- Between networks (administrative domains, Autonomous Systems)
- Path-vector routing
- Consideration of business agreements

# Routing criteria

---

## Correctness ✓

Every packet is delivered to its destination

## Efficiency ⚙️

Choose paths with small delay and high throughput (network wide)

## Complexity ⚙️

Setting up routing tables  
Making routing decisions

## Robustness 🛡️

Cope with topology changes  
No network reboot

## Adaptiveness ⚖️

Load balancing  
and traffic control

## Fairness ⚖️

All users get the same degree of service

# Path costs, Routing metrics

## Minimum hop

Number of channels  
traversed



## Shortest path



Channels have statically  
assigned weights

Cost of a path is the sum of  
costs of the edges

Presumes no negative-cost  
cycles

## Minimum delay

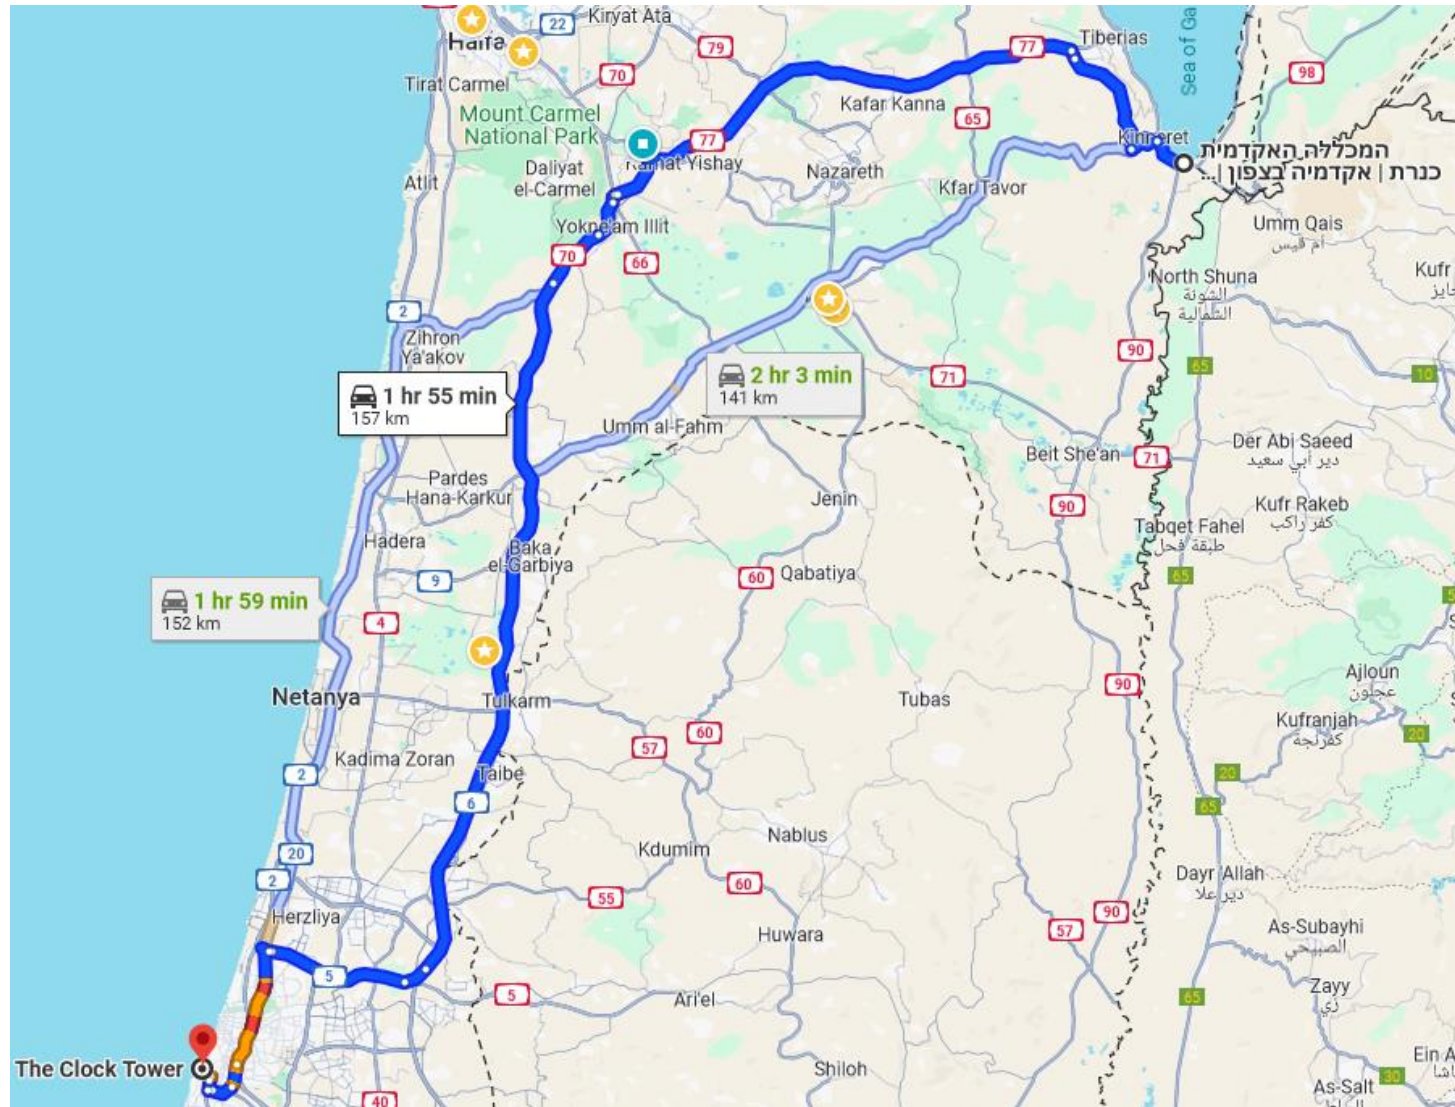


Channels have dynamically  
assigned weights based on the  
traffic on the channel

Routing tables are repeatedly  
revised such that paths with close  
to minimal delays are chosen



# Shortest vs Fastest



# So Far



---

- Routing
  - Introduction and Goals
  - RIP
  - OSPF
- UDP



# Spanning Trees (Abstractly)


- Given a connected graph  $G$
- A *spanning tree* is an acyclic, connected subgraph of  $G$  that contains all the nodes.



# Spanning Tree Algorithm (Abstractly)

---

- Pick a root node
  - Compute shortest paths to root
  - Need to break ties




# Distance Vector Algorithm (RIP)

---


- Similar to the Spanning Tree Algorithm
  - Except that information about distance to ALL nodes is forwarded (not just info. about root.)
  - Sometimes called **Bellman-Ford algorithm**
- Each node constructs a *Distance Vector*
  - Contains distances (costs) to reach all other node
  - Initially:
    - Distance to neighbors (a simplification for now) = 1
    - Distance to others =  $\infty$
  - Routing table reflects node's beliefs

# Example Network Graph



A's initial information:

Dest	Cost	NextHop
B	1	B
C	1	C
D	$\infty$	-
E	1	E
F	1	F
G	$\infty$	-



# Iteration Steps

Each host  
sends its  
DV (cost) to  
its  
neighbors

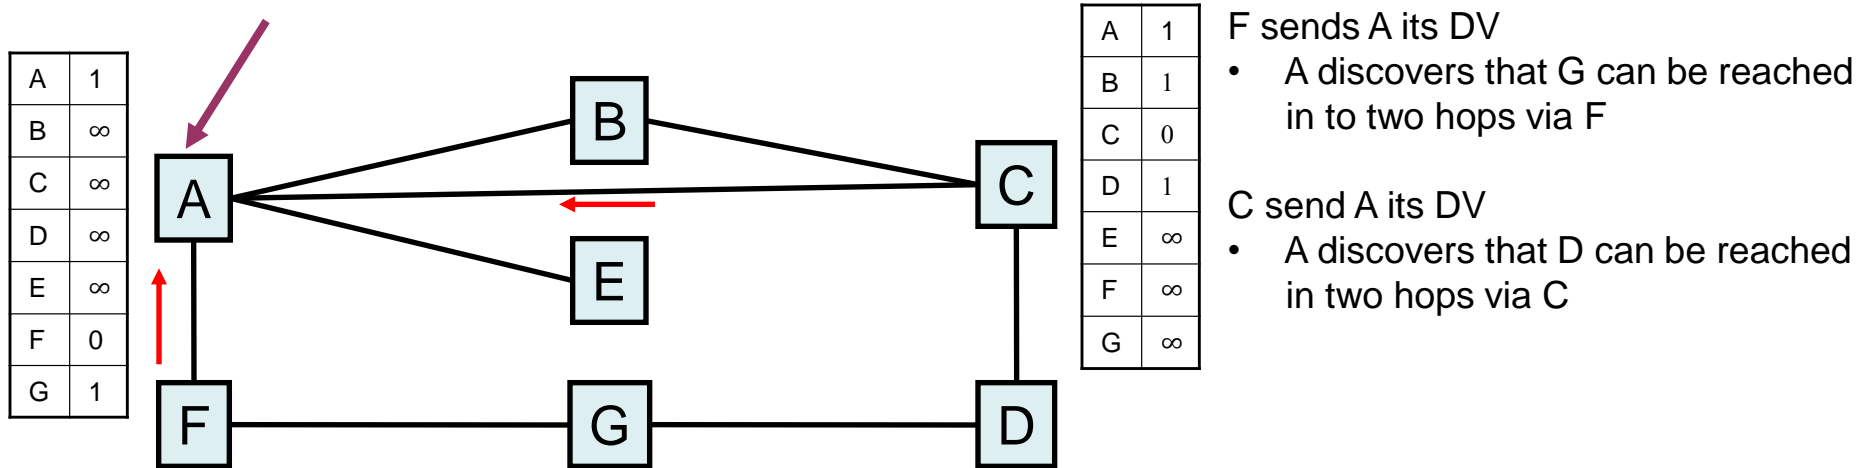
Neighbors  
update distance  
vectors and  
routing  
information  
accordingly.

- Ignore worse information
- Update any better routes

If host  
changed  
its tables,  
send new  
DV to  
neighbors

After a few  
iterations,  
routing  
information  
*converges*

# Example Iteration Steps



Dest	Cost	NextHop
B	1	B
C	1	C
D	$\infty$	-
E	1	E
F	1	F
G	$\infty$	-

Dest	Cost	NextHop
B	1	B
C	1	C
D	$\infty$	-
E	1	E
F	1	F
G	2	F

Dest	Cost	NextHop
B	1	B
C	1	C
D	2	C
E	1	E
F	1	F
G	2	F

# Details

---



- Note: No single host has all routing information.
- When to send update vectors?
  - When your routing table changes (triggered)
  - Periodically (“I’m alive!”)
- Detecting link/node failure
  - (1) Periodically exchange “I’m alive!” messages.
  - (2) Timeout mechanism
- In a **static network**, if all weights are 1, once A discovers a path to B ( $\text{cost} < \infty$ ), **no subsequent round will ever discover a better path to B**
  - All weights are 1 is called “Hop Count” (Default RIP action)



# Dynamic Networks


What about a dynamic network?

- A node enters
  - The new neighbors will eventually inform everyone else (easy)
- A node/link fails
  - Is the network partitioned?
  - How do nodes discover that a node/link is gone?



# Link Failure Example 1


A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	2	F	G	2	D	G	1	G	G	1	G



A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	2	F	G	2	D	G	1	G	G	$\infty$	-

# Link Failure Example 1

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	2	F	G	2	D	G	1	G	G	$\infty$	-




Case 1:  
1. F sends to A

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	$\infty$	-	G	2	D	G	1	G	G	$\infty$	-

# Link Failure Example 1

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	$\infty$	-	G	2	D	G	1	G	G	$\infty$	-




Case 1:

1. F sends to A
2. A sends to C
  - (nothing)
3. C sends to A

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	3	C	G	2	D	G	1	G	G	$\infty$	-

# Link Failure Example 1

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	3	C	G	2	D	G	1	G	G	$\infty$	-




Case 1:

1. F sends to A
2. A sends to C
  - (nothing)
3. C sends to A
4. A sends to F

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	3	C	G	2	D	G	1	G	G	4	A

# Link Failure Example 2

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	2	F	G	2	D	G	1	G	G	$\infty$	-




Case 2:  
1. A sends to F

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	2	F	G	2	D	G	1	G	G	3	A

# Link Failure Example 2

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	2	F	G	2	D	G	1	G	G	3	A



Case 2:


1. A sends to F
2. F sends to A

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	4	F	G	2	D	G	1	G	G	3	A



# Link Failure Example 2

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	4	F	G	2	D	G	1	G	G	3	A




Case 2:

1. A sends to F
2. F sends to A
3. A sends to F

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	4	F	G	2	D	G	1	G	G	5	A

# Link Failure Example 2

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	4	F	G	2	D	G	1	G	G	5	A



Case 2:

1. A sends to F
2. F sends to A
3. A sends to F
4. ...
5. C sends to A
6. A sends to F

A			C			D			F		
Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop	Dest	C	Next hop
A	-	-	A	1	A	A	2	C	A	1	A
C	1	C	C	-	-	C	1	C	C	2	A
D	2	C	D	1	D	D	-	-	D	2	G
F	1	F	F	2	A	F	2	G	F	-	-
G	3	C	G	2	D	G	1	G	G	4	A

# Network Partition

---

- The loop between A and F is broken only by C's update
  - C has **real information** about another path to G
- What if C's update **never** comes? What if C's information is **false**?
  - The network is partitioned –OR–
  - G is completely offline
- The other nodes will continue counting until infinity (or the distance field reaches its max)

# Three RIP Solutions

---

## Option 1: Infinity is small



- Don't let the distance fields go above a predefined maximum
    - Say 15, 20, etc. (15 is the actual value)
    - The max must be greater than the diameter of the network
- 

## Option 2: Don't send routes to the one who sent it to you



Example: F receives an update from A about a path to G.


- In the example before, F didn't know that's A's "path" went via F
- A doesn't send F an advertisement to get to G since it goes via F
- Called "split horizon"
- "Split horizon with poison reverse"- A sends route for G to F, but with infinite distance

# Three RIP Solutions

---

## Option 3: Route Poisoning and Holddowns

- Instead of just stopping to announce a lost destination, a router sends a poison message for the destination
  - E.g. the router announces distance 16 ( $\equiv \infty$ ) for the destination
- For a fixed period of time (say 3 minutes) any new offers for the destination via the sender will be rejected
  - E.g. Wait until everyone has heard the poison and then accept new offers
- Gives a chance for the route to be deleted by everyone.
- Cisco specific



# RIP v1 (1986) Details

## Classful routing only

- No subnet masks
- Assumes everyone has same mask

Supports multiple address families

## Commands

- Request: Ask for someone to send their info
- Response: Respond with routing table

## Updates sent

- Timer about 30 seconds between updates
- Due to changes

Split Horizon, Poison Reverse, Small Infinity

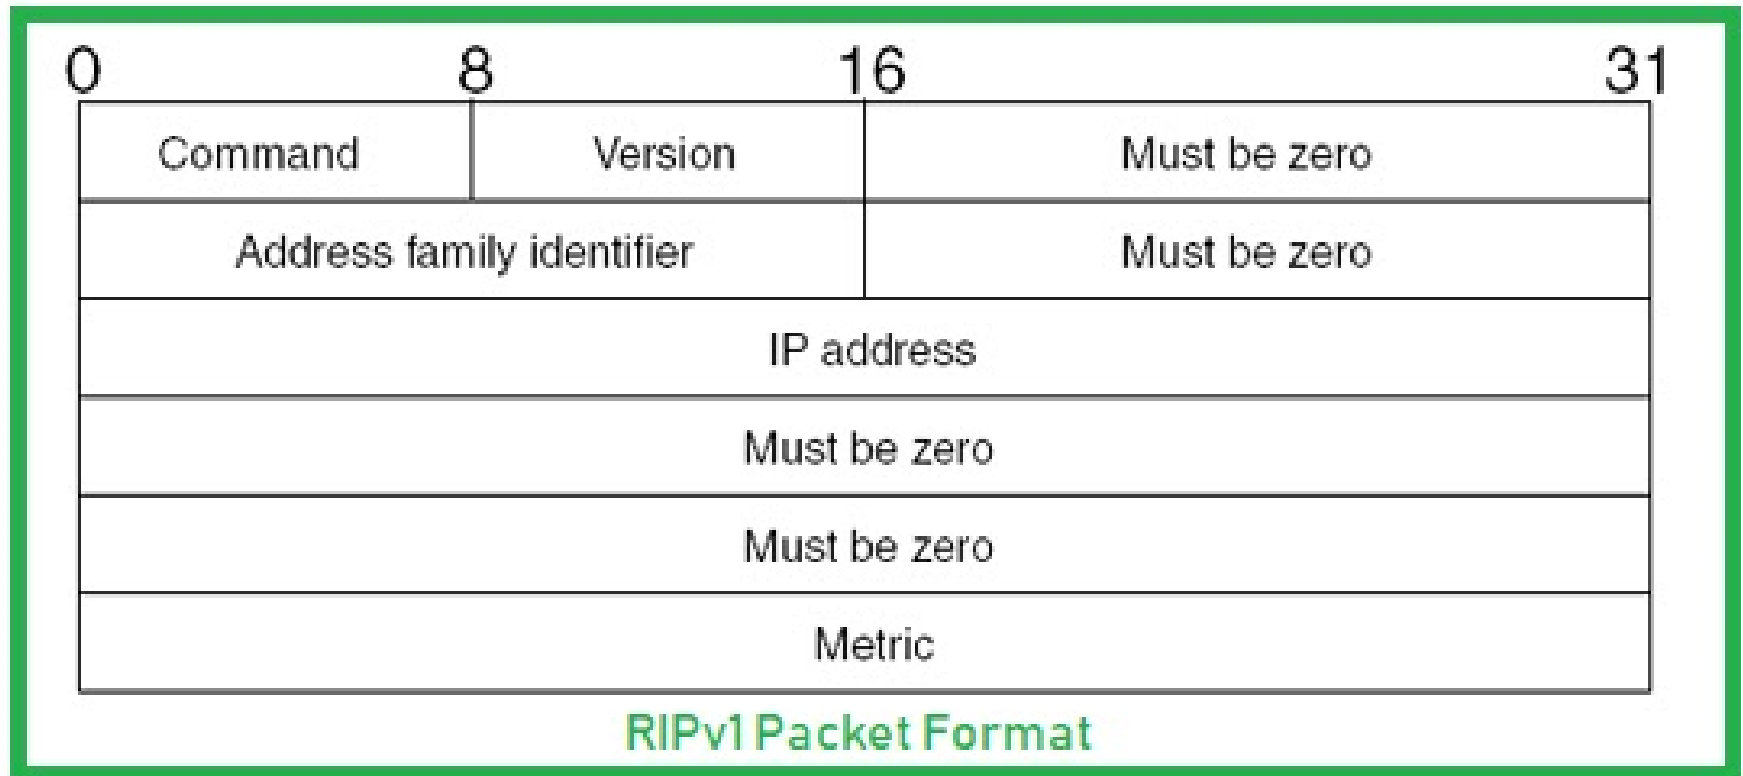
## Takes shortest path

- If two are shortest, splits traffic on them equally

## Route deletion

- If no update in 180 seconds or reaches infinity
- Marked as deleted
- Lives 120 seconds longer to send to others

# RIP v1 (1986) Details



<https://www.networkurge.com/2020/05/ripv1.html>



# RIP v2 (1998) Details

## Added authentication



- Simple password
- Address family ID 0xFFFF

## Route tag for exterior routing

## Next hop

- 0.0.0.0 for the sender
- Another address must be on the subnet, allows just some routers to run RIP

## Message sent via Multicast 224.0.0.9



# So Far

---

- Routing
  - Introduction and Goals
  - RIP
  - OSPF
- UDP

# Open Shortest Path First (OSPF)

---

## OSPFv2 (1998)

- RFC 2328
- IPv4 ✓
- In protocol authentication of routers and completeness of messages
- Atoms are networks and subnets
- Routers identified by Router ID and/or IP address

## OSPFv3 (2008)

- RFC 5340
- IPv4 + IPv6 ✓
- Authentication at IP layer (IPSec)
- Completeness at IP layer
- Atoms are links (may have multiple subnets)
- Router IDs and IP addresses are disjoint

# Open Shortest Path First (OSPF)

---

- Each node sends a *reliable flood* of information to all other nodes
  - In v3, can limit the flooding scope
- **Link-State Packets (LSPs)** contain

ID of the node that created the packet

List of (neighbor, cost) pairs associated with the source node

Sequence Number (64bits—no wrapping)

Time To Live (ensure old info is eventually removed)


Authentication (v2)  
– password or message authenticity code (MAC)

# OSPF Steps

---

## Reliable Flooding

- Adjacent routers reliably send updates (ACKs)
- Source sends to all neighbors
- Recipients
  - Send to all neighbors except the one it got the message from
  - Ignores duplicates



## Local Calculations

1. Once all of the link-state info has been flooded each node has complete network topology
2. Compute routing information using Dijkstra's shortest-path algorithm
3. Periodic updates and failure detection are like RIP.

# OSPF Features

## Authentication of routing messages (v2)

- Misconfigured or malicious host could advertise bad route info (i.e. reach anywhere in 0 hops)
- Prevent routers from accidentally joining a network
- Prevent malicious or accidental changes to route information
- (Eventually added to RIP too.)

## Additional Hierarchy

- Partitions domains into *areas*
- Reduces transmission & storage overhead

## Load Balancing

- Multiple routes with same cost
- Traffic evenly distributed

# Dijkstra's Algorithm

---

Node has two lists – Confirmed and Tentative - pairs of (Destination, Cost, Next-Hop)


## Algorithm:

1. Initialize Confirmed with an entry for self with cost 0.
2. For the node just added to Confirmed in the previous step, (*Next*), examine its Link State Packet (LPS).
  - a) If *Neigh* isn't on Confirmed or Tentative, add (*Neigh*, *cost*, *NH*) to Tentative, where *NH* is the way to reach *Next* in Confirmed.
  - a) If *Neigh* is on Tentative and *cost* is better than the old cost, replace it with (*Neigh*, *cost*, *NH*), where *NH* is the way to reach *Next* in Confirmed.
3. For each *Neigh* of *Next*, calculate the distance to *Neigh* via *Next*  
 $cost = (self \rightarrow Next) + (Next \rightarrow Neigh)$ 
  - a) If *Neigh* isn't on Confirmed or Tentative, add (*Neigh*, *cost*, *NH*) to Tentative, where *NH* is the way to reach *Next* in Confirmed.
  - a) If *Neigh* is on Tentative and *cost* is better than the old cost, replace it with (*Neigh*, *cost*, *NH*), where *NH* is the way to reach *Next* in Confirmed.
4. If Tentative is empty, **stop**.
5. Choose the node on Tentative with the lowest cost, move it to Confirmed, and go to step 2.




# Dijkstra Example

---




# Dijkstra Example



Step	Confirmed	Tentative	Comments
1	(D,0,-)		Since D is the only new member of the confirmed list, look at its LSP.
2	(D,0,-)	(B,11,B) (C,2,C)	D's LSP says we can reach B through B at cost 11, which is better than anything else on either list, so put it on Tentative list; same for C.
3	(D,0,-) (C,2,C)	(B,11,B)	Put lowest-cost member of Tentative (C) onto Confirmed list. Next, examine LSP of newly confirmed member (C).
4	(D,0,-) (C,2,C)	(B,5,C) (A,12,C)	Cost to reach B through C is 5, so replace (B,11,B). C's LSP tells us that we can reach A at cost 12.
5	(D,0,-) (C,2,C) (B,5,C)	(A,12,C)	Move lowest-cost member of Tentative (B) to Confirmed, then look at its LSP.
6	(D,0,-) (C,2,C) (B,5,C)	(A,10,C)	Since we can reach A at cost 5 through B, replace the Tentative entry.
7	(D,0,-) (C,2,C) (B,5,C) (A,10,C)		Move lowest-cost member of Tentative (A) to Confirmed, and we are all done.

# Dijkstra Example 2



Step	Confirmed	Tentative
1	(A,0,-)	
2	(A,0,-)	(B,1,B) (D,5,D)
3	(A,0,-) (B,1,B)	(D,4,B) (C,7,B)
4	(A,0,-) (B,1,B) (D,4,B)	(C,5,B) (E,7,B)
5	(A,0,-) (B,1,B) (D,4,B) (C,5,B)	(E,6,B)
6	(A,0,-) (B,1,B) (D,4,B) (C,5,B) (E,6,B)	

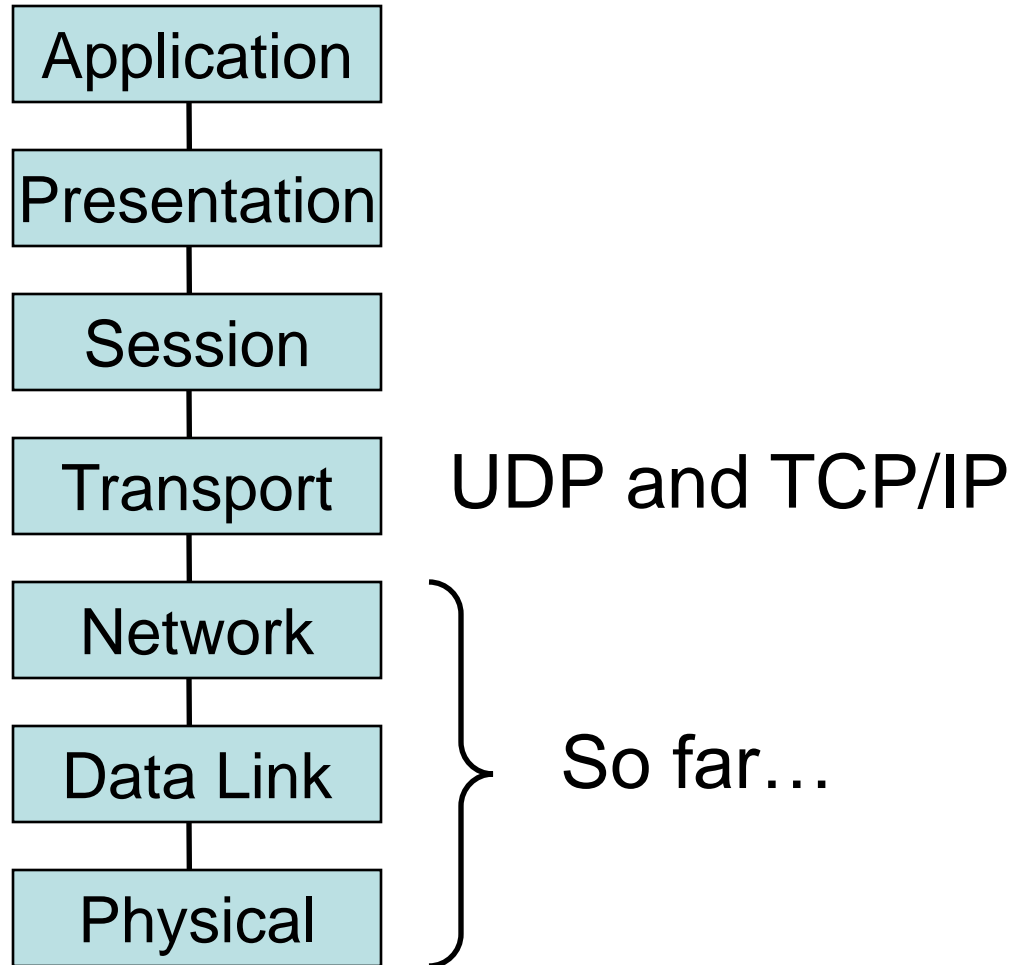
# So Far

---

- Routing
  - Introduction and Goals
  - RIP
  - OSPF
- UDP

# Protocol Stack Revisited


---



# Application vs. Network

Application Needs	Network Characteristics
Reliable, Ordered, Single-Copy Message Delivery	Drops, Duplicates and Reorders Messages
Arbitrarily large messages	Finite message size (MTU)
Timely receipt of messages	Arbitrary delay
Supports multiple applications per host	Host-level addressing and delivery mechanisms (no ability to discern conversations)
Flow control by Receiver	No end-to-end flow control mechanisms

# User Datagram Protocol (UDP)

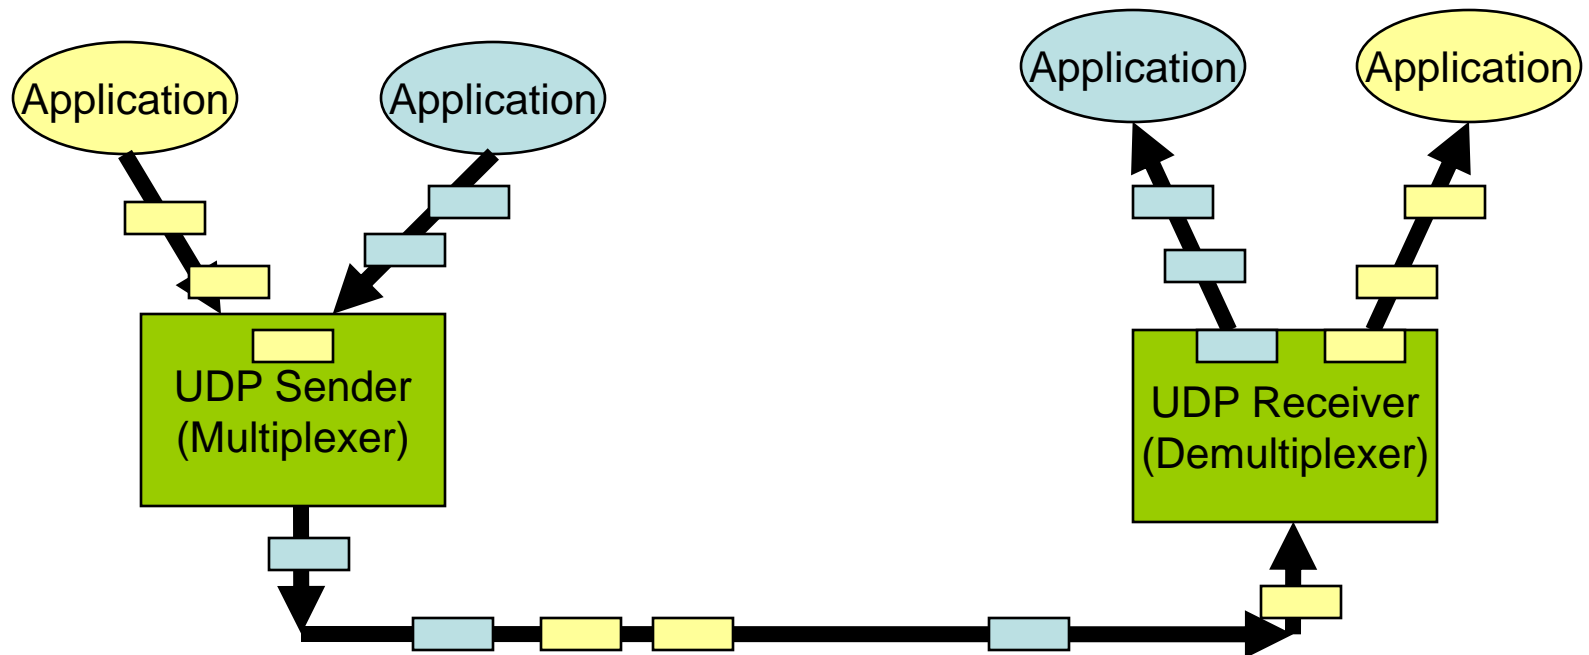


- Simplest transport-layer protocol
- Just exposes IP packet functionality to application level
- *Ports* identify sending/receiving process
  - Demultiplexing information
  - (port, host) pair identifies a network process

**Question:** Why is there another length field?

# UDP End-to-End Model

- Multiplexing/Demultiplexing with Port number





# Using Ports

---

- Client contacts Server at a *well-known port*
  - SMTP: port 25
  - DNS: port 53
  - POP3: port 110
  - HTTP: port 80
  - Unix talk : port 517
  - In unix, ports are listed in /etc/services
  - In Windows: C:\Windows\System32\drivers\etc\services
- Client and Server agree on a different port for subsequent communication
- Ports are an abstraction
  - Implemented differently on different OS's
  - Typically a message queue

# Conclusion

---

- Routing
  - Introduction and Goals
  - RIP
  - OSPF
- UDP